



The OpenGL ES® Shading Language

Language Version: 3.20
Document Revision: 2
6 August 2015

Editor: Robert J. Simpson, Qualcomm

OpenGL GLSL editor: John Kessenich, LunarG
GLSL version 1.1 Authors: John Kessenich, Dave Baldwin, Randi Rost



Copyright (c) 2013-2015 The Khronos Group Inc. All Rights Reserved.

This specification is protected by copyright laws and contains material proprietary to the Khronos Group, Inc. It or any components may not be reproduced, republished, distributed, transmitted, displayed, broadcast, or otherwise exploited in any manner without the express prior written permission of Khronos Group. You may use this specification for implementing the functionality therein, without altering or removing any trademark, copyright or other notice from the specification, but the receipt or possession of this specification does not convey any rights to reproduce, disclose, or distribute its contents, or to manufacture, use, or sell anything that it may describe, in whole or in part.

Khronos Group grants express permission to any current Promoter, Contributor or Adopter member of Khronos to copy and redistribute UNMODIFIED versions of this specification in any fashion, provided that NO CHARGE is made for the specification and the latest available update of the specification for any version of the API is used whenever possible. Such distributed specification may be reformatted AS LONG AS the contents of the specification are not changed in any way. The specification may be incorporated into a product that is sold as long as such product includes significant independent work developed by the seller. A link to the current version of this specification on the Khronos Group website should be included whenever possible with specification distributions.

Khronos Group makes no, and expressly disclaims any, representations or warranties, express or implied, regarding this specification, including, without limitation, any implied warranties of merchantability or fitness for a particular purpose or non-^o©-infringement of any intellectual property. Khronos Group makes no, and expressly disclaims any, warranties, express or implied, regarding the correctness, accuracy, completeness, timeliness, and reliability of the specification. Under no circumstances will the Khronos Group, or any of its Promoters, Contributors or Members or their respective partners, officers, directors, employees, agents, or representatives be liable for any damages, whether direct, indirect, special or consequential damages for lost revenues, lost profits, or otherwise, arising from or in connection with these materials.

Khronos, Vulkan, SYCL, SPIR, WebGL, EGL, COLLADA, StreamInput, OpenVX, OpenKCam, glTF, OpenKODE, OpenVG, OpenWF, OpenSL ES, OpenMAX, OpenMAX AL, OpenMAX IL and OpenMAX DL are trademarks and WebCL is a certification mark of the Khronos Group Inc. OpenCL is a trademark of Apple Inc. and OpenGL and OpenML are registered trademarks and the OpenGL ES and OpenGL SC logos are trademarks of Silicon Graphics International used under license by Khronos. All other product names, trademarks, and/or company names are used solely for identification and belong to their respective owners.

Table of Contents

1 Introduction.....	9
1.1 Changes.....	9
1.1.1 Changed from GLSL ES 3.2 revision 1.....	9
1.1.2 Changes from GLSL ES 3.1 revision 4.....	9
1.1.3 Changes from GLSL ES 3.1 revision 3.....	9
1.1.4 Changes from GLSL ES 3.1 revision 2.....	10
1.1.5 Changes from GLSL ES 3.1 revision 1.....	10
1.1.6 Changes from GLSL ES 3.0:.....	11
1.2 Overview.....	11
1.3 Error Handling.....	11
1.4 Typographical Conventions.....	12
1.5 Compatibility.....	13
2 Overview of OpenGL ES Shading.....	14
2.1 Vertex Processor.....	14
2.2 Tessellation Control Processor.....	14
2.3 Tessellation Evaluation Processor.....	15
2.4 Geometry Processor.....	15
2.5 Fragment Processor.....	15
2.6 Compute Processor.....	15
3 Basics.....	17
3.1 Character Set.....	17
3.2 Source Strings.....	17
3.3 Version Declaration.....	18
3.4 Preprocessor.....	19
3.5 Comments.....	23
3.6 Tokens.....	24
3.7 Keywords.....	24
3.8 Identifiers.....	26
3.9 Definitions.....	27
3.9.1 Static Use.....	27
3.9.2 Uniform and Non-Uniform Control Flow.....	27
3.9.3 Dynamically Uniform Expressions.....	28
3.10 Logical Phases of Compilation.....	28
4 Variables and Types.....	30
4.1 Basic Types.....	30
4.1.1 Void.....	33
4.1.2 Booleans.....	33
4.1.3 Integers.....	34
4.1.4 Floats.....	35
4.1.5 Vectors.....	37

4.1.6 Matrices.....	37
4.1.7 Opaque Types.....	37
4.1.7.1 Samplers.....	38
4.1.7.2 Images.....	38
4.1.7.3 Atomic Counters.....	38
4.1.8 Structures.....	39
4.1.9 Arrays.....	40
4.2 Scoping.....	43
4.2.1 Definition of Terms.....	43
4.2.2 Types of Scope.....	43
4.2.3 Redeclaring Names.....	45
4.2.4 Global Scope.....	47
4.2.5 Shared Globals.....	47
4.3 Storage Qualifiers.....	47
4.3.1 Default Storage Qualifier.....	48
4.3.2 Constant Qualifier.....	48
4.3.3 Constant Expressions.....	49
4.3.4 Input Variables.....	49
4.3.5 Uniform Variables.....	52
4.3.6 Output Variables.....	53
4.3.7 Buffer Variables.....	55
4.3.8 Shared Variables.....	56
4.3.9 Interface Blocks.....	57
4.4 Layout Qualifiers.....	62
4.4.1 Input Layout Qualifiers.....	64
4.4.1.1 Tessellation Evaluation Inputs.....	66
4.4.1.2 Geometry Shader Inputs.....	68
4.4.1.3 Fragment Shader Inputs.....	69
4.4.1.4 Compute Shader Inputs.....	69
4.4.2 Output Layout Qualifiers.....	70
4.4.2.1 Tessellation Control Outputs.....	72
4.4.2.2 Geometry Outputs.....	72
4.4.2.3 Fragment Outputs.....	73
4.4.3 Uniform Variable Layout Qualifiers.....	74
4.4.4 Uniform and Shader Storage Block Layout Qualifiers.....	75
4.4.5 Opaque Uniform Layout Qualifiers.....	77
4.4.6 Atomic Counter Layout Qualifiers.....	78
4.4.7 Format Layout Qualifiers.....	79
4.5 Interpolation Qualifiers.....	80
4.6 Parameter Qualifiers.....	81
4.7 Precision and Precision Qualifiers.....	81
4.7.1 Range and Precision.....	81

4.7.2	Conversion between precisions.....	84
4.7.3	Precision Qualifiers.....	84
4.7.4	Default Precision Qualifiers.....	85
4.7.5	Available Precision Qualifiers.....	87
4.8	Variance and the Invariant Qualifier.....	88
4.8.1	The Invariant Qualifier.....	88
4.8.2	Invariance Within a Shader.....	89
4.8.3	Invariance of Constant Expressions.....	90
4.8.4	Invariance of Undefined Values.....	90
4.9	The Precise Qualifier.....	91
4.10	Memory Access Qualifiers.....	93
4.11	Order of Qualification.....	96
4.12	Empty Declarations.....	96
5	Operators and Expressions.....	97
5.1	Operators.....	97
5.2	Array Operations.....	98
5.3	Function Calls.....	98
5.4	Constructors.....	98
5.4.1	Conversion and Scalar Constructors.....	98
5.4.2	Vector and Matrix Constructors.....	99
5.4.3	Structure Constructors.....	101
5.4.4	Array Constructors.....	102
5.5	Vector Components.....	102
5.6	Matrix Components.....	104
5.7	Structure and Array Operations.....	104
5.8	Assignments.....	105
5.9	Expressions.....	106
5.10	Vector and Matrix Operations.....	109
5.11	Evaluation of Expressions.....	110
6	Statements and Structure.....	111
6.1	Function Definitions.....	112
6.1.1	Function Calling Conventions.....	114
6.2	Selection.....	115
6.3	Iteration.....	116
6.4	Jumps.....	117
7	Built-in Variables.....	119
7.1	Built-in Language Variables.....	119
7.1.1	Vertex Shader Special Variables.....	119
7.1.2	Tessellation Control Shader Special Variables.....	120
7.1.2.1	Tessellation Control Input Variables.....	120
7.1.2.2	Tessellation Control Output Variables.....	120
7.1.3	Tessellation Evaluation Shader Special Variables.....	120

7.1.3.1 Tessellation Evaluation Input Variables.....	121
7.1.3.2 Tessellation Evaluation Output Variables.....	121
7.1.4 Geometry Shader Special Variables.....	121
7.1.4.1 Geometry Shader Input Variables.....	122
7.1.4.2 Geometry Shader Output Variables.....	122
7.1.5 Fragment Shader Special Variables.....	123
7.1.6 Compute Shader Special Variables.....	126
7.2 Built-In Constants.....	127
7.3 Built-In Uniform State.....	129
7.4 Redeclaring Built-in Blocks.....	130
8 Built-in Functions.....	131
8.1 Angle and Trigonometry Functions.....	133
8.2 Exponential Functions.....	134
8.3 Common Functions.....	135
8.4 Floating-Point Pack and Unpack Functions.....	140
8.5 Geometric Functions.....	143
8.6 Matrix Functions.....	145
8.7 Vector Relational Functions.....	146
8.8 Integer Functions.....	147
8.9 Texture Functions.....	149
8.9.1 Texture Query Functions.....	150
8.9.2 Texel Lookup Functions.....	151
8.9.3 Texture Gather Functions.....	157
8.10 Atomic-Counter Functions.....	159
8.11 Atomic Memory Functions.....	160
8.12 Image Functions.....	161
8.13 Geometry Shader Functions.....	165
8.14 Fragment Processing Functions.....	165
8.14.1 Interpolation Functions.....	167
8.15 Shader Invocation Control Functions.....	168
8.16 Shader Memory Control Functions.....	169
9 Shader Interface Matching.....	172
9.1 Input Output Matching by Name in Linked Programs.....	172
9.2 Matching of Qualifiers.....	173
9.2.1 Linked Shaders.....	174
9.2.2 Separable Programs.....	175
10 Shading Language Grammar.....	176
11 Errors.....	189
11.1 Preprocessor Errors.....	189
11.2 Lexer/Parser Errors.....	189
11.3 Semantic Errors.....	189
11.4 Linker.....	191

12 Counting of Inputs and Outputs.....	193
13 Issues.....	196
13.1 Compatibility with OpenGL ES 2.0.....	196
13.2 Convergence with OpenGL.....	196
13.3 Numeric Precision.....	196
13.4 Floating Point Representation and Functionality.....	197
13.5 Precision Qualifiers.....	198
13.6 Function and Variable Name Spaces.....	201
13.7 Local Function Declarations and Function Hiding.....	202
13.8 Overloading main().....	202
13.9 Error Reporting.....	202
13.10 Structure Declarations.....	202
13.11 Embedded Structure Definitions.....	203
13.12 Redefining Built-in Functions.....	203
13.13 Global Scope.....	204
13.14 Constant Expressions.....	204
13.15 Varying Linkage.....	204
13.16 gl_Position.....	205
13.17 Preprocessor.....	205
13.18 Character set.....	206
13.19 Line Continuation.....	207
13.20 Phases of Compilation.....	207
13.21 Maximum Number of Varyings.....	207
13.22 Array Declarations.....	209
13.23 Invariance.....	209
13.24 Invariance Within a shader.....	211
13.25 While-loop Declarations.....	212
13.26 Cross Linking Between Shaders.....	212
13.27 Visibility of Declarations.....	212
13.28 Language Version.....	213
13.29 Samplers.....	213
13.30 Dynamic Indexing.....	213
13.31 Maximum Number of Texture Units.....	214
13.32 On-target Error Reporting.....	214
13.33 Rounding of Integer Division.....	214
13.34 Undefined Return Values.....	214
13.35 Precisions of Operations.....	215
13.36 Compiler Transforms.....	216
13.37 Expansion of Function-like Macros in the Preprocessor.....	216
13.38 Should Extension Macros be Globally Defined?.....	216
13.39 Minimum Requirements.....	217
13.40 Packing Functions.....	217

13.41 Boolean logical vector operations.....	217
13.42 Range Checking of literals.....	218
13.43 Sequence operator and constant expressions.....	218
13.44 Version Directive.....	219
13.45 Use of Unsigned Integers.....	219
13.46 Vertex Attribute Aliasing.....	220
13.47 Does a vertex input Y collide with a fragment uniform Y?.....	221
13.48 Counting Rules for Flat and Smooth Varyings.....	222
13.49 Array of Arrays: Ordering of Indices.....	222
13.50 Precision of Evaluation of Compile-time Expressions.....	223
13.51 Matching of Memory Qualifiers in Function Parameters.....	224
14 Acknowledgments.....	226
15 Normative References.....	228

1 Introduction

This document specifies only version 3.2 of the OpenGL ES Shading Language. It requires `__VERSION__` to substitute 320, and requires `#version` to accept only 320 es. If `#version` is declared with a smaller number, the language accepted is a previous version of the shading language, which will be supported depending on the version and type of context in the OpenGL ES API. See the OpenGL ES Graphics System Specification, Version 3.2, for details on what language versions are supported.

All OpenGL ES Graphics System Specification references in this specification are to version 3.2

1.1 Changes

This specification is derived from GLSL ES 3.1 revision 4.

1.1.1 Changed from GLSL ES 3.2 revision 1

- Signed zeros must be supported

1.1.2 Changes from GLSL ES 3.1 revision 4

- Added the following extensions:
 - [KHR_blend_equation_advanced](#)
 - [OES_sample_variables](#)
 - [OES_shader_image_atomic](#)
 - [OES_shader_multisample_interpolation](#)
 - [OES_texture_storage_multisample_2d_array](#)
 - [OES_geometry_shader](#)
 - [OES_gpu_shader5](#)
 - [OES_primitive_bounding_box](#)
 - [OES_shader_io_blocks](#)
 - [OES_tessellation_shader](#)
 - [OES_texture_buffer](#)
 - [OES_texture_cube_map_array](#)
 - [KHR_robustness](#)

1.1.3 Changes from GLSL ES 3.1 revision 3

- Added section on the memory and execution modeling

- Clarified that **barrier()** does not order memory transactions
- Added GL_FRAGMENT_PRECISION_HIGH macro
- Matching of memory access qualifiers when calling a function
- Matching rules for anonymous blocks
- Opaque uniforms and uniform block members are not initialized to 0
- Matching rules for qualifiers in shader interfaces
- Precision of evaluation of constant expressions
- Clarified input/output counting rules for separable programs
- Removed memory qualifiers from formal parameters in built-in functions
- Clarified that atomic_uint cannot be used inside a structure

1.1.4 Changes from GLSL ES 3.1 revision 2

- Source character set is now UTF-8
- Built-in functions with an out parameter are not constant expressions
- Accessing packed uniform or shader storage buffer from multiple shader stages is an error
- Use of double underscore in macro names
- ldexp, frexp behavior for boundary conditions
- Invariant pragma allowed in fragment shaders
- Usage of *interface-qualifiers*
- Clarified that GLSL ES 3.1 shaders cannot be linked with shaders declaring a previous version
- Atomic counters must be highp
- Fragment outputs cannot be or contain array of arrays
- Array of arrays syntax extended
- glMaxFragmentAtomicCounterBuffers changed to '0'

1.1.5 Changes from GLSL ES 3.1 revision 1

- Updated 'Logical Phases of Compilation'
- Assignment to runtime-sized arrays disallowed.
- Replaced glMaxCombinedImageUnitsAndFragmentOutputs with gl_MaxCombinedShaderOutputResources
- Atomic functions: memory parameters are qualified with **coherent**
- Lexing and parsing errors must be reported at compile-time
- Removed Image atomics

1.1.6 Changes from GLSL ES 3.0:

Removed:

- (None)

Added:

- Compute shaders
- Shader storage buffer objects
- Arrays of arrays
- Atomic counters
- Images
- Separate program objects (also known as separate shader objects)
- Explicit uniform locations
- Texture gather
- Bitfield operations
- Integer mix function

1.2 Overview

This document describes *The OpenGL ES Shading Language, version 3.20*

Independent compilation units written in this language are called *shaders*. A *program* is a set of shaders that are compiled and linked together. The aim of this document is to thoroughly specify the programming language. The OpenGL ES Graphics System Specification will specify the OpenGL ES entry points used to manipulate and communicate with programs and shaders.

1.3 Error Handling

Compilers, in general, accept programs that are ill-formed, due to the impossibility of detecting all ill-formed programs. Portability is only ensured for well-formed programs, which this specification describes. Compilers are encouraged to detect ill-formed programs and issue diagnostic messages, but are not required to do so for all cases. The compilation process is implementation-dependent but is generally split into a number of stages, each of which occurs at one of the following times:

- A call to *glCompileShader*
- A call to *glLinkProgram*
- A draw call or a call to *glValidateProgram*

The implementation should report errors as early as possible but in any case must satisfy the following:

- All lexical, grammatical and semantic errors must have been detected following a call to *glLinkProgram*

- Errors due to mismatch between the shaders (link errors) must have been detected following a call to `glLinkProgram`
- Errors due to exceeding resource limits must have been detected following any draw call or a call to `glValidateProgram`
- A call to `glValidateProgram` must report all errors associated with a program object given the current GL state.

Where the specification uses the terms ***required***, ***must/must not***, ***does/does not***, ***disallowed*** or ***not supported***, the compiler or linker is required to detect and report any violations. Similarly when a condition or situation is an ***error***, it must be reported. Use of any feature marked as ***reserved*** is an error. Where the specification uses the terms ***should/should not*** or ***undefined behavior*** there is no such requirement but compilers are encouraged to report possible violations.

A distinction is made between ***undefined behavior*** and an ***undefined value*** (or ***result***). Undefined behavior includes system instability and/or termination of the application. It is expected that systems will be designed to handle these cases gracefully but specification of this is outside the scope of OpenGL ES.

If a value or result is undefined, the system may behave as if the value or result had been assigned a random value. For example, an undefined `gl_Position` may cause a triangle to be drawn with a random size and position. The value may not be consistent. For example an undefined boolean value may cause both sub-statements in an if-then-else statement to be executed (see section 4.8.4 “Invariance of Undefined Values”). The implementation may also detect the generation and/or use of undefined values and behave accordingly (for example causing a trap). Undefined values must not by themselves cause system instability. However undefined values may lead to other more serious conditions such as infinite loops or out of bounds array accesses.

Implementations may not in general support functionality beyond the mandated parts of the specification without use of the relevant extension. The only exceptions are:

1. If a feature is marked as optional.
2. Where a maximum value is stated (e.g. the maximum number of vertex outputs), the implementation may support a higher value than that specified.

Where the implementation supports more than the mandated specification, off-target compilers are encouraged to issue warnings if these features are used.

The compilation process is split between the compiler and linker. The allocation of tasks between the compiler and linker is implementation dependent. Consequently there are many errors which may be detected either at compile or link time, depending on the implementation.

1.4 Typographical Conventions

Italic, bold, and font choices have been used in this specification primarily to improve readability. Code fragments use a fixed width font. Identifiers embedded in text are italicized. Keywords embedded in text are bold. Operators are called by their name, followed by their symbol in bold in parentheses. The clarifying grammar fragments in the text use bold for literals and italics for non-terminals. The official grammar in section 9 “Shading Language Grammar” uses all capitals for terminals and lower case for non-terminals.

1.5 Compatibility

The OpenGL ES 3.2 API is designed to work with GLSL ES v1.00, GLSL ES 3.00, GLSL ES 3.10 and GLSL ES 3.20. In general a shader written for versions prior to OpenGL ES 3.2 should work without modification in OpenGL ES 3.2.

When porting applications from an earlier to later version of the API, the following points should be noted:

- Not all language constructs present in earlier versions of the language are available in later versions e.g. attribute and varying qualifiers are present in v1.00 but not v3.00. However, the functionality of GLSL ES 3.20 is a super-set of GLSL ES 3.10.
- Some features of later versions of the API require language features that are not present in earlier version of the language.
- It is an error to link shaders if they are written in different versions of the language.
- The OpenGL ES 2.0 and 3.0 APIs do not support shaders written in GLSL ES 3.2.
- Using GLSL ES 1.00 shaders within OpenGL ES 3.x may extend the resources available beyond the minima specified in GLSL ES 1.0. Shaders which make use of this will not necessarily run on an OpenGL ES 2.0 implementation: Similarly for GLSL ES 3.00 shaders running within OpenGL ES 3.2.

Uniforms

The number of uniforms specified by `gl_MaxVertexUniformVectors` and returned by the corresponding API query is the same for GLSL ES versions 1.00 and 3.x when used as part of OpenGL ES 3.2.

Varyings, vertex outputs and fragment inputs

These are specified differently in the two versions of the language and may be different. For GLSL ES 1.00, the maximum number of varyings is specified by `gl_MaxVaryingVectors`. For GLSL ES 3.00 and GLSL ES 3.10, the maximum number of vertex outputs and fragment inputs is independently specified by `gl_MaxVertexOutputVectors` and `gl_MaxFragmentInputVectors`.

In GLSL ES 1.00, only varyings which are statically used in both the vertex and fragment shaders are counted. This applies when GLSL ES 1.00 is used in OpenGL ES 3.2.

Multiple Render Targets

Although `gl_FragData` is declared as an array in GLSL ES 1.00, multiple render targets are not supported in OpenGL ES 2.0 and are therefore not available when using GLSL ES 1.00 in OpenGL ES 3.x.

- Support of line continuation and support of UTF-8 characters within comments is optional in GLSL ES 1.00 when used with the OpenGL ES 2.0 API. However, support is mandated for both of these when a GLSL ES 1.00 shader is used with the OpenGL ES 3.x APIs.

2 Overview of OpenGL ES Shading

The OpenGL ES Shading Language is actually several closely related languages. These languages are used to create shaders for each of the programmable processors contained in the OpenGL ES processing pipeline. Currently, these processors are the vertex, tessellation control, tessellation evaluation, geometry and fragment and compute processors.

Compilation units for these processors are referred to as *shaders* and the processors themselves are also referred to as *shader stages*. Only one shader can be run on a processor at any one time; there is no support for linking multiple compilation units together for a single shader stage. One or more shaders are linked together to form a single program, and each program contains shader *executables* for one or more consecutive shader stages.

Unless otherwise noted in this paper, a language feature applies to all languages, and common usage will refer to these languages as a single language. The specific languages will be referred to by the name of the processor they target: vertex, tessellation control, tessellation evaluation, geometry or fragment or compute.

Most OpenGL ES state is not tracked or made available to shaders. Typically, user-defined variables will be used for communicating between different stages of the OpenGL ES pipeline. However, a small amount of state is still tracked and automatically made available to shaders, and there are a few built-in variables for interfaces between different stages of the OpenGL ES pipeline.

2.1 Vertex Processor

The *vertex processor* is a programmable unit that operates on incoming vertices and their associated data. Compilation units written in the OpenGL ES Shading Language to run on this processor are called *vertex shaders*.

The vertex processor operates on one vertex at a time. It does not replace graphics operations that require knowledge of several vertices at a time.

2.2 Tessellation Control Processor

The *tessellation control processor* is a programmable unit that operates on a patch of incoming vertices and their associated data, emitting a new output patch. Compilation units written in the OpenGL ES Shading Language to run on this processor are called tessellation control shaders.

The tessellation control processor is invoked for each vertex of the output patch. Each invocation can read the attributes of any vertex in the input or output patches, but can only write per-vertex attributes for the corresponding output patch vertex. The shader invocations collectively produce a set of per-patch attributes for the output patch.

After all tessellation control shader invocations have completed, the output vertices and per-patch attributes are assembled to form a patch to be used by subsequent pipeline stages.

Tessellation control shader invocations run mostly independently, with undefined relative execution order. However, the built-in function **barrier()** can be used to control execution order by synchronizing invocations, effectively dividing tessellation control shader execution into a set of phases. Tessellation control shaders will get undefined results if one invocation reads a per-vertex or per-patch attribute written by another invocation at any point during the same phase, or if two invocations attempt to write different values to the same per-patch output in a single phase.

2.3 Tessellation Evaluation Processor

The *tessellation evaluation processor* is a programmable unit that evaluates the position and other attributes of a vertex generated by the tessellation primitive generator, using a patch of incoming vertices and their associated data. Compilation units written in the OpenGL ES Shading Language to run on this processor are called tessellation evaluation shaders.

Each invocation of the tessellation evaluation executable computes the position and attributes of a single vertex generated by the tessellation primitive generator. The executable can read the attributes of any vertex in the input patch, plus the tessellation coordinate, which is the relative location of the vertex in the primitive being tessellated. The executable writes the position and other attributes of the vertex.

2.4 Geometry Processor

The geometry processor is a programmable unit that operates on data for incoming vertices for a primitive assembled after vertex processing and outputs a sequence of vertices forming output primitives. Compilation units written in the OpenGL ES Shading Language to run on this processor are called geometry shaders.

A single invocation of the geometry shader executable on the geometry processor will operate on a declared input primitive with a fixed number of vertices. This single invocation can emit a variable number of vertices that are assembled into primitives of a declared output primitive type and passed to subsequent pipeline stages.

2.5 Fragment Processor

The *fragment processor* is a programmable unit that operates on fragment values and their associated data. Compilation units written in the OpenGL ES Shading Language to run on this processor are called *fragment shaders*.

A fragment shader cannot change a fragment's (x, y) position. Access to neighboring fragments is not allowed. The values computed by the fragment shader are ultimately used to update framebuffer memory or texture memory, depending on the current OpenGL ES state and the OpenGL ES command that caused the fragments to be generated.

2.6 Compute Processor

The *compute processor* is a programmable unit that operates independently from the other shader processors. Compilation units written in the OpenGL ES Shading Language to run on this processor are called *compute shaders*.

2 Overview of OpenGL ES Shading

A compute shader has access to many of the same resources as fragment and other shader processors, such as textures, buffers, image variables, atomic counters, and so on. It does not have any predefined inputs nor any fixed-function outputs. It is not part of the graphics pipeline and its visible side effects are through actions on images, storage buffers, and atomic counters.

A compute shader operates on a group of work items called a work group.

A work group is a collection of shader invocations that execute the same code, potentially in parallel. An invocation within a work group may share data with other members of the same work group through shared variables and issue memory and control flow barriers to synchronize with other members of the same work group.

3 Basics

3.1 Character Set

The source character set used for the OpenGL ES shading languages is Unicode in the UTF-8 encoding scheme. Invalid UTF-8 characters are ignored. After preprocessing, only the following characters are allowed in the resulting stream of GLSL tokens:

The letters **a-z**, **A-Z**, and the underscore (`_`).

The numbers **0-9**.

The symbols period (`.`), plus (`+`), dash (`-`), slash (`/`), asterisk (`*`), percent (`%`), angled brackets (`<` and `>`), square brackets (`[` and `]`), parentheses (`(` and `)`), braces (`{` and `}`), caret (`^`), vertical bar (`|`), ampersand (`&`), tilde (`~`), equals (`=`), exclamation point (`!`), colon (`:`), semicolon (`;`), comma (`,`), and question mark (`?`).

The number sign (`#`) for preprocessor use.

Backslash (`\`), used to indicate line continuation when immediately preceding a new-line.

White space: the space character, horizontal tab, vertical tab, form feed, carriage-return, and line-feed.

There are no digraphs or trigraphs. There are no escape sequences or other uses of the backslash beyond use as the line-continuation character.

Lines are relevant for compiler diagnostic messages and the preprocessor. They are terminated by carriage-return or line-feed. If both are used together, it will count as only a single line termination. For the remainder of this document, any of these combinations is simply referred to as a new-line. Lines may be of arbitrary length.

In general, the language's use of this character set is case sensitive.

There are no character or string data types, so no quoting characters are included.

There is no end-of-file character.

3.2 Source Strings

The source for a single shader is an array of strings of characters from the character set. A single shader is made from the concatenation of these strings. Each string can contain multiple lines, separated by new-lines. No new-lines need be present in a string; a single line can be formed from multiple strings. No new-lines or other characters are inserted by the implementation when it concatenates the strings to form a single shader.

Diagnostic messages returned from compiling a shader must identify both the line number within a string and which source string the message applies to. Source strings are counted sequentially with the first string being string 0. Line numbers are one more than the number of new-lines that have been processed, including counting the new lines that will be removed by the line-continuation character (`\`).

Lines separated by the line-continuation character preceding a new line are concatenated together before either comment processing or preprocessing. This means that no white space is substituted for the line-continuation character. That is, a single token could be formed by the concatenation by taking the characters at the end of one line concatenating them with the characters at the beginning of the next line.

```
float f\
oo;
// forms a single line equivalent to "float foo;"
// (assuming '\\' is the last character before the new line and "oo" are
// the first two characters of the next line)
```

3.3 Version Declaration

Shaders must declare the version of the language they are written to. The version is specified in the first line of a shader by a character string:

```
#version number es
```

where *number* must be a version of the language, following the same convention as `__VERSION__` above. The directive “**#version 320 es**” is required in any shader that uses version 3.20 of the language. Any *number* representing a version of the language a compiler does not support will cause an error to be generated. Version 1.00 of the language does not require shaders to include this directive, and shaders that do not include a **#version** directive will be treated as targeting version 1.00.

Shaders declaring version 3.20 of the shading language cannot be linked with shaders declaring a previous version.

The **#version** directive must be present in the first line of a shader and must be followed by a newline. It may contain optional white-space as specified below but no other characters are allowed. The directive is only permitted in the first line of a shader.

Processing of the **#version** directive occurs before all other preprocessing, including line concatenation and comment processing.

version-declaration:

```
whitespaceopt POUND whitespaceopt VERSION whitespace number whitespace ES whitespaceopt
```

Tokens:

POUND	#
VERSION	version
ES	es

3.4 Preprocessor

There is a preprocessor that processes the source strings as part of the compilation process.

The complete list of preprocessor directives is as follows.

```
#
#define
#undef

#if
#ifdef
#ifndef
#else
#elif
#endif

#error
#pragma

#extension

#line
```

The following operator is also available:

```
defined
```

Note that the version directive is not considered to be a preprocessor directive and so is not listed here.

Each number sign (#) can be preceded in its line only by spaces or horizontal tabs. It may also be followed by spaces and horizontal tabs, preceding the directive. Each directive is terminated by a new-line. Preprocessing does not change the number or relative location of new-lines in a source string.

The number sign (#) on a line by itself is ignored. Any directive not listed above will cause an error.

#define and **#undef** functionality are defined as is standard for C++ preprocessors for macro definitions both with and without macro parameters.

The following predefined macros are available:

```
__LINE__
__FILE__
__VERSION__
GL_ES
```

__LINE__ will substitute a decimal integer constant that is one more than the number of preceding new-lines in the current source string.

__FILE__ will substitute a decimal integer constant that says which source string number is currently being processed.

`__VERSION__` will substitute a decimal integer reflecting the version number of the OpenGL ES shading language. The version of the shading language described in this document will have `__VERSION__` substitute the decimal integer 320.

`GL_ES` will be defined and set to 1. This is not true for the non-ES OpenGL Shading Language, so it can be used to do a compile time test to determine if a shader is running on an ES system.

By convention, all macro names containing two consecutive underscores (`__`) are reserved for use by underlying software layers. Defining such a name in a shader does not itself result in an error, but may result in unintended behaviors that stem from having multiple definitions of the same name. All macro names prefixed with “GL_” (“GL” followed by a single underscore) are also reserved, and defining such a name results in a compile-time error.

It is an error to undefine or to redefine a built-in (pre-defined) macro name.

The maximum length of a macro name is 1024 characters. It is an error to declare a name with a length greater than this.

`#if`, `#ifdef`, `#ifndef`, `#else`, `#elif`, and `#endif` are defined to operate as for C++ except for the following:

- Expressions following `#if` and `#elif` are restricted to expressions operating on literal integer constants, plus identifiers consumed by the **defined** operator.
- Undefined identifiers not consumed by the **defined** operator do not default to '0'. Use of such identifiers causes an error.
- Character constants are not supported.

As in C++, a macro name defined with an empty replacement list does not default to '0' when used in a preprocessor expression.

The operators available are as follows:

Precedence	Operator class	Operators	Associativity
1 (highest)	parenthetical grouping	()	NA
2	unary	defined + - ~ !	Right to Left
3	multiplicative	* / %	Left to Right
4	additive	+ -	Left to Right
5	bit-wise shift	<< >>	Left to Right
6	relational	< > <= >=	Left to Right
7	equality	== !=	Left to Right
8	bit-wise and	&	Left to Right
9	bit-wise exclusive or	^	Left to Right
10	bit-wise inclusive or		Left to Right
11	logical and	&&	Left to Right
12 (lowest)	logical inclusive or		Left to Right

The **defined** operator can be used in either of the following ways:

```
defined identifier
defined ( identifier )
```

There are no number sign based operators (e.g. no # or #@), no ## operator, nor is there a **sizeof** operator.

The semantics of applying operators in the preprocessor match those standard in the C++ preprocessor with the following exceptions:

- The 2nd operand in a logical and ('&&') operation is evaluated if and only if the 1st operand evaluates to non-zero.
- The 2nd operand in a logical or ('||') operation is evaluated if and only if the 1st operand evaluates to zero.

If an operand is not evaluated, the presence of undefined identifiers in the operand will not cause an error.

Preprocessor expressions will be evaluated at compile time.

#error will cause the implementation to put a diagnostic message into the shader object's information log (see section 7.12 "Shader, Program and Program Pipeline Queries" in the OpenGL ES Graphics System Specification for how to access a shader object's information log). The message will be the tokens following the **#error** directive, up to the first new-line. The implementation must treat the presence of a **#error** directive as a compile-time error.

#pragma allows implementation dependent compiler control. Tokens following **#pragma** are not subject to preprocessor macro expansion. If an implementation does not recognize the tokens following **#pragma**, then it will ignore that pragma. The following pragmas are defined as part of the language.

```
#pragma STDGL
```

The **STDGL** pragma is used to reserve pragmas for use by this and future revisions of the language. No implementation may use a pragma whose first token is **STDGL**.

```
#pragma optimize(on)
#pragma optimize(off)
```

can be used to turn off optimizations as an aid in developing and debugging shaders. It can only be used outside function definitions. By default, optimization is turned on for all shaders. The debug pragma

```
#pragma debug(on)
#pragma debug(off)
```

can be used to enable compiling and annotating a shader with debug information, so that it can be used with a debugger. It can only be used outside function definitions. By default, debug is turned off.

The scope as well as the effect of the optimize and debug pragmas is implementation-dependent except that their use must not generate an error.

By default, compilers of this language must issue compile time syntactic, grammatical, and semantic errors for shaders that do not conform to this specification. Any extended behavior must first be enabled. Directives to control the behavior of the compiler with respect to extensions are declared with the **#extension** directive

```
#extension extension_name : behavior
#extension all : behavior
```

where *extension_name* is the name of an extension. Extension names are not documented in this specification. The token **all** means the behavior applies to all extensions supported by the compiler. The *behavior* can be one of the following:

<i>behavior</i>	Effect
require	Behave as specified by the extension <i>extension_name</i> . Give an error on the #extension if the extension <i>extension_name</i> is not supported, or if all is specified.
enable	Behave as specified by the extension <i>extension_name</i> . Warn on the #extension if the extension <i>extension_name</i> is not supported. Give an error on the #extension if all is specified.
warn	Behave as specified by the extension <i>extension_name</i> , except issue warnings on any detectable use of that extension, unless such use is supported by other enabled or required extensions. If all is specified, then warn on all detectable uses of any extension used. Warn on the #extension if the extension <i>extension_name</i> is not supported.
disable	Behave (including issuing errors and warnings) as if the extension <i>extension_name</i> is not part of the language definition. If all is specified, then behavior must revert back to that of the non-extended core version of the language being compiled to. Warn on the #extension if the extension <i>extension_name</i> is not supported.

The **extension** directive is a simple, low-level mechanism to set the behavior for each extension. It does not define policies such as which combinations are appropriate, those must be defined elsewhere. Order of directives matters in setting the behavior for each extension: Directives that occur later override those seen earlier. The **all** variant sets the behavior for all extensions, overriding all previously issued **extension** directives, but only for the *behaviors* **warn** and **disable**.

The initial state of the compiler is as if the directive

```
#extension all : disable
```

was issued, telling the compiler that all error and warning reporting must be done according to this specification, ignoring any extensions.

Each extension can define its allowed granularity of scope. If nothing is said, the granularity is a shader (that is, a single compilation unit), and the extension directives must occur before any non-preprocessor tokens. If necessary, the linker can enforce granularities larger than a single compilation unit, in which case each involved shader will have to contain the necessary extension directive.

Macro expansion is not done on lines containing **#extension** and **#version** directives.

For each extension there is an associated macro. The macro is always defined in an implementation that supports the extension. This allows the following construct to be used:

```
#ifdef OES_extension_name
    #extension OES_extension_name : enable
    // code that requires the extension
#else
    // alternative code
#endif
```

#line must have, after macro substitution, one of the following forms:

```
#line line
#line line source-string-number
```

where *line* and *source-string-number* are constant integral expressions. After processing this directive (including its new-line), the implementation will behave as if it is compiling at line number *line* and source string number *source-string-number*. Subsequent source strings will be numbered sequentially, until another **#line** directive overrides that numbering.

If during macro expansion a preprocessor directive is encountered, the results are undefined; the compiler may or may not report an error in such cases.

3.5 Comments

Comments are delimited by `/*` and `*/`, or by `//` and a newline. `/*` style comments include the initial `/*` marker and continue up to, but not including, the terminating newline. `/*...*/` comments include both the start and end marker. The begin comment delimiters (`/*` or `//`) are not recognized as comment delimiters inside of a comment, hence comments cannot be nested. Comments are treated syntactically as a single space.

3.6 Tokens

The language is a sequence of tokens. A token can be

token:
keyword
identifier
integer-constant
floating-constant
operator
 ; { }

3.7 Keywords

The following are the keywords in the language, and cannot be used for any other purpose than that defined by this document:

const uniform buffer shared
coherent volatile restrict readonly writeonly
atomic_uint
layout
centroid flat smooth
patch sample
precise
break continue do for while switch case default
if else
in out inout
float int void bool true false
invariant
discard return
mat2 mat3 mat4
mat2x2 mat2x3 mat2x4
mat3x2 mat3x3 mat3x4
mat4x2 mat4x3 mat4x4
vec2 vec3 vec4 ivec2 ivec3 ivec4 bvec2 bvec3 bvec4
uint uvec2 uvec3 uvec4

lowp mediump highp precision
 sampler2D sampler3D samplerCube
 sampler2DShadow samplerCubeShadow
 sampler2DArray
 sampler2DArrayShadow
 isampler2D isampler3D isamplerCube
 isampler2DArray
 usampler2D usampler3D usamplerCube
 usampler2DArray
 sampler2DMS isampler2DMS usampler2DMS
 samplerBuffer isamplerBuffer usamplerBuffer
 imageBuffer iimageBuffer uimageBuffer
 imageCubeArray iimageCubeArray uimageCubeArray
 samplerCubeArray isamplerCubeArray usamplerCubeArray
 samplerCubeArrayShadow
 sampler2DMSArray isampler2DMSArray usampler2DMSArray
 image2DArray iimage2DArray uimage2DArray
 image2D iimage2D uimage2D
 image3D iimage3D uimage3D
 imageCube iimageCube uimageCube
 struct

The following are the keywords reserved for future use. Using them will result in an error:

attribute varying
 resource
 noperspective
 subroutine
 common partition active
 asm

```

class union enum typedef template this
goto
inline noinline public static extern external interface
long short double half fixed unsigned superp
input output
hvec2 hvec3 hvec4 dvec2 dvec3 dvec4 fvec2 fvec3 fvec4
sampler3DRect
filter
image1D
iimage1D
uimage1D
image1DArray
iimage1DArray uimage1DArray
sampler1D sampler1DShadow sampler1DArray sampler1DArrayShadow
isampler1D isampler1DArray usampler1D usampler1DArray
sampler2DRect sampler2DRectShadow isampler2DRect usampler2DRect
sizeof cast
namespace using
dmat2 dmat3 dmat4
dmat2x2 dmat2x3 dmat2x4
dmat3x2 dmat3x3 dmat3x4
dmat4x2 dmat4x3 dmat4x4
image2DRect iimage2DRect uimage2DRect
image2DMS iimage2DMS uimage2DMS
image2DMSArray iimage2DMSArray uimage2DMSArray

```

In addition, all identifiers containing two consecutive underscores (__) are reserved for use by underlying software layers. Defining such a name in a shader does not itself result in an error, but may result in unintended behaviors that stem from having multiple definitions of the same name.

3.8 Identifiers

Identifiers are used for variable names, function names, structure names, and field selectors (field selectors select components of vectors and matrices similar to structure fields, as discussed in section 5.5 “Vector Components” and section 5.6 “Matrix Components”). Identifiers have the form

identifier:

nondigit

identifier nondigit

identifier digit

nondigit: one of

_ a b c d e f g h i j k l m n o p q r s t u v w x y z

A B C D E F G H I J K L M N O P Q R S T U V W X Y Z

digit: one of

0 1 2 3 4 5 6 7 8 9

Identifiers starting with “gl_” are reserved for use by OpenGL ES, and in general, may not be declared in a shader as either a variable or a function; this results in an error. However, as noted in the specification, there are certain cases where previously declared variables can be redeclared, and predeclared “gl_” names are allowed to be redeclared in a shader only for these specific purposes.

The maximum length of an identifier is 1024 characters. It is an error if the length exceeds this value.

3.9 Definitions

Some language rules described below depend on the following definitions.

3.9.1 Static Use

A shader contains a *static use* of a variable x if, after preprocessing, the shader contains a statement that would read or write x (or part of x), whether or not run-time flow of control will cause that statement to be executed. Such a variable is referred to as being *statically used*.

3.9.2 Uniform and Non-Uniform Control Flow

When executing statements in a fragment shader, control flow starts as *uniform control flow*; all fragments enter the same control path into *main()*. Control flow becomes *non-uniform* when different fragments take different paths through control-flow statements (selection, iteration, and jumps). Control flow subsequently returns to being uniform after such divergent sub-statements or skipped code completes, until the next time different control paths are taken.

For example:

```
main()
{
    float a = ...; // this is uniform control flow
    if (a < b) { // this expression is true for some fragments, not all
        ...;    // non-uniform control flow
    } else {
        ...;    // non-uniform control flow
    }
    ...;        // uniform control flow again
}
```

Other examples of non-uniform control flow can occur within switch statements and after conditional breaks, continues, early returns, and after fragment discards, when the condition is true for some fragments but not others. Loop iterations that only some fragments execute are also non-uniform control flow.

This is similarly defined for other shader stages, based on the per-instance data items they process.

3.9.3 Dynamically Uniform Expressions

A fragment-shader expression is *dynamically uniform* if all fragments evaluating it get the same resulting value. When loops are involved, this refers to the expression's value for the same loop iteration. When functions are involved, this refers to calls from the same call point.

This is similarly defined for other shader stages, based on the per-instance data they process.

Note that constant expressions are trivially dynamically uniform. It follows that typical loop counters based on these are also dynamically uniform.

3.10 Logical Phases of Compilation

The compilation process is based on a subset of the C++ standard (see section 15: Normative References). The compilation units for the shader processors are processed separately before optionally being linked together in the final stage of compilation. The logical phases of compilation are:

1. Source strings are input as byte sequences. The value 'zero' is interpreted as a terminator.
2. Source strings are concatenated to form a single input. Zero bytes are discarded but all other values are retained.
3. Each string is interpreted according to the UTF-8 standard, with the exception that all invalid byte sequences are retained in their original form for subsequent processing.
4. Each {carriage-return, line-feed} and {line-feed, carriage return} sequence is replaced by a single newline. All remaining carriage-return and line-feed characters are then each replaced by a newline.

5. Line numbering for each character, which is equal to the number of preceding newlines plus one, is noted. Note this can only be subsequently changed by the `#line` directive and is not affected by the removal of newlines in phase 6 of compilation.
6. Wherever a backslash (`\`) occurs immediately before a newline, both are deleted. Note that no whitespace is substituted, thereby allowing a single preprocessing token to span a newline. This operation is not recursive; any new {backslash newline} sequences generated are not removed.
7. All comments are replaced with a single space. All (non-zero) characters and invalid UTF-8 byte sequences are allowed within comments. `/*` style comments include the initial `/*` marker and continue up to, but not including, the terminating newline. `/*...*/` comments include both the start and end marker.
8. The source string is converted into a sequence of preprocessing tokens. These tokens include preprocessing numbers, identifiers and preprocessing operations. The line number associated with each token is copied from the line number of the first character of the token.
9. The preprocessor is run. Directives are executed and macro expansion is performed.
10. White space and newlines are discarded.
11. Preprocessing tokens are converted into tokens.
12. The syntax is analyzed according to the GLSL ES grammar.
13. The result is checked according to the semantic rules of the language.
14. Optionally, the shaders are linked together to form one or more programs or separable programs. When a pair of shaders from consecutive stages are linked into the same program, any outputs and corresponding inputs not used in both shaders may be discarded.
15. The binary is generated.

4 Variables and Types

All variables and functions must be declared before being used. Variable and function names are identifiers.

There are no default types. All variable and function declarations must have a declared type, and optionally qualifiers. A variable is declared by specifying its type followed by one or more names separated by commas. In many cases, a variable can be initialized as part of its declaration by using the assignment operator (=). The grammar near the end of this document provides a full reference for the syntax of declaring variables.

User-defined types may be defined using **struct** to aggregate a list of existing types into a single name.

The OpenGL ES Shading Language is type safe. There are no implicit conversions between types.

4.1 Basic Types

Definition:

A basic type is a type defined by a keyword in the language.

The OpenGL ES Shading Language supports the following basic data types, grouped as follows.

Transparent types

Type	Meaning
void	for functions that do not return a value
bool	a conditional type, taking on values of true or false
int	a signed integer
uint	an unsigned integer
float	a single floating-point scalar
vec2	a two-component floating-point vector
vec3	a three-component floating-point vector
vec4	a four-component floating-point vector
bvec2	a two-component Boolean vector
bvec3	a three-component Boolean vector
bvec4	a four-component Boolean vector
ivec2	a two-component signed integer vector
ivec3	a three-component signed integer vector
ivec4	a four-component signed integer vector

4 Variables and Types

Type	Meaning
uvec2	a two-component unsigned integer vector
uvec3	a three-component unsigned integer vector
uvec4	a four-component unsigned integer vector
mat2	a 2×2 floating-point matrix
mat3	a 3×3 floating-point matrix
mat4	a 4×4 floating-point matrix
mat2x2	same as a mat2
mat2x3	a floating-point matrix with 2 columns and 3 rows
mat2x4	a floating-point matrix with 2 columns and 4 rows
mat3x2	a floating-point matrix with 3 columns and 2 rows
mat3x3	same as a mat3
mat3x4	a floating-point matrix with 3 columns and 4 rows
mat4x2	a floating-point matrix with 4 columns and 2 rows
mat4x3	a floating-point matrix with 4 columns and 3 rows
mat4x4	same as a mat4

Floating Point Opaque Types

Type	Meaning
sampler2D image2D	a handle for accessing a 2D texture/image
sampler3D image3D	a handle for accessing a 3D texture/image
samplerCube imageCube	a handle for accessing a cube mapped texture/image
samplerCubeShadow	a handle for accessing a cube map depth texture with comparison
sampler2DShadow	a handle for accessing a 2D depth texture with comparison
sampler2DArray image2DArray	a handle for accessing a 2D array texture/image
sampler2DArrayShadow	a handle for accessing a 2D array depth texture with comparison
sampler2DMS	a handle for accessing a 2D multisample texture
samplerBuffer imageBuffer	a handle for accessing a buffer texture/image

4 Variables and Types

Type	Meaning
samplerCubeArray imageCubeArray	a handle for accessing a cube map array texture/image
samplerCubeArrayShadow	a handle for accessing a cube map array depth texture with comparison
sampler2DMSArray	handle for accessing a 2D multisample array texture

Signed Integer Opaque Types

Type	Meaning
isampler2D iimage2D	a handle for accessing an integer 2D texture/image
isampler3D iimage3D	a handle for accessing an integer 3D texture/image
isamplerCube iimageCube	a handle for accessing an integer cube mapped texture/image
isampler2DArray iimage2DArray	a handle for accessing an integer 2D array texture/image
isampler2DMS	a handle for accessing an integer 2D multisample texture
isamplerBuffer iimageBuffer	a handle for accessing an integer buffer texture/image
isamplerCubeArray iimageCubeArray	a handle for accessing an integer cube map array texture/image
isampler2DMSArray	handle for accessing an integer 2D multisample array texture

Unsigned Integer Opaque Types

Type	Meaning
usampler2D uimage2D	a handle for accessing an unsigned integer 2D texture/image
usampler3D uimage3D	a handle for accessing an unsigned integer 3D texture/image
usamplerCube uimageCube	a handle for accessing an unsigned integer cube mapped texture/image
usampler2DArray uimage2DArray	a handle for accessing an unsigned integer 2D array texture/image
atomic_uint	a handle for accessing an unsigned atomic counter
usampler2DMS	a handle for accessing an unsigned integer 2D multisample texture
usamplerBuffer uimageBuffer	a handle for accessing an unsigned integer buffer texture/image
usamplerCubeArray uimageCubeArray	a handle for accessing an unsigned integer cube map array texture/image
usampler2DMSArray	handle for accessing an unsigned integer 2D multisample array texture/image

In addition, a shader can aggregate these using arrays and structures to build more complex types.

There are no pointer types.

4.1.1 Void

Functions that do not return a value must be declared as **void**. There is no default function return type. The keyword **void** cannot be used in any other declarations (except for empty formal or actual parameter lists).

4.1.2 Booleans

Definition:

A boolean type is any boolean scalar or vector type.

To make conditional execution of code easier to express, the type **bool** is supported. There is no expectation that hardware directly supports variables of this type. It is a genuine Boolean type, holding only one of two values meaning either true or false. Two keywords **true** and **false** can be used as literal Boolean constants. Booleans are declared and optionally initialized as in the follow example:

```
bool success;          // declare "success" to be a Boolean
bool done = false;     // declare and initialize "done"
```

The right side of the assignment operator (=) must be an expression whose type is **bool**.

Expressions used for conditional jumps (**if**, **for**, **?:**, **while**, **do-while**) must evaluate to the type **bool**.

4.1.3 Integers

Definitions:

An *integral type* is any signed or unsigned, scalar or vector integer type. It excludes arrays and structures.

A *scalar integral type* is a scalar signed or unsigned integer type:

A *vector integral types* is a vector of signed or unsigned integers:

Signed and unsigned integer variables are fully supported. In this document, the term *integer* is meant to generally include both signed and unsigned integers. Highp unsigned integers have exactly 32 bits of precision. Highp signed integers use 32 bits, including a sign bit, in two's complement form. Mediump and lowp integers have implementation-defined numbers of bits. Operations resulting in overflow or underflow will not cause any exception, nor will they saturate, rather they will “wrap” to yield the low-order n bits of the result where n is the size in bits of the integer.

See section 4.7.1 “Range and Precision” for details.

Integers are declared and optionally initialized with integer expressions, as in the following example:

```
int i, j = 42;  // default integer literal type is int
uint k = 3u;   // “u” establishes the type as uint
```

Literal integer constants can be expressed in decimal (base 10), octal (base 8), or hexadecimal (base 16) as follows.

integer-constant:

decimal-constant integer-suffix_{opt}

octal-constant integer-suffix_{opt}

hexadecimal-constant integer-suffix_{opt}

integer-suffix: one of

u U

decimal-constant:

nonzero-digit

decimal-constant digit

octal-constant:

0

octal-constant octal-digit

hexadecimal-constant:

0x *hexadecimal-digit*

0X *hexadecimal-digit*

hexadecimal-constant hexadecimal-digit

digit:

0

nonzero-digit

nonzero-digit: one of

1 2 3 4 5 6 7 8 9

octal-digit : one of

0 1 2 3 4 5 6 7

hexadecimal-digit: one of

0 1 2 3 4 5 6 7 8 9

a b c d e f

A B C D E F

No white space is allowed between the digits of an integer constant, including after the leading **0** or after the leading **0x** or **0X** of a constant, or before the suffix **u** or **U**. When the suffix **u** or **U** is present, the literal has type **uint**, otherwise the type is **int**. A leading unary minus sign (-) is interpreted as an arithmetic unary negation, not as part of the constant.

It is an error to provide a literal integer whose value would be too large to store in a highp uint variable. Note that this only applies to literals; no error checking is performed on the result of a constant expression.

Examples

```
1           // signed integer, value 1

1u          // unsigned integer, value 1

-1          // unary minus applied to signed integer.
            // result is a signed integer, value -1

-1u         // unary minus applies to unsigned integer
            // result is an unsigned integer, value 0xffffffff

0xffffffff // signed integer, value -1

0xffffffffu // unsigned integer, value 0xffffffff

0xffffffff0 // error: values of signed integer is too large
```

4.1.4 Floats

Definition:

4 Variables and Types

A *floating point type* is any floating point scalar, vector or matrix type. It excludes arrays and structures.

Floats are available for use in a variety of scalar calculations. Floating-point variables are defined as in the following example:

```
float a, b = 1.5;
```

As an input value to one of the processing units, a floating-point variable is expected to match the IEEE 754 single precision floating-point definition for precision and dynamic range. Highp floating-point variables within a shader are encoded according to the IEEE 754 specification for single-precision floating-point values (logically, not necessarily physically). While encodings are logically IEEE 754, operations (addition, multiplication, etc.) are not necessarily performed as required by IEEE 754. See section 4.7.1 “Range and Precision” for more details on precision and usage of NaNs (Not a Number) and Infs (positive or negative infinities).

Floating-point constants are defined as follows.

floating-constant:

fractional-constant *exponent-part*_{opt} *floating-suffix*_{opt}
digit-sequence *exponent-part* *floating-suffix*_{opt}

fractional-constant:

digit-sequence . *digit-sequence*
digit-sequence .
. *digit-sequence*

exponent-part:

e *sign*_{opt} *digit-sequence*
E *sign*_{opt} *digit-sequence*

sign: one of

+ -

digit-sequence:

digit
digit-sequence digit

floating-suffix: one of

f F

A decimal point (.) is not needed if the exponent part is present. No white space may appear anywhere within a floating-point constant, including before a suffix. A leading unary minus sign (-) is interpreted as a unary operator and is not part of the floating-point constant.

There is no limit on the number of digits in any *digit-sequence*. If the value of the floating point number is too large (small) to be stored as a single precision value, it is converted to positive (negative) infinity. A value with a magnitude too small to be represented as a mantissa and exponent is converted to zero. Implementations may also convert subnormal (denormalized) numbers to zero.

4.1.5 Vectors

The OpenGL ES Shading Language includes data types for generic 2-, 3-, and 4-component vectors of floating-point values, integers, and Booleans. Floating-point vector variables can be used to store colors, normals, positions, texture coordinates, texture lookup results and the like. Boolean vectors can be used for component-wise comparisons of numeric vectors. Some examples of vector declarations are:

```
vec2 texcoord1, texcoord2;
vec3 position;
vec4 myRGBA;
ivec2 textureLookup;
bvec3 less;
```

Initialization of vectors can be done with constructors. See section 5.4.2 (“Vector and Matrix Constructors”).

4.1.6 Matrices

The OpenGL ES Shading Language has built-in types for 2×2, 2×3, 2×4, 3×2, 3×3, 3×4, 4×2, 4×3, and 4×4 matrices of floating-point numbers. The first number in the type is the number of columns, the second is the number of rows. Example matrix declarations:

```
mat2 mat2D;
mat3 optMatrix;
mat4 view, projection;
mat4x4 view; // an alternate way of declaring a mat4
mat3x2 m;    // a matrix with 3 columns and 2 rows
```

Initialization of matrix values is done with constructors (described in section 5.4.2 “Vector and Matrix Constructors”) in column-major order.

mat2 is an alias for **mat2x2**, not a distinct type. Similarly for **mat3** and **mat4**. The following is legal:

```
mat2 a;
mat2x2 b = a;
```

4.1.7 Opaque Types

Definition:

An opaque type is a type where the internal structure of the type is hidden from the language.

The opaque types, as listed in the following sections, declare variables that are effectively opaque handles to other objects. These objects are accessed through built-in functions, not through direct reading or writing of the declared variable. They can only be declared as function parameters or in **uniform**-qualified variables (see section 4.3.5 “Uniform Variables”). The only opaque types that take memory qualifiers are the image types. Except for array indexing, structure member selection, and parentheses, opaque variables are not allowed to be operands in expressions; such use results in a compile-time error.

When aggregated into arrays within a shader, opaque types can only be indexed with a dynamically uniform integral expression (see section 3.9.3) unless otherwise noted; otherwise, results are undefined.

Opaque variables cannot be treated as l-values; hence cannot be used as **out** or **inout** function parameters, nor can they be assigned into. Any such use results in a compile-time error. However, they can be passed as in parameters with matching types and memory qualifiers. They cannot be declared with an initializer.

Because a single opaque type declaration effectively declares two objects, the opaque handle itself and the object it is a handle to, there is room for both a storage qualifier and a memory qualifier. The storage qualifier will qualify the opaque handle, while the memory qualifier will qualify the object it is a handle to.

4.1.7.1 Samplers

Sampler types (e.g., **sampler2D**) are opaque types, declared and behaving as described above for opaque types.

Sampler variables are handles to two- and three- dimensional textures, cube maps, depth textures (shadowing), etc., as enumerated in the basic types tables. There are distinct sampler types for each texture target, and for each of float, integer, and unsigned integer data types. Texture accesses are done through built-in texture functions (described in section 8.9 “Texture Functions”) and samplers are used to specify which texture to access and how it is to be filtered.

4.1.7.2 Images

Image types are opaque types, declared and behaving as described above for opaque types. They can be further qualified with memory qualifiers. When aggregated into arrays within a shader, images can only be indexed with a constant integral expression.

Image variables are handles to two- or three-dimensional images corresponding to all or a portion of a single level of a texture image bound to an image unit. There are distinct image variable types for each texture target, and for each of float, integer, and unsigned integer data types. Image accesses should use an image type that matches the target of the texture whose level is bound to the image unit, or for non-layered bindings of 3D or array images should use the image type that matches the dimensionality of the layer of the image (i.e. a layer of 3D, 2DArray, Cube or CubeArray should use image2D). If the image target type does not match the bound image in this manner, if the data type does not match the bound image, or if the format layout qualifier does not match the image unit format as described in Section 8.22 “Texture Image Loads and Stores” of the OpenGL ES Specification, the results of image accesses are undefined but cannot include program termination.

Image variables are used in the image load, store and atomic functions described in section 8.12 (“Image Functions”) to specify an image to access.

4.1.7.3 Atomic Counters

Atomic Counter types (e.g. **atomic_uint**) are opaque handles to counters, declared and behaving as described above for opaque types. The variables they declare specify which counter to access when using the built-in atomic counter functions as described in section 8.10 (“Atomic-Counter Functions”). They are bound to buffers as described in section 4.4.6 (“Atomic Counter Layout Qualifiers”).

Atomic counters aggregated into arrays within a shader can only be indexed with dynamically uniform integral expressions, otherwise results are undefined.

The default precision of all atomic types is `highp`. It is an error to declare an atomic type with a different precision or to specify the default precision for an atomic type to be `lowp` or `mediump`.

4.1.8 Structures

User-defined types can be created by aggregating other types into a structure using the **struct** keyword. For example,

```
struct light {
    float intensity;
    vec3 position;
} lightVar;
```

In this example, *light* becomes the name of the new type, and *lightVar* becomes a variable of type *light*. To declare variables of the new type, use its name (without the keyword **struct**).

```
light lightVar2;
```

More formally, structures are declared as follows. However, the definitive grammar is as given in section 9 (“Shading Language Grammar”).

struct-definition:

*qualifier*_{opt} **struct** *name*_{opt} { *member-list* } *declarators*_{opt} ;

member-list:

member-declaration;

member-declaration member-list;

member-declaration:

basic-type declarators;

where *name* becomes the user-defined type, and can be used to declare variables to be of this new type. The *name* shares the same name space as other variables, types, and functions. All previously visible variables, types, constructors, or functions with that name are hidden. The optional *qualifier* only applies to any *declarators*, and is not part of the type being defined for *name*.

Structures must have at least one member declaration. Member declarators may contain precision qualifiers, but may not contain any other qualifiers. Bit fields are not supported. Member types must be already defined (there are no forward references). Member declarations cannot contain initializers. Member declarators can contain arrays. Such arrays must have a size specified, and the size must be a constant integral expression that's greater than zero (see section 4.3.3 “Constant Expressions”). Each level of structure has its own name space for names given in member declarators; such names need only be unique within that name space.

Anonymous structures are not supported. Embedded structure definitions are not supported.

```
struct S { float f; }; // Allowed: S is defined as a structure.

struct T {
    S;                // Error: anonymous structures disallowed
    struct { ... };   // Error: embedded structures disallowed
    S s;              // Allowed: nested structure with a name.
};
```

Structures can be initialized at declaration time using constructors, as discussed in section 5.4.3 (“Structure Constructors”).

Any restrictions on the usage of a type or qualifier also apply to a structure that contains that type or qualifier. This applies recursively.

Structures can contain variables of any type except:

- `atomic_uint` (since there is no mechanism to specify the location)

4.1.9 Arrays

Variables of the same type can be aggregated into arrays by declaring a name followed by brackets (`[]`) enclosing an optional size. When present, the array size must be a constant integral expression (see section 4.3.3 “Constant Expressions”) greater than zero. The type of the size parameter can be a signed or unsigned integer and the choice of type does not affect the type of the resulting array. Arrays only have a single dimension (a single number within “`[]`”), however, arrays of arrays can be declared. Any type can be formed into an array.

Arrays are sized either at compile-time or at run-time. To size an array at compile-time, either the size must be specified within the brackets as above or must be inferred from the type of the initializer.

If an array is declared as the last member of a shader storage block and the size is not specified at compile-time, it is sized at run-time. In all other cases, arrays are sized only at compile-time. An array declaration sized at compile-time which leaves the size of the array unspecified is an error.

For compile-time sized arrays, it is illegal to index an array with a constant integral expression greater than or equal to the declared size or with a negative constant expression. Arrays declared as formal parameters in a function declaration must also specify a size. Undefined behavior results from indexing an array with a non-constant expression that’s greater than or equal to the array’s size or less than 0. If robust buffer access is enabled (see section 10.3.5 “Robust Buffer Access” of the OpenGL ES 3.2 API Specification), such indexing must not result in abnormal program termination. The results are still undefined, but implementations are encouraged to produce zero values for such accesses.

Some examples are:

4 Variables and Types

```
float frequencies[3];
uniform vec4 lightPosition[4u];
const int numLights = 2;
light lights[numLights];

vec4 a[3][2];
a.length()      // this is 3
a[x].length()   // this is 2

// a shader storage block, introduced in section 4.3.7 "Buffer Variables"
buffer b {
    float u[]; // an error
    vec4 v[];  // okay, v will be sized at run-time
} name[3];    // when the block is arrayed, all u will be the same size,
              // but not necessarily all v, if sized dynamically
```

When the **length** method will return a compile-time constant, the expression in brackets (x above) will be evaluated and subject to the rules required for array indices, but the array will not be dereferenced. Thus, behavior is well defined even if the run-time value of the expression is out of bounds.

An array type can be formed by specifying a non-array type (`type_specifier_nonarray`) followed by an `array_specifier`.

```
float[5]
```

Note that the construct *type [size]* does not always result in an array of length *size* of type *type*:

```
float[2][3] // an array of size [2] of array of size [3] of float,
            // not size [3] of float[2]
```

This type can be used anywhere any other type can be used, including as the return value from a function

```
float[5] foo() { }
```

as a constructor of an array:

```
float[5](3.4, 4.2, 5.0, 5.2, 1.1)
```

as an unnamed parameter:

```
void foo(float[5])
```

and as an alternate way of declaring a variable or function parameter:

```
float[5] a;
```

An array type can also be formed without specifying a size if the definition includes an initializer:

4 Variables and Types

```
float x[] = float[2] (1.0, 2.0);    // declares an array of size 2
float y[] = float[] (1.0, 2.0, 3.0); // declares an array of size 3

float a[5];
float b[] = a;
```

Note that the initializer itself does not need to be a constant expression but the length of the initializer will be a constant expression.

Arrays can have initializers formed from array constructors:

```
float a[5] = float[5](3.4, 4.2, 5.0, 5.2, 1.1);
float a[5] = float[](3.4, 4.2, 5.0, 5.2, 1.1); // same thing
```

An array of arrays can be declared as

```
vec4 a[3][2]; // size-3 array of size-2 array of vec4
```

which declares a one-dimensional array of size 3 of one-dimensional arrays of size 2 of vec4s. The following declarations do the same thing:

```
vec4[2] a[3]; // size-3 array of size-2 array of vec4
vec4[3][2] a; // size-3 array of size-2 array of vec4
```

When in transparent memory (like in a uniform block), the layout is that the 'inner' (right-most in declaration) dimensions iterate faster than the outer dimensions. That is, for the above, the order in memory would be:

low address...a[0][0] a[0][1] a[1][0] a[1][1] a[2][0] a[2][1]...high address

The last member of a shader storage block (section 4.3.7 “Buffer Variables”), may be declared without specifying a size. For such arrays, the effective array size is inferred at run-time from the size of the data store backing the shader storage block. Such *runtime-sized* arrays may be indexed with general integer expressions, but may not be passed as an argument to a function or indexed with a negative constant expression.

```
struct S { float f; };

buffer ShaderStorageBlock1
{
    vec4 a[];           // illegal
    vec4 b[];           // legal, runtime-sized arrays are last member
};

buffer ShaderStorageBlock2
{
    vec4 a[4];          // legal, size declared
    S b[];              // legal, runtime-sized arrays are allowed,
                        // including arrays of structures
};
```

However, it is a compile-time error to assign to a runtime-sized array. Assignments to individual elements of the array is allowed.

Arrays have a fixed number of elements. This can be obtained by using the length method:

```
float a[5];  
a.length(); // returns 5
```

The return value is a signed integral expression. For compile-time sized arrays, the value returned by the length method is a constant expression. For run-time sized arrays, the value returned will not be constant expression and will be determined at run time based on the size of the buffer object providing storage for the block.

The precision is determined using the same rules as for other cases where there is no intrinsic precision. See section 4.7.3 (“Precision Qualifiers”).

Any restrictions on the usage of a type also apply to arrays of that type. This applies recursively.

4.2 Scoping

The scope of a declaration determines where the declaration is visible. GLSL ES uses a system of statically nested scopes. This allows names to be redefined within a shader.

4.2.1 Definition of Terms

The term *scope* refers to a specified region of the program where names are guaranteed to be visible. For example, a *compound_statement_with_scope* (`{ 'statement statement ... ' }`) defines a scope.

A *nested scope* is a scope defined within an outer scope.

The terms '*same scope*' and '*current scope*' are equivalent to the term '*scope*' but used to emphasize that nested scopes are excluded.

The *scope of a declaration* is the region or regions of the program where that declaration is visible.

A *name space* defines where names may be defined. Within a single *name space*, a name has at most one entry, specifying it to be one of: structure, variable, or function.

In general, each scope has an associated name space. However, in certain cases e.g. for uniforms, multiple scopes share the same name space. In these cases, conflicting declarations are an error, even though the name is only visible in the scopes where it is declared.

4.2.2 Types of Scope

The scope of a variable is determined by where it is declared. If it is declared outside all function definitions, it has global scope, which starts from where it is declared and persists to the end of the shader it is declared in. If it is declared in a **while** test or a **for** statement, then it is scoped to the end of the following sub-statement (specified as *statement-no-new-scope* in the grammar). Otherwise, if it is declared as a statement within a compound statement, it is scoped to the end of that compound statement. If it is declared as a parameter in a function definition, it is scoped until the end of that function definition. A function's parameter declarations and body together form a single scope.

4 Variables and Types

```
int f( /* nested scope begins here */ int k)
{
    int k = k + 3;  // redeclaration error of the name k
    ...
}

int f(int k)
{
    {
        int k = k + 3; // 2nd k is parameter, initializing nested first k
        int m = k      // use of new k, which is hiding the parameter
    }
}
```

For both for and while loops, the sub-statement itself does not introduce a new scope for variable names, so the following has a redeclaration compile-time error:

```
for ( /* nested scope begins here */ int i = 0; i < 10; i++)
{
    int i; // redeclaration error
}
```

The body of a do-while loop introduces a new scope lasting only between the do and while (not including the while test expression), whether or not the body is simple or compound:

```
int i = 17;
do
    int i = 4; // okay, in nested scope
while (i == 0); // i is 17, scoped outside the do-while body
```

The statement following a **switch** (...) forms a nested scope.

Representing the if construct as:

if if-expression **then** if-statement **else** else-statement,

a variable declared in the if-statement is scoped to the end of the if-statement. A variable declared in the else-statement is scoped to the end of the else-statement. This applies both when these statements are simple statements and when they are compound statements. The if-expression does not allow new variables to be declared, hence does not form a new scope.

Within a declaration, the scope of a name starts immediately after the initializer if present or immediately after the name being declared if not. Several examples:

```
int x = 1;
{
    int x = 2, /* 2nd x visible here */ y = x; // y is initialized to 2
    int z = z;          // error if z not previously defined.
}
{
    int x = x;          // x is initialized to '1'
}
```

A structure name declaration is visible at the end of the *struct_specifier* in which it was declared:

```
struct S
{
    int x;
};

{
    S S = S(0); // 'S' is only visible as a struct and constructor
    S;          // 'S' is now visible as a variable
}

int x = x;      // Error if x has not been previously defined.
```

4.2.3 Redeclaring Names

All variable names, structure type names, and function names in a given scope share the same name space. Function names can be redeclared in the same scope, with the same or different parameters, without error. Otherwise, within a shader, a declared name cannot be redeclared in the same scope; doing so results in a redeclaration error. If a nested scope redeclares a name used in an outer scope, it hides all existing uses of that name. There is no way to access the hidden name or make it unhidden, without exiting the scope that hid it.

Names of built-in functions cannot be redeclared as functions. Therefore overloading or redefining built-in functions is an error.

4 Variables and Types

A *declaration* is considered to be a statement that adds a name or signature to the symbol table. A *definition* is a statement that fully defines that name or signature. E.g.

```
int f();                // declaration;
int f() {return 0;}      // declaration and definition
int x;                  // declaration and definition
int a[4];               // array declaration and definition
struct S {int x;};      // structure declaration and definition
```

The determination of equivalence of two declarations depends on the type of declaration. For functions, the whole function signature must be considered (see section 6.1 Function Definitions). For variables (including arrays) and structures only the names must match.

Within each scope, a name may be declared either as a variable declaration *or* as function declarations *or* as a structure.

Examples of combinations that are allowed:

1.

```
void f(int) {...}
void f(float) {...} // function overloading allowed
```

2.

```
void f(int);          // 1st declaration (allowed)
void f(int);          // repeated declaration (allowed)
void f(int) {...}     // single definition (allowed)
```

Examples of combinations that are disallowed:

1.

```
void f(int) {...}
void f(int) {...} // Error: repeated definition
```

2.

```
void f(int);
struct f {int x;}; // Error: type 'f' conflicts with function 'f'
```

3.

```
struct f {int x;};
int f; //Error: conflicts with the type 'f'
```

4.

```
int a[3];
int a[3]; // Error: repeated array definition
```

5.

```
int x;
int x; // Error: repeated variable definition
```

4.2.4 Global Scope

The built-in functions are scoped in the global scope users declare global variables in. That is, a shader's global scope, available for user-defined functions and global variables, is the same as the scope containing the built-in functions. Function declarations (prototypes) cannot occur inside of functions; they must be at global scope. Hence it is not possible to hide a name with a function.

4.2.5 Shared Globals

Shared globals are variables that can be accessed by multiple compilation units. In GLSL ES the only shared globals are uniforms. Vertex shader outputs are not considered to be shared globals since they must pass through the rasterization stage before they are used as input by the fragment shader.

Shared globals share the same name space, and must be declared with the same type and precision. They will share the same storage. Shared global arrays must have the same base type and the same explicit size. Scalars must have exactly the same precision, type name and type definition. Structures must have the same name, sequence of type names, and type definitions, and field names to be considered the same type. This rule applies recursively for nested or embedded types.

4.3 Storage Qualifiers

Variable declarations may have one storage qualifier specified in front of the type. These are summarized as

Qualifier	Meaning
< none: default >	local read/write memory, or an input parameter to a function
const	a compile-time constant
in	linkage into a shader from a previous stage, variable is copied in
out	linkage out of a shader to a subsequent stage, variable is copied out
uniform	value does not change across the primitive being processed, uniforms form the linkage between a shader, OpenGL ES, and the application
buffer	value is stored in a buffer object, and can be read or written both by shader invocations and the OpenGL ES API
shared	compute shader only; variable storage is shared across all work items in a local work group

Some input and output qualified variables can be qualified with at most one additional auxiliary storage qualifier:

Auxiliary Storage Qualifier	Meaning
centroid	centroid-based interpolation
sample	per-sample interpolation
patch	per-tessellation-patch attributes

Local variables can only use the **const** storage qualifier.

Note that function parameters can use **const**, **in**, and **out** qualifiers, but as *parameter qualifiers*. Parameter qualifiers are discussed in section 6.1.1 (“Function Calling Conventions”). Function return types and structure fields do not use storage qualifiers.

Data types for communication from one run of a shader executable to its next run (to communicate between fragments or between vertices) do not exist. This would prevent parallel execution of the same shader executable on multiple vertices or fragments.

In declarations of global variables with no storage qualifier or with a **const** qualifier, any initializer must be a constant expression. Declarations of global variables with other storage qualifiers may not contain initializers. Global variables without storage qualifiers that are not initialized in their declaration or by the application will not be initialized by OpenGL ES, but rather will enter *main()* with undefined values.

4.3.1 Default Storage Qualifier

If no qualifier is present on a global variable, then the variable has no linkage to the application or shaders running on other pipeline stages. For either global or local unqualified variables, the declaration will appear to allocate memory associated with the processor it targets. This variable will provide read/write access to this allocated memory.

4.3.2 Constant Qualifier

Named compile-time constants can be declared using the **const** qualifier. Any variables qualified as constant are read-only variables for that shader. Declaring variables as constant allows more descriptive shaders than using hard-wired numerical constants. The **const** qualifier can be used with any of the non-void transparent basic data types as well as structures and arrays of these. It is an error to write to a **const** variable outside of its declaration, so they must be initialized when declared. For example,

```
const vec3 zAxis = vec3 (0.0, 0.0, 1.0);
```

Structure fields may not be qualified with **const**. Structure variables can be declared as **const**, and initialized with a structure constructor.

Initializers for **const** declarations must be constant expressions, as defined in section 4.3.3 (“Constant Expressions”).

4.3.3 Constant Expressions

A *constant expression* is one of

- a literal value (e.g. **5** or **true**)
- a global or local variable qualified as **const** (i.e., not including function parameters)
- an expression formed by an operator on operands that are all constant expressions, including getting an element of a constant array, or a field of a constant structure, or components of a constant vector. However, the sequence operator (,) and the assignment operators (=, +=, ...) are not included in the operators that can create a constant expression.
- the length() method on a compile-time sized array, whether or not the object itself is constant.
- a constructor whose arguments are all constant expressions
- a built-in function call whose arguments are all constant expressions, with the exception of the texture lookup functions. This rule excludes functions with a **void** return or functions that have an **out** parameter. The built-in functions **dFdx**, **dFdy**, and **fwidth** must return 0 when evaluated inside an initializer with an argument that is a constant expression.

Function calls to user-defined functions (non-built-in functions) cannot be used to form constant expressions.

Scalar, vector, matrix, array and structure variables are constant expressions if qualified as **const**. Opaque types cannot be constant expressions.

A *constant integral expression* is a constant expression that evaluates to a scalar signed or unsigned integer.

Constant expressions will be evaluated in an invariant way so as to create the same value in multiple shaders when the same constant expressions appear in those shaders. See section 4.8.1 (“The Invariant Qualifier”) for more details on how to create invariant expressions and section 4.7.3 (“Precision Qualifiers”) for detail on how expressions are evaluated.

4.3.4 Input Variables

Shader input variables are declared with the **in** storage qualifier. They form the input interface between previous stages of the OpenGL ES pipeline and the declaring shader. Input variables must be declared at global scope. Values from the previous pipeline stage are copied into input variables at the beginning of shader execution. Variables declared as **in** may not be written to during shader execution. Only the input variables that are actually read need to be written by the previous stage; it is allowed to have superfluous declarations of input variables.

See section 7 (“Built-in Variables”) for a list of the built-in input names.

Vertex shader input variables (or attributes) receive per-vertex data. It is an error to use auxiliary storage or interpolation qualifiers in a vertex shader input. The values copied in are established by the OpenGL ES API or through the use of the layout identifier *location*.

It is a compile-time error to declare a vertex shader input with, or that contains, any of the following types:

- A *boolean type*
- An *opaque type*
- An array
- A structure

Example declarations in a vertex shader:

```
in vec4 position;  
in vec3 normal;
```

It is expected that graphics hardware will have a small number of fixed vector locations for passing vertex inputs. Therefore, the OpenGL ES Shading language defines each non-matrix input variable as taking up one such vector location. There is an implementation dependent limit on the number of locations that can be used, and if this is exceeded it will cause a link error. (Declared input variables that are not statically used do not count against this limit.) A scalar input counts the same amount against this limit as a **vec4**, so applications may want to consider packing groups of four unrelated float inputs together into a vector to better utilize the capabilities of the underlying hardware. A matrix input will use up multiple locations. The number of locations used will equal the number of columns in the matrix.

Tessellation control, evaluation, and geometry shader input variables get the per-vertex values written out by output variables of the same names in the previous active shader stage. For these inputs, centroid and interpolation qualifiers are allowed, but have no effect. Since tessellation control, tessellation evaluation, and geometry shaders operate on a set of vertices, each input variable (or input block, see interface blocks below) needs to be declared as an array. For example,

```
in float foo[]; // geometry shader input for vertex "out float foo"
```

Each element of such an array corresponds to one vertex of the primitive being processed. Each array can optionally have a size declared. For geometry shaders, the array size will be set by, (or if provided must be consistent with) the input **layout** declaration(s) establishing the type of input primitive, as described later in section 4.4.1 (“Input Layout Qualifiers”).

Some inputs and outputs are arrayed, meaning that for an interface between two shader stages either the input or output declaration requires an extra level of array indexing for the declarations to match. For example, with the interface between a vertex shader and a geometry shader, vertex shader output variables and geometry shader input variables of the same name must have matching types, except that the geometry shader will have one more array dimension than the vertex shader, to allow for vertex indexing. If such an arrayed interface variable is not declared with the necessary additional input or output array dimension, a link-time error will result. Geometry shader inputs, tessellation control shader inputs and outputs, and tessellation evaluation inputs all have an additional level of arrayness relative to other shader inputs and outputs.

For non-arrayed interfaces (meaning array dimensionally stays the same between stages), it is a link-time error if the input variable is not declared with the same type, including array dimensionality, as the matching output variable.

The link-time type-matching rules apply to all declared input and output variables, whether or not they are used.

Additionally, tessellation evaluation shaders support per-patch input variables declared with the `patch` and in qualifiers. Per-patch input variables are filled with the values of per-patch output variables written by the tessellation control shader. Per-patch inputs may be declared as one-dimensional arrays, but are not indexed by vertex number. Applying the `patch` qualifier to inputs can only be done in tessellation evaluation shaders. As with other input variables, per-patch inputs must be declared using the same type and qualification as per-patch outputs from the previous (tessellation control) shader stage. It is a compile-time error to use `patch` with inputs in any other stage.

It is a compile-time error to declare a tessellation control, tessellation evaluation or geometry shader input with, or that contains, any of the following types:

- A *boolean type*
- An *opaque type*
- An array, except for
 - per-vertex variables in the tessellation control, tessellation evaluation and geometry shaders
 - per-sample variables in the fragment shader
- A structure

Fragment shader inputs get per-fragment values, typically interpolated from a previous stage's outputs.

It is a compile-time error to declare a fragment shader input with, or that contains, any of the following types:

- A *boolean type*
- An *opaque type*
- An array of arrays
- An array of structures
- A structure containing an array
- A structure containing a structure

Fragment shader inputs that are, or contain, *integer types* must be qualified with the interpolation qualifier **flat**.

Fragment inputs are declared as in the following examples:

```
in vec3 normal;
centroid in vec2 TexCoord;
flat in vec3 myColor;
```

The fragment shader inputs form an interface with the last active shader in the vertex processing pipeline. For this interface, the last active shader stage output variables and fragment shader input variables of the same name must match in type and qualification, with a few exceptions: The storage qualifiers must, of course, differ (one is **in** and one is **out**). Also, auxiliary qualification (e.g. **centroid**) may differ. When auxiliary qualifiers do not match, those provided in the fragment shader supersede those provided in previous stages. If any such qualifiers are completely missing in the fragment shaders, then the default is used, rather than any qualifiers that may have been declared in previous stages. That is, what matters is what is declared in the fragment shaders, not what is declared in shaders in previous stages.

When an interface between shader stages is formed using shaders from two separate program objects, it is not possible to detect mismatches between inputs and outputs when the programs are linked. When there are mismatches between inputs and outputs on such interfaces, attempting to use the two programs in the same program pipeline will result in program pipeline validation failures, as described in section 7.4.1 (“Shader Interface Matching”) of the OpenGL ES 3.2 Specification.

Shaders can ensure matches across such interfaces either by using input and output layout qualifiers (sections 4.4.1 “Input Layout Qualifiers” and 4.4.2 “Output Layout Qualifiers”) or by using identical input and output declarations of blocks or variables. Complete rules for interface matching are found in section 7.4.1 “Shader Interface Matching” of the OpenGL ES 3.2 Graphics System Specification.

Compute shaders do not permit user-defined input variables and do not form a formal interface with any other shader stage. See section 7.1.6 (“Compute Shader Special Variables”) for a description of built-in compute shader input variables. All other input to a compute shader is retrieved explicitly through image loads, texture fetches, loads from uniforms or uniform buffers, or other user supplied code.

4.3.5 Uniform Variables

The **uniform** qualifier is used to declare global variables whose values are the same across the entire primitive being processed. All **uniform** variables are read-only. Except for variables declared within a uniform block, all uniform variables are initialized to 0 at link time and may be updated through the API.

Example declarations are:

```
uniform vec4 lightPosition;
```

The **uniform** qualifier can be used with any of the basic data types, or when declaring a variable whose type is a structure, or an array of any of these.

There is an implementation dependent limit on the amount of storage for uniforms that can be used for each type of shader and if this is exceeded it will cause a compile-time or link-time error. Uniform variables that are declared but not statically used do not count against this limit. The number of user-defined uniform variables and the number of built-in uniform variables that are used within a shader are added together to determine whether available uniform storage has been exceeded.

Uniforms in shaders all share a single global name space when linked into a program or separable program. Hence, the types, precisions and any location specifiers of all declared uniform variables with the same name must match across shaders that are linked into a single program. However it is not required to repeat the location specifier in all the linked shaders. While this single uniform name space is cross stage, a uniform variable name's scope is per stage: If a uniform variable name is declared in one stage (e.g., a vertex shader) but not in another (e.g., a fragment shader), then that name is still available in the other stage for a different use.

A compile or link error is generated if any of the explicitly given or compiler generated uniform locations is greater than the implementation-defined maximum number of uniform locations minus one.

Unlike locations for inputs and outputs, uniform locations are logical values, not register locations, and there is no concept of overlap. For example:

```
layout (location = 2) uniform mat4 x;
layout (location = 3) uniform mat4 y; // No overlap with x

layout(location = 2) in mat4 x;
layout(location = 3) in mat4 y; // Error, locations conflict with x
```

4.3.6 Output Variables

Shader output variables are declared with the **out** storage qualifier. They form the output interface between the declaring shader and the subsequent stages of the OpenGL ES pipeline. Output variables must be declared at global scope. During shader execution they will behave as normal unqualified global variables. Their values are copied out to the subsequent pipeline stage on shader exit. Only output variables that are read by the subsequent pipeline stage need to be written; it is allowed to have superfluous declarations of output variables.

There is *not* an **inout** storage qualifier for declaring a single variable name as both input and output to a shader. Also, a variable cannot be declared with both the in and the out qualifiers, this will result in an error. Output variables must be declared with different names than input variables, where the names include the instance name of any enclosing block.

Vertex, tessellation evaluation and geometry output variables output per-vertex data. and are declared using the **out** storage qualifier. Applying **patch** to an output can only be done in a tessellation control shader. It is a compile-time error to use **patch** on outputs in any other stage.

It is a compile-time error to declare a vertex, tessellation evaluation, tessellation control or geometry shader output with, or that contains, any of the following types:

- A *boolean type*
- An *opaque type*
- An array of arrays
- An array of structures (except for redeclaring the built-in `gl_PerVertex` structure)
- A structure containing an array
- A structure containing a structure

Vertex shader outputs that are, or contain, integer types must be qualified with the interpolation qualifier **flat**.

Individual outputs are declared as in the following examples:

```
out vec3 normal;
centroid out vec2 TexCoord;
invariant centroid out vec4 Color;
flat out vec3 myColor;
sample out vec4 perSampleColor;
```

These can also appear in interface blocks, as described in section 4.3.9 (“Interface Blocks”). Interface blocks allow simpler addition of arrays to the interface from vertex to geometry shader. They also allow a fragment shader to have the same input interface as a geometry shader for a given vertex shader.

Tessellation control shader output variables are used to output per-vertex and per-patch data. Per-vertex output variables are arrayed (see arrayed under section 4.3.4 “Input Variables”) and declared using the **out** qualifier without the **patch** qualifier. Per-patch output variables are declared using the **patch** and **out** qualifiers.

Since tessellation control shaders produce an arrayed primitive comprising multiple vertices, each per-vertex output variable (or output block, see interface blocks below) needs to be declared as an array. For example,

```
out float foo[]; // feeds next stage input "in float foo[]"
```

Each element of such an array corresponds to one vertex of the primitive being produced. Each array can optionally have a size declared. The array size will be set by (or if provided must be consistent with) the output layout declaration(s) establishing the number of vertices in the output patch, as described later in section 4.4.2.1 (“Tessellation Control Outputs”).

Each tessellation control shader invocation has a corresponding output patch vertex, and may assign values to per-vertex outputs only if they belong to that corresponding vertex. If a per-vertex output variable is used as an l-value, it is a compile-time or link-time error if the expression indicating the vertex index is not the identifier *gl_InvocationID*.

The order of execution of a tessellation control shader invocation relative to the other invocations for the same input patch is undefined unless the built-in function `barrier()` is used. This provides some control over relative execution order. When a shader invocation calls `barrier()`, its execution pauses until all other invocations have reached the same point of execution. Output variable assignments performed by any invocation executed prior to calling `barrier()` will be visible to any other invocation after the call to `barrier()` returns.

Because tessellation control shader invocations execute in undefined order between barriers, the values of per-vertex or per-patch output variables will sometimes be undefined. Consider the beginning and end of shader execution and each call to `barrier()` as synchronization points. The value of an output variable will be undefined in any of the three following cases:

1. At the beginning of execution.
2. At each synchronization point, unless
 - the value was well-defined after the previous synchronization point and was not written by any invocation since, or
 - the value was written by exactly one shader invocation since the previous synchronization point, or

- the value was written by multiple shader invocations since the previous synchronization point, and the last write performed by all such invocations wrote the same value.
3. When read by a shader invocation, if
- the value was undefined at the previous synchronization point and has not been written by the same shader invocation since, or
 - the output variable is written to by any other shader invocation between the previous and next synchronization points, even if that assignment occurs in code following the read.

Fragment outputs output per-fragment data and are declared using the **out** storage qualifier. It is an error to use auxiliary storage or interpolation qualifiers in a fragment shader output declaration. It is a compile-time error to declare a fragment shader output with, or that contains, any of the following types :

- A *boolean type*
- An *opaque type*
- A matrix
- A structure
- An array of arrays

Fragment shader outputs declared as arrays may only be indexed by a constant integral expression.

Fragment outputs are declared as in the following examples:

```
out vec4 FragmentColor;
out uint Luminosity;
```

Compute shaders have no built-in output variables, do not support user-defined output variables and do not form a formal interface with any other shader stage. All outputs from a compute shader take the form of the side effects such as image stores and operations on atomic counters.

4.3.7 Buffer Variables

The **buffer** qualifier is used to declare global variables whose values are stored in the data store of a buffer object bound through the OpenGL ES API. Buffer variables can be read and written, with the underlying storage shared among all active shader invocations. Buffer variable memory reads and writes within a single shader invocation are processed in order. However, the order of reads and writes performed in one invocation relative to those performed by another invocation is largely undefined. Buffer variables may be qualified with memory access qualifiers affecting how the underlying memory is accessed, as described in 4.10 (“Memory Access Qualifiers”).

The **buffer** qualifier can be used with any type except:

- An opaque type

Buffer variables may only be declared inside interface blocks (section 4.3.9 “Interface Blocks”), which are then referred to as *shader storage blocks*¹. It is a compile-time error to declare buffer variables at global scope (outside a block). Buffer variables cannot have initializers.

```
// use buffer to create a buffer block (shader storage block)
buffer BufferName { // externally visible name of buffer
    int count;      // typed, shared memory...
    ...             // ...
    vec4 v[];       // last element may be an array that is not sized
                   // until after link time (dynamically sized)
} Name;            // name of block within the shader
```

There are implementation-dependent limits on the number of the shader storage blocks used for each type of shader, the combined number of shader storage blocks used for a program, and the amount of storage required by each individual shader storage block. If any of these limits are exceeded, it will cause a compile-time or link-time error.

If multiple shaders are linked together, then they will share a single global buffer variable name space. Hence, the types of declared buffer variables with the same name must match across all shaders that are linked into a single program.

Precision qualifiers for such variables need not match.

4.3.8 Shared Variables

The **shared** qualifier is used to declare variables that have storage shared between all work items in a compute shader local work group. Variables declared as **shared** may only be used in compute shaders (see section 2.6 “Compute Processor”). Shared variables are implicitly coherent (see section 4.10 “Memory Access Qualifiers”).

Variables declared as **shared** may not have initializers and their contents are undefined at the beginning of shader execution. Any data written to shared variables will be visible to other work items (executing the same shader) within the same local work group.

In the absence of synchronization, the order of reads and writes to the same shared variable by different invocations of a shader is not defined.

In order to achieve ordering with respect to reads and writes to shared variables, a combination of control flow and memory barriers must be employed using the **barrier()** and **memoryBarrier()** functions (see section 8.15 “Shader Invocation Control Functions”).

There is a limit to the total size of all variables declared as shared in a single program. This limit, expressed in units of basic machine units may be determined by using the OpenGL API to query the value of `MAX_COMPUTE_SHARED_MEMORY_SIZE`.

¹ The terms *shader storage buffer (object)*, and *shader storage block* are often used interchangeably. The former generally refers to the underlying storage whereas the latter refers only to the interface definition in the shader. Similarly for *uniform buffer objects* and *uniform blocks*.

4.3.9 Interface Blocks

Input, output, uniform and buffer variable declarations can be grouped into named interface blocks to provide coarser granularity backing than is achievable with individual declarations. They can have an optional instance name, used in the shader to reference their members. An output block of one programmable stage is backed by a corresponding input block in the following programmable stage. A *uniform block* is backed by the application with a buffer object. A block of buffer variables, called a *shader storage block*, is also backed by the application with a buffer object. It is a compile-time error to have an input block in a vertex shader or an output block in a fragment shader.

An interface block is started by an **in**, **out**, **uniform** or **buffer** keyword, followed by a block name, followed by an open curly brace ({) as follows:

interface-block:

*layout-qualifier*_{opt} *interface-storage-qualifier* *block-name* { *member-list* } *instance-name*_{opt} ;

layout-qualifier:

layout (*layout-qualifier-id-list*)

layout-qualifier-id-list:

layout-qualifier-id

layout-qualifier-id , *layout-qualifier-id-list*

interface-storage-qualifier:

in

out

patch in // Note: qualifiers can be in any order

patch out

uniform

buffer

member-list:

member-declaration

member-declaration *member-list*

member-declaration:

*layout-qualifier*_{opt} *qualifiers*_{opt} *type declarators* ; // Note: qualifiers can be in any order

instance-name:

identifier

identifier [*constant-integral-expression*]

Each of the above elements is discussed below, with the exception of layout qualifiers (*layout-qualifier*), which are defined in the next section.

First, an example,

4 Variables and Types

```
uniform Transform {
    mat4 ModelViewMatrix;
    mat4 ModelViewProjectionMatrix;
    uniform mat3 NormalMatrix;          // allowed restatement of qualifier
    float Deformation;
};
```

The above establishes a uniform block named “Transform” with four uniforms grouped inside it.

Types and declarators are the same as for other input, output, uniform and buffer variable declarations outside blocks, with these exceptions:

- Opaque types are not allowed
- Structure definitions cannot be nested inside a block
- Arrays of arrays of blocks are not allowed

If no optional qualifier is used in a member-declaration, the qualification of the member includes all **in**, **out**, **patch**, **uniform**, or **buffer** as determined by *interface-qualifier*. If optional qualifiers are used, they can include interpolation qualifiers, auxiliary storage qualifiers, and storage qualifiers and they must declare an input, output, or uniform member consistent with the interface qualifier of the block: Input variables, output variables, uniform variables, and **buffer** members can only be in in blocks, out blocks, uniform blocks, and shader storage blocks, respectively.

Repeating the **in**, **out**, **uniform** or **buffer** interface qualifier for a member's storage qualifier is optional. For example,

```
uniform Transform
{
    uniform mat4 model_view; // legal, uniform inside a uniform block.
    mat4 projection;        // legal, 'uniform' inherited from block.
    in bool transform_flag; // illegal, conflicting storage qualifiers
};
```

A *shader interface* is defined to be one of these:

- All the uniform variables and uniform blocks declared in a program. This spans all compilation units linked together within one program.
- All shader storage blocks.
- The boundary between adjacent programmable pipeline stages: This spans all the outputs declared in the first stage and all the inputs declared in the second stage. Note that for the purposes of this definition, the fragment shader and the preceding shader are considered to have a shared boundary even though in practice, all values passed to the fragment shader first pass through the rasterizer and interpolator.

4 Variables and Types

The block name (*block-name*) is used to match across shader interfaces: an output block of one pipeline stage will be matched to an input block with the same name in the subsequent pipeline stage. For uniform or shader storage blocks, the application uses the block name to identify the block. Block names have no other use within a shader beyond interface matching; it is an error to use a block name at global scope for anything other than as a block name (e.g., use of a block name for a global variable name or function name is currently reserved). It is a compile-time error to use the same block name for more than one block declaration in the same *shader interface* (as defined above) within one shader, even if the block contents are identical.

Matched block names within a *shader interface* (as defined above) must match in terms of having the same number of declarations with the same sequence of types, precisions and the same sequence of member names, as well as having the same member-wise layout qualification (see next section). Furthermore, if a matching block is declared as an array, then the array sizes must also match (or follow array matching rules for the shader interface between consecutive shader stages). A block name is allowed to have different definitions in different shader interfaces within the same shader, allowing, for example, an input block and output block to have the same name..

4 Variables and Types

If an instance name (*instance-name*) is not used, the names declared inside the block are scoped at the global level and accessed as if they were declared outside the block. If an instance name (*instance-name*) is used, then it puts all the members inside a scope within its own name space, accessed with the field selector (.) operator (analogously to structures). For example,

```
uniform Transform_1
{
    mat4 modelview;
};

uniform Transform_2
{
    mat4 projection;
} transform_2;

mat4 modelview; // illegal as modelview already defined at this scope
mat4 projection; // legal as projection and transform_2.projection are
                // distinct.

in Light
{
    vec4 LightPos;
    vec3 LightColor;
};

in ColoredTexture
{
    vec4 Color;
    vec2 TexCoord;
} Material;      // instance name

vec3 Color;      // different Color than Material.Color
vec4 LightPos;   // illegal, already defined
...
... = LightPos;  // accessing LightPos
... = Material.Color; // accessing Color in ColoredTexture block
```

Matched uniform or shader storage block names (but not input or output block names) must also either all be lacking an instance name or all having an instance name, thereby putting their members at the same scoping level. When instance names are present on matched block names, it is allowed for the instance names to differ; they need not match for the blocks to match.

Outside the shading language (i.e., in the API), members are similarly identified except the block name is always used in place of the instance name (API accesses are to interfaces, not to shaders). If there is no instance name, then the API does not use the block name to access a member, just the member name. For example:

4 Variables and Types

```
uniform Transform_1
{
    mat4 modelview; // API will use "modelview"
};

uniform Transform_2
{
    mat4 projection; // API will use "Transform_2.projection"
} transform_2;
```

Within a *shader interface*, all declarations of the same global name must be for the same object and must match in type and in whether they declare a variable or member of a block with no instance name. The API also needs this name to uniquely identify an object in the shader interface. It is a link-time error if any particular shader interface contains

- two different blocks, each having no instance name, and each having a member of the same name

or

- a variable outside a block, and a block with no instance name, where the variable has the same name as a member in the block.

For blocks declared as arrays, the array index must also be included when accessing members, as in this example

```
uniform Transform { // API uses "Transform[2]" to refer to instance 2
    mat4          ModelViewMatrix;
    mat4          ModelViewProjectionMatrix;
    float         Deformation;
} transforms[4];

...
... = transforms[2].ModelViewMatrix; // shader access of instance 2
// API uses "Transform.ModelViewMatrix" to query an offset or other query
```

For uniform or shader storage blocks declared as an array, each individual array element corresponds to a separate buffer object bind range, backing one instance of the block. As the array size indicates the number of buffer objects needed, uniform and shader storage block array declarations must specify an array size. All indices used to index a shader storage block array must be constant integral expressions. A uniform block array can only be indexed with a dynamically uniform integral expression, otherwise results are undefined.

When using OpenGL ES API entry points to identify the name of an individual block in an array of blocks, the name string may include an array index (e.g., *Transform[2]*). When using OpenGL ES API entry points to refer to offsets or other characteristics of a block member, an array index must not be specified (e.g., *Transform.ModelViewMatrix*). See section 7.3.1 ("Program Interfaces") in the OpenGL ES 3.2 specification for details.

Tessellation control, tessellation evaluation and geometry shader input blocks must be declared as arrays and follow the array declaration and linking rules for all shader inputs for the respective stages. All other input and output block arrays must specify an array size.

There are implementation-dependent limits on the number of uniform blocks and the number of shader storage blocks that can be used per stage. If either limit is exceeded, it will cause a link error.

4.4 Layout Qualifiers

Layout qualifiers can appear in several forms of declaration. They can appear as part of an interface block definition or block member, as shown in the grammar in the previous section. They can also appear with just an *interface-qualifier* to establish layouts of other declarations made with that qualifier:

layout-qualifier interface-qualifier ;

Or, they can appear with an individual variable declared with an interface qualifier:

layout-qualifier interface-qualifier declaration ;

Declarations of layouts can only be made at global scope, and only where indicated in the following subsections; their details are specific to what the interface qualifier is, and are discussed individually.

Interface qualifiers are a subset of storage qualifiers:

interface-qualifier:

uniform
buffer
in
out

As shown in the previous section, *layout-qualifier* expands to:

layout-qualifier :

layout (layout-qualifier-id-list)

layout-qualifier-id-list :

layout-qualifier-id
layout-qualifier-id , layout-qualifier-id-list

layout-qualifier-id:

layout-qualifier-name
layout-qualifier-name = layout-qualifier-value
shared

The tokens in any *layout-qualifier-id-list* are identifiers, not keywords and they have their own name space. Generally, they can be listed in any order. Order-dependent meanings exist only if explicitly called out below. As for other identifiers, they are case sensitive.

The set of allowed layout qualifiers depends on the shader, the interface and the variable type as specified in the following sections. For example, a sampler in the default uniform block in a fragment shader can have *location* and *binding* layout qualifiers but no others. Invalid use of layout qualifiers is an error.

The following table summarizes the use of layout qualifiers. It shows for each one what kinds of declarations it may be applied to. These are all discussed in detail in the following sections.

4 Variables and Types

Layout Qualifier	Qualifier Only	Individual Variable	Block	Block Member	Allowed interfaces
shared packed std140 std430	X		X		uniform/buffer
row_major column_major	X		X	X	
binding =		opaque types only	X		
offset =		atomic_uint only			
location =		X			uniform/buffer
location =		X	X	X	all in/out , except for compute
triangles quads isolines	X				tessellation evaluation in
equal_spacing fractional_even_spacing fractional_odd_spacing	X				tessellation evaluation in
cw ccw	X				tessellation evaluation in
point_mode	X				tessellation evaluation in
points	X				geometry in/out
points lines lines_adjacency triangles triangles_adjacency	X				geometry in
invocations =	X				geometry in
early_fragment_tests	X				fragment in
local_size_x = local_size_y = local_size_z =	X				compute in
vertices =	X				tessellation control out

Layout Qualifier	Qualifier Only	Individual Variable	Block	Block Member	Allowed interfaces
[points] line_strip triangle_strip	X				geometry out
max_vertices =	X				geometry out
rgba32f rgba16f r32f rgba8 rgba8_snorm rgba32i rgba16i rgba8i r32i rgba32ui rgba16ui rgba8ui r32ui		image types only			uniform
blend_support_multiply blend_support_screen blend_support_overlay blend_support_darken blend_support_lighten blend_support_colordodge blend_support_colorburn blend_support_hardlight blend_support_softlight blend_support_difference blend_support_exclusion blend_support_hsl_hue blend_support_hsl_saturation blend_support_hsl_color blend_support_hsl_luminosity blend_support_all_equations	X				fragment out

4.4.1 Input Layout Qualifiers

All shaders except compute shaders allow input layout qualifiers on input variable declarations. The layout qualifier identifier for shader inputs is:

layout-qualifier-id:

location = *integer-constant*

Only one argument is accepted. For example,

```
layout(location = 3) in vec4 normal;
```

will establish that the shader input *normal* is copied in from vector location number 3. For vertex shader inputs, the location specifies the number of the vertex attribute from which input values are taken. For inputs of all other shader types, the location specifies a vector number that can be used to match against outputs from a previous shader stage, even if that shader is in a different program object.

The following language describes how many locations are consumed by a given type. However, geometry shader inputs, tessellation control shader inputs and outputs, and tessellation evaluation inputs all have an additional level of arrayness relative to other shader inputs and outputs. This outer array level is removed from the type before considering how many locations the type consumes.

If a shader input is any scalar or vector type, it will consume a single location.

If the declared input (after potentially removing an outer array level as just described above) is an array of size n and each of the elements takes m locations, it will be assigned $m * n$ consecutive locations starting with the location specified. For example,

```
layout (location = 6) in vec4 colors[3];
```

will establish that the shader input *colors* is assigned to vector location numbers 6, 7, and 8.

If the declared input is an $n \times m$ matrix, it will be assigned multiple locations starting with the location specified. The number of locations assigned for each matrix will be the same as for an n -element array of m -component vectors. For example,

```
layout (location = 9) in mat4 transforms[2];
```

will establish that shader input *transforms* is assigned to vector locations 9-16, with *transforms[0]* being assigned to locations 9-12, and *transforms[1]* being assigned to locations 13-16.

If the declared input is a structure or block, its members will be assigned consecutive locations in their order of declaration, with the first member assigned the location provided in the layout qualifier. For a structure, this process applies to the entire structure. It is a compile-time error to use a location qualifier on a member of a structure. For a block, this process applies to the entire block, or until the first member is reached that has a location layout qualifier.

When a block member is declared with a location qualifier, its location comes from that qualifier: The member's location qualifier overrides the block-level declaration. Subsequent members are again assigned consecutive locations, based on the newest location, until the next member declared with a location qualifier. The values used for locations do not have to be declared in increasing order.

If a block has no block-level location layout qualifier, it is required that either all or none of its members have a location layout qualifier, or a compile-time error results.

4 Variables and Types

The locations consumed by block and structure members are determined by applying the rules above recursively as though the structure member were declared as an input variable of the same type. For example:

```
layout(location = 3) in struct S
{
    vec3 a;                // gets location 3
    mat2 b;                // gets locations 4 and 5
    vec4 c[2];             // gets locations 6 and 7
    layout (location = 8) vec2 A; // ERROR, can't use on struct member
} s;

layout(location = 4) in block
{
    vec4 d;                // gets location 4
    vec4 e;                // gets location 5
    layout(location = 7) vec4 f; // gets location 7
    vec4 g;                // gets location 8
    layout (location = 1) vec4 h; // gets location 1
    vec4 i;                // gets location 2
    vec4 j;                // gets location 3
    vec4 k;                // ERROR, location 4 already used
};
```

The number of input locations available to a shader is limited. For vertex shaders, the limit is the advertised number of vertex attributes. For all other shaders, the limit is implementation-dependent and must be no less than one fourth of the advertised maximum input component count.

A program will fail to link if any attached shader uses a location greater than or equal to the number of supported locations, unless device-dependent optimizations are able to make the program fit within available hardware resources.

A program will fail to link if explicit location assignments leave the linker unable to find space for other variables without explicit assignments. For the purposes of determining if a non-vertex input matches an output from a previous shader stage, the location layout qualifier (if any) must match.

If a vertex shader input variable with no location assigned in the shader text has a location specified through the OpenGL ES API, the API-assigned location will be used. Otherwise, such variables will be assigned a location by the linker. See section 11.1.1 “Vertex Attributes” of the OpenGL ES 3.2 Graphics System Specification for more details.

It is an error if more than one input or element of a matrix input is bound to the same location.

4.4.1.1 Tessellation Evaluation Inputs

Additional input layout qualifier identifiers allowed for tessellation evaluation shaders are described below.

layout-qualifier-id:
primitive_mode
vertex_spacing
ordering
point_mode

The *primitive-mode* is used to specify a tessellation primitive mode to be used by the tessellation primitive generator.

primitive-mode:
triangles
quads
isolines

If present, the *primitive-mode* specifies that the tessellation primitive generator should subdivide a triangle into smaller triangles, a quad into triangles, or a quad into a collection of lines, respectively.

A second group of layout identifiers, *vertex spacing*, is used to specify the spacing used by the tessellation primitive generator when subdividing an edge.

vertex-spacing:
equal_spacing
fractional_even_spacing
fractional_odd_spacing

equal_spacing specifies that edges should be divided into a collection of equal-sized segments;

fractional_even_spacing specifies that edges should be divided into an even number of equal-length segments plus two additional shorter “fractional” segments; or

fractional_odd_spacing specifies that edges should be divided into an odd number of equal-length segments plus two additional shorter “fractional” segments.

A third group of layout identifiers, *ordering*, specifies whether the tessellation primitive generator produces triangles in clockwise or counter-clockwise order, according to the coordinate system depicted in the OpenGL ES Specification.

ordering:
cw
ccw

The identifiers **cw** and **ccw** indicate clockwise and counter-clockwise triangles, respectively. If the tessellation primitive generator does not produce triangles, the order is ignored.

Finally, *point mode* indicates that the tessellation primitive generator should produce one point for each distinct vertex in the subdivided primitive, rather than generating lines or triangles.

point-mode:
point_mode

Any or all of these identifiers may be specified one or more times in a single input layout declaration. The tessellation evaluation shader object in a program must declare a primitive mode in its input layout. Declaring vertex spacing, ordering or point mode identifiers is optional. If spacing or vertex order declarations are omitted, the tessellation primitive generator will use equal spacing or counter-clockwise vertex ordering, respectively. If a point mode declaration is omitted, the tessellation primitive generator will produce lines or triangles according to the primitive mode.

4.4.1.2 Geometry Shader Inputs

Additional layout qualifier identifiers for geometry shader inputs include primitive identifiers and an invocation count identifier:

```
layout-qualifier-id
    points
    lines
    lines_adjacency
    triangles
    triangles_adjacency
    invocations = integer-constant
```

The identifiers **points**, **lines**, **lines_adjacency**, **triangles**, and **triangles_adjacency** are used to specify the type of input primitive accepted by the geometry shader, and only one of these is accepted. The geometry shader must declare this input primitive layout.

The identifier **invocations** is used to specify the number of times the geometry shader executable is invoked for each input primitive received. Invocation count declarations are optional. If no invocation count is declared in the geometry shader, it will be run once for each input primitive. If an invocation count is declared, all such declarations must specify the same count. If a shader specifies an invocation count greater than the implementation-dependent maximum, it will fail to compile.

For example,

```
layout(triangles, invocations = 6) in;
```

will establish that all inputs to the geometry shader are triangles and that the geometry shader executable is run six times for each triangle processed.

All geometry shader input unsized array declarations will be sized by an earlier input primitive layout qualifier, when present, as per the following table.

Layout	Size of Input Arrays
points	1
lines	2
lines_adjacency	4
triangles	3
triangles_adjacency	6

The intrinsically declared input array `gl_in[]` will also be sized by any input primitive-layout declaration. Hence, the expression

```
gl_in.length()
```

will return the value from the table above.

An input can be declared without an array size if there is a previous layout which specifies the size. For built-in inputs (e.g. `gl_in[]`), a layout must be declared before any use.

It is a compile-time error if a layout declaration's array size (from the table above) does not match all the explicit array sizes specified in declarations of an input variables in the same shader. The following includes examples of compile-time errors:

```
// code sequence within one shader...
in vec4 Color2[2];      // legal, size is 2
in vec4 Color3[3];      // illegal, input sizes are inconsistent
layout(lines) in;       // legal for Color2, input size is 2, matching Color2
in vec4 Color4[3];      // illegal, contradicts layout of lines
layout(lines) in;       // legal, matches other layout() declaration
layout(triangles) in;   // illegal, does not match earlier layout()
                        // declaration
```

It is a link-time error if not all provided sizes (sized input arrays and layout size) match in the geometry shader of a program.

4.4.1.3 Fragment Shader Inputs

Fragment shaders allow the following layout qualifier on **in** only (not with variable declarations):

layout-qualifier-id:
early_fragment_tests

to request that fragment tests be performed before fragment shader execution, as described in Section 13.6 (“Early Fragment Tests”) of the OpenGL ES Specification.

For example,

```
layout(early_fragment_tests) in;
```

Specifying this will make per-fragment tests be performed before fragment shader execution. In addition it is an error to statically write to `gl_FragDepth` in the fragment shader. If this is not declared, per-fragment tests will be performed after fragment shader execution.

4.4.1.4 Compute Shader Inputs

There are no layout location qualifiers for compute shader inputs.

Layout qualifier identifiers for compute shader inputs are the work-group size qualifiers:

layout-qualifier-id:

local_size_x = *integer-constant*

local_size_y = *integer-constant*

local_size_z = *integer-constant*

The **local_size_x**, **local_size_y**, and **local_size_z** qualifiers are used to declare a fixed local group size by the compute shader in the first, second, and third dimension, respectively. The default size in each dimension is 1.

For example, the following declaration in a compute shader

```
layout (local_size_x = 32, local_size_y = 32) in;
```

is used to declare a two-dimensional compute shader with a local size of 32 X 32 elements, which is equivalent to a three-dimensional compute shader where the third dimension has size one.

As another example, the declaration

```
layout (local_size_x = 8) in;
```

effectively specifies that a one-dimensional compute shader is being compiled, and its size is 8 elements.

If the fixed local group size of the shader in any dimension is greater than the maximum size supported by the implementation for that dimension, a compile-time error results. Also, if such a layout qualifier is declared more than once in the same shader, all those declarations must set the same set of local work-group sizes and set them to the same values; otherwise a compile-time error results.

Furthermore, if a program object contains a compute shader, that shader must contain an input layout qualifier specifying a fixed local group size for the program, or a link-time error will occur.

4.4.2 Output Layout Qualifiers

Some output layout qualifiers apply to all shader stages and some apply only to specific stages. The latter are discussed in separate sections below.

As with input layout qualifiers, all shaders except compute shaders allow location layout qualifiers on output variable declarations, output block declarations, and output block member declarations.

The layout qualifier identifier for shader outputs is:

layout-qualifier-id:

location = *integer-constant*

The usage and rules for applying the location qualifier to blocks and structures are exactly as described in section 4.4.1 (“Input Layout Qualifiers”).

The qualifier may appear at most once within a declaration. For example, in a fragment shader,

```
layout(location = 3) out vec4 color;
```

will establish that the fragment shader output *color* is copied out to draw buffer 3.

4 Variables and Types

For fragment shader outputs, the location specifies the color output number receiving the values of the output. For outputs of all other shader stages, the location specifies a vector number that can be used to match against inputs in a subsequent shader stage, even if that shader is in a different program object.

Declared outputs of scalar or vector type consume a single location.

If the declared output is an array, it will be assigned consecutive locations starting with the location specified. For example, in a fragment shader,

```
layout(location = 2) out vec4 colors[3];
```

will establish that *colors* is copied out to locations (i.e. draw buffers) 2, 3, and 4.

If the declared output is an $n \times m$ matrix, it will be assigned multiple locations starting with the location specified. The number of locations assigned will be the same as for an n -element array of m -component vectors.

If the declared output is a structure, its members will be assigned consecutive locations in the order of declaration, with the first member assigned the location specified for the structure. The number of locations consumed by a structure member is determined by applying the rules above recursively as though the structure member were declared as an output variable of the same type.

Location layout qualifiers may be used on output variables declared as structures, but it is a compile-time error to use them on individual members. Location layout qualifiers may be used on output blocks and output block members.

The number of output locations available to a shader is limited. For fragment shaders, the limit is the advertised number of draw buffers.

For all other shaders, the limit is implementation-dependent and must be no less than one fourth of the advertised maximum output component count (compute shaders have no outputs.) A program will fail to link if any attached shader uses a location greater than or equal to the number of supported locations, unless device-dependent optimizations are able to make the program fit within available hardware resources.

Compile-time errors may also be given if at compile time it is known the link will fail. A negative output location will result in an error.

A program will fail to link if any of the following occur:

- any two fragment shader output variables are assigned to the same location
- if any two output variables from the same vertex, tessellation or geometry shader stage are assigned to the same location.

For all shader types, a program will fail to link if explicit location assignments leave the linker unable to find space for other variables without explicit assignments.

If an output variable has no location assigned in the shader text, it will be assigned a location by the linker. See section 3.9.2 “Shader Execution” of the OpenGL ES Specification for more details.

For the purposes of determining if a non-fragment output matches an input from a subsequent shader stage, the location layout qualifier (if any) must match.

4.4.2.1 Tessellation Control Outputs

Tessellation control shaders allow output layout qualifiers only on the interface qualifier **out**, not on an output block, block member, or variable declaration. The output layout qualifier identifiers allowed for tessellation control shaders are:

layout-qualifier-id
vertices = *integer-constant*

The identifier **vertices** specifies the number of vertices in the output patch produced by the tessellation control shader, which also specifies the number of times the tessellation control shader is invoked. It is a compile- or link-time error for the output vertex count to be less than or equal to zero, or greater than the implementation-dependent maximum patch size.

The intrinsically declared tessellation control output array `gl_out[]` will also be sized by any output layout declaration. Hence, the expression

```
gl_out.length()
```

will return the output patch vertex count specified in a previous output layout qualifier. For outputs declared without an array size, including intrinsically declared outputs (i.e., `gl_out`), a layout must be declared before any use of the method `length()` or other array use that requires its size to be known.

It is a compile-time error if the output patch vertex count specified in an output layout qualifier does not match the array size specified in any output variable declaration in the same shader.

All tessellation control shader layout declarations in a program must specify the same output patch vertex count. There must be at least one layout qualifier specifying an output patch vertex count in any program containing a tessellation control shader.

4.4.2.2 Geometry Outputs

Geometry shaders can have two additional types of output layout identifiers: an output primitive type and a maximum output vertex count. The primitive type and vertex count identifiers are allowed only on the interface qualifier **out**, not on an output block, block member, or variable declaration.

The layout qualifier identifiers for geometry shader outputs are

layout-qualifier-id
points
line_strip
triangle_strip
max_vertices = *integer-constant*

The primitive type identifiers `points`, `line_strip`, and `triangle_strip` are used to specify the type of output primitive produced by the geometry shader, and only one of these is accepted. The geometry shader object in a program must declare an output primitive type, and all geometry shader output primitive type declarations in a program must declare the same primitive type.

The vertex count identifier `max_vertices` is used to specify the maximum number of vertices the shader will ever emit in a single invocation. The geometry shader object in a program must declare a maximum output vertex count, and all geometry shader output vertex count declarations in a program must declare the same count.

In this example,

```
layout(triangle_strip, max_vertices = 60) out; // order does not matter
layout(max_vertices = 60) out;               // redeclaration okay
layout(triangle_strip) out;                  // redeclaration okay
layout(points) out;                          // error, contradicts triangle_strip
layout(max_vertices = 30) out;               // error, contradicts 60
```

all outputs from the geometry shader are triangles and at most 60 vertices will be emitted by the shader. It is an error for the maximum number of vertices to be greater than `gl_MaxGeometryOutputVertices`.

All geometry shader output layout declarations in a program must declare the same layout and same value for `max_vertices`. If geometry shaders are in a program, there must be at least one geometry output layout declaration somewhere in that program.

4.4.2.3 Fragment Outputs

Fragment shaders can have an output layout for redeclaring the built-in variable `gl_FragDepth`:

```
out float gl_FragDepth; // redeclaration that changes nothing is allowed
```

The built-in `gl_FragDepth` is only predeclared in fragment shaders, so redeclaring it in any other shader language results in an error.

If there is only a single output, the location does not need to be specified, in which case it defaults to zero. This applies for all output types, including arrays. For example,

```
out vec4 my_FragColor; // must be the only output declaration
```

will establish that the fragment shader output `my_FragColor` is copied out to draw buffer 0. Likewise,

```
out vec4 my_FragData[4]; // must be the only output declaration
```

will establish that the fragment shader outputs `my_FragData[0]` to `my_FragData[3]` is copied out to draw buffers 0 through 3 respectively.

If there is more than one fragment output, the location must be specified for all outputs. It is an error if any of the following occur:

- The location of any fragment output or element of a fragment array output, is greater or equal to the value of `MAX_DRAW_BUFFERS`.
- More than one fragment output or element of a fragment array output is bound to the same location.

See section 11.1.3 “Shader Execution” of the OpenGL ES 3.2 Graphics System Specification for more details.

Fragment shaders additionally support layout qualifiers specifying a set of advanced blend equations supported when the fragment shader is used. These layout qualifiers are only permitted on the interface qualifier out, and use the identifiers specified in the “Layout Qualifier” column of the table below.

If a layout qualifier in the table below is specified in the fragment shader, the fragment shader may be used with the corresponding advanced blend equation in the “Blend Equation(s) Supported” column. Additionally, the special qualifier **blend_support_all_equations** indicates that the shader may be used with any advanced blending equation supported by the OpenGL ES Specification. It is not an error to specify more than one of these identifiers in any fragment shader. Specifying more than one qualifier or **blend_support_all_equations** means that the fragment shader may be used with multiple advanced blend equations. Additionally, it is not an error to specify any single one of these layout qualifiers more than once.

Layout Qualifier	Blend Equations(s) Supported
blend_support_multiply	MULTIPLY
blend_support_screen	SCREEN
blend_support_overlay	OVERLAY
blend_support_darken	DARKEN
blend_support_lighten	LIGHTEN
blend_support_colordodge	COLORDODGE
blend_support_colorburn	COLORBURN
blend_support_hardlight	HARDLIGHT
blend_support_softlight	SOFTLIGHT
blend_support_difference	DIFFERENCE
blend_support_exclusion	EXCLUSION
blend_support_hsl_hue	HSL_HUE
blend_support_hsl_saturation	HSL_SATURATION
blend_support_hsl_color	HSL_COLOR
blend_support_hsl_luminosity	HSL_LUMINOSITY
blend_support_all_equations	all blend equations

4.4.3 Uniform Variable Layout Qualifiers

The following layout qualifier can be used for all default-block uniform variables but not for variables in uniform or shader storage blocks. The layout qualifier identifier for uniform variables is:

layout-qualifier-id:

location = *integer-constant*

The location specifies the location by which the OpenGL ES API can reference the uniform and update its value. Individual elements of a uniform array are assigned consecutive locations with the first element taking location *location*. No two default-block uniform variables in the program can have the same location, even if they are unused, otherwise a compiler or linker error will be generated. Valid locations for default-block uniform variable locations are in the range of 0 to the implementation-defined maximum number of uniform locations minus one.

Locations can be assigned to default-block uniform arrays and structures. The first inner-most scalar, vector or matrix member or element takes the specified *location* and the compiler assigns the next inner-most member or element the next incremental location value. Each subsequent inner-most member or element gets incremental locations for the entire structure or array. This rule applies to nested structures and arrays and gives each inner-most scalar, vector, or matrix member a unique location. When the linker generates locations for uniforms without an explicit location, it assumes for all uniforms with an explicit location all their array elements and structure members are used and the linker will not generate a conflicting location, even if that element or member is deemed unused.

4.4.4 Uniform and Shader Storage Block Layout Qualifiers

Layout qualifiers can be used for uniform and shader storage blocks, but not for non-block uniform declarations. The layout qualifier identifiers for uniform and shader storage blocks are:

layout-qualifier-id:

- shared**
- packed**
- std140**
- std430**
- row_major**
- column_major**
- binding** = *integer-constant*

None of these have any semantic effect at all on the usage of the variables being declared; they only describe how data is laid out in memory. For example, matrix semantics are always column-based, as described in the rest of this specification, no matter what layout qualifiers are being used.

Uniform and shader storage block layout qualifiers can be declared for global scope, on a single uniform or shader storage block, or on a single block member declaration.

Default layouts are established at global scope for uniform blocks as:

```
layout (layout-qualifier-id-list) uniform;
```

and for shader storage blocks as:

```
layout (layout-qualifier-id-list) buffer;
```

When this is done, the previous default qualification is first inherited and then overridden as per the override rules listed below for each qualifier listed in the declaration. The result becomes the new default qualification scoped to subsequent uniform or shader storage block definitions.

The initial state of compilation is as if the following were declared:

```
layout(shared, column_major) uniform;
layout(shared, column_major) buffer;
```

Uniform and shader storage blocks can be declared with optional layout qualifiers, and so can their individual member declarations. Such block layout qualification is scoped only to the content of the block. As with global layout declarations, block layout qualification first inherits from the current default qualification and then overrides it. Similarly, individual member layout qualification is scoped just to the member declaration, and inherits from and overrides the block's qualification.

The *shared* qualifier overrides only the *std140*, *std430* and *packed* qualifiers; other qualifiers are inherited. The compiler/linker will ensure that multiple programs and programmable stages containing this definition will share the same memory layout for this block, as long as they also matched in their *row_major* and/or *column_major* qualifications. This allows use of the same buffer to back the same block definition across different programs.

The *packed* qualifier overrides only *std140*, *std430* and *shared*; other qualifiers are inherited. When *packed* is used, no shareable layout is guaranteed. The compiler and linker can optimize memory use based on what variables actively get used and on other criteria. Offsets must be queried, as there is no other way of guaranteeing where (and which) variables reside within the block.

It is a link-time error to access the same packed uniform or shader storage block in multiple stages within a program. Attempts to access the same packed uniform or shader storage block across programs can result in conflicting member offsets and in undefined values being read. However, implementations may aid application management of packed blocks by using canonical layouts for packed blocks.

The *std140* and *std430* qualifiers override only the *packed*, *shared*, *std140* and *std430* qualifiers; other qualifiers are inherited. The *std430* qualifier is supported only for shader storage blocks; a shader using the *std430* qualifier on a uniform block will fail to compile.

The layout is explicitly determined by this, as described in section 7.6.2.2 “Standard Uniform Block Layout” of the OpenGL ES Graphics System Specification. Hence, as in *shared* above, the resulting layout is shareable across programs.

Layout qualifiers on member declarations cannot use the *shared*, *packed*, *std140*, or *std430* qualifiers. These can only be used at global scope or on a block declaration.

The *row_major* and *column_major* qualifiers affect the layout of only matrices, including all matrices contained in structures and arrays they are applied to, to all depths of nesting. These qualifiers can be applied to other types, but will have no effect.

The *row_major* qualifier overrides only the *column_major* qualifier; other qualifiers are inherited. Elements within a matrix row will be contiguous in memory.

The *column_major* qualifier overrides only the *row_major* qualifier; other qualifiers are inherited. Elements within a matrix column will be contiguous in memory.

The *binding* qualifier specifies the binding point corresponding to the uniform or shader storage block, which will be used to obtain the values of the member variables of the block. It is a compile-time error to specify the *binding* qualifier for the global scope or for block member declarations. Any uniform or shader storage block declared without a *binding* qualifier is initially assigned to block binding point zero. After a program is linked, the binding points used for uniform (but not shader storage) blocks declared with or without a *binding* qualifier can be updated by the OpenGL ES API.

If the *binding* qualifier is used with a uniform block or shader storage block instantiated as an array, the first element of the array takes the specified block binding and each subsequent element takes the next consecutive binding point.

If the binding point for any uniform or shader storage block instance is less than zero, or greater than or equal to the implementation-dependent corresponding maximum number of buffer bindings, a compile-time error will occur. When the *binding* qualifier is used with a uniform or shader storage block instantiated as an array of size N , all elements of the array from *binding* through *binding* + $N - 1$ must be within this range.

When multiple arguments are listed in a **layout** declaration, the effect will be the same as if they were declared one at a time, in order from left to right, each in turn inheriting from and overriding the result from the previous qualification.

For example

```
layout(row_major, column_major)
```

results in the qualification being *column_major*. Other examples:

```
layout(shared, row_major) uniform; // default is now shared and row_major

layout(std140) uniform Transform { // layout of this block is std140
    mat4 M1;                        // row_major
    layout(column_major) mat4 M2;   // column major
    mat3 N1;                        // row_major
};

uniform T2 { // layout of this block is shared
    ...
};

layout(column_major) uniform T3 { // shared and column_major
    mat4 M3;                        // column_major
    layout(row_major) mat4 m4;     // row major
    mat3 N2;                        // column_major
};
```

4.4.5 Opaque Uniform Layout Qualifiers

Uniform layout qualifiers can be used to bind opaque uniform variables to specific buffers or units.

Samplers can be bound to texture image units, images can be bound to image units, and atomic counters can be bound to buffers.

Sampler, image and atomic counter types take the uniform layout qualifier identifier for binding:

```
layout-qualifier-id:
    binding = integer-constant
```

The identifier binding specifies which unit will be bound. Any uniform sampler, image or atomic counter variable declared without a binding qualifier is initially bound to unit zero. After a program is linked, the unit referenced by a sampler uniform variable declared with or without a binding qualifier can be updated by the OpenGL ES API.

If the binding qualifier is used with an array, the first element of the array takes the specified unit and each subsequent element takes the next consecutive unit.

If the binding is less than zero, or greater than or equal to the implementation-dependent maximum supported number of units, a compile-time error will occur. When the binding qualifier is used with an array of size N , all elements of the array from binding through binding + $N - 1$ must be within this range.

A link-time error will result if two shaders in a program specify different *integer-constant* bindings for the same opaque-uniform name. However, it is not an error to specify a binding on some but not all declarations for the same name, as shown in the examples below.

```
// in one shader...
layout(binding=3) uniform sampler2D s; // s bound to unit 3

// in another shader...
uniform sampler2D s; // okay, s still bound at 3

// in another shader...
layout(binding=4) uniform sampler2D s; // ERROR: contradictory bindings
```

4.4.6 Atomic Counter Layout Qualifiers

Atomic counter layout qualifiers can be used on atomic counter declarations. The atomic counter qualifiers are

layout-qualifier-id:

binding = *integer-constant*

offset = *integer-constant*

For example,

```
layout (binding = 2, offset = 4) uniform atomic_uint a;
```

will establish that the opaque handle to the atomic counter *a* will be bound to atomic counter buffer binding point 2 at an offset of 4 basic machine units into that buffer. The default *offset* for binding point 2 will be post incremented by 4 (the size of an atomic counter).

A subsequent atomic counter declaration will inherit the previous (post incremented) offset. For example, a subsequent declaration of

```
layout (binding = 2) uniform atomic_uint bar;
```

will establish that the atomic counter *bar* has a binding to buffer binding point 2 at an offset of 8 basic machine units into that buffer. The offset for binding point 2 will again be post-incremented by 4 (the size of an atomic counter).

When multiple variables are listed in a layout declaration, the effect will be the same as if they were declared one at a time, in order from left to right.

Binding points are not inherited, only offsets. Each binding point tracks its own current default *offset* for inheritance of subsequent variables using the same *binding*. The initial state of compilation is that all binding points have an *offset* of 0. The *offset* can be set per binding point at global scope (without declaring a variable). For example,

```
layout (binding = 2, offset = 4) uniform atomic_uint;
```

Establishes that the next **atomic_uint** declaration for binding point 2 will inherit *offset* 4 (but does not establish a default *binding*):

```
layout (binding = 2) uniform atomic_uint bar; // offset is 4
layout (offset = 8) uniform atomic_uint bar; // error, no default binding
```

Atomic counters may share the same binding point, but if a binding is shared, their offsets must be either explicitly or implicitly (from inheritance) unique and non overlapping.

Example valid uniform declarations, assuming top of shader:

```
layout (binding=3, offset=4) uniform atomic_uint a; // offset = 4
layout (binding=2) uniform atomic_uint b;           // offset = 0
layout (binding=3) uniform atomic_uint c;           // offset = 8
layout (binding=2) uniform atomic_uint d;           // offset = 4
```

Example of an invalid uniform declaration:

```
layout (offset=4) ... // error, must include binding
layout (binding=1, offset=0) ... a; // okay
layout (binding=2, offset=0) ... b; // okay
layout (binding=1, offset=0) ... c; // error, offsets must not be shared
// between a and c
layout (binding=1, offset=2) ... d; // error, overlaps offset 0 of a
```

It is a compile-time error to bind an atomic counter with a binding value greater than or equal to *gl_MaxAtomicCounterBindings*.

4.4.7 Format Layout Qualifiers

Format layout qualifiers can be used on image variable declarations (those declared with a basic type having “**image**” in its keyword). The format layout qualifier identifiers for image variable declarations are:

```
layout-qualifier-id:
    float-image-format-qualifier
    int-image-format-qualifier
    uint-image-format-qualifier
    binding = integer-constant
```

float-image-format-qualifier:

rgba32f
rgba16f
r32f
rgba8
rgba8_snorm

int-image-format-qualifier:

rgba32i
rgba16i
rgba8i
r32i

uint-image-format-qualifier:

rgba32ui
rgba16ui
rgba8ui
r32ui

A format layout qualifier specifies the image format associated with a declared image variable. Only one format qualifier may be specified for any image variable declaration. For image variables with floating-point component types (**image***), signed integer component types (**iimage***), or unsigned integer component types (**uimage***), the format qualifier used must match the *float-image-format-qualifier*, *int-image-format-qualifier*, or *uint-image-format-qualifier* grammar rules, respectively. It is an error to declare an image variable where the format qualifier does not match the image variable type.

The *binding* qualifier was described in section 4.4.5 “Opaque Uniform Layout Qualifiers”.

Any image variable must specify a format layout qualifier.

4.5 Interpolation Qualifiers

Inputs and outputs that could be interpolated can be further qualified by at most one of the following interpolation qualifiers:

Qualifier	Meaning
smooth	perspective correct interpolation
flat	no interpolation

The presence of and type of interpolation is controlled by the above interpolation qualifiers as well as the auxiliary storage qualifiers **centroid** and **sample**. When no interpolation qualifier is present, smooth interpolation is used. It is a compile-time error to use more than one interpolation qualifier. The auxiliary storage qualifier patch is not used for interpolation; it is a compile-time error to use interpolation qualifiers with **patch**.

A variable qualified as **flat** will not be interpolated. Instead, it will have the same value for every fragment within a primitive. This value will come from a single provoking vertex, as described by the OpenGL ES Graphics System Specification. A variable qualified as **flat** may also be qualified as **centroid** or **sample**, which will mean the same thing as qualifying it only as **flat**.

A variable qualified as **smooth** will be interpolated in a perspective-correct manner over the primitive being rendered. Interpolation in a perspective correct manner is specified in equations 13.4 in the OpenGL ES 3.2 Graphics System Specification, section 13.4.1 “Line Segments” and equation 13.7, section 13.5.1 “Polygon Interpolation”.

When single-sampling, or for fragment shader input variables qualified with neither **centroid** nor **sample**, the value is interpolated to the pixel's center and a single value may be assigned to each sample within the pixel, to the extent permitted by the OpenGL ES Specification.

When multi-sampling, **centroid** and **sample** may be used to control the location and frequency of the sampling of the qualified fragment shader input. If a fragment shader input is qualified with **centroid**, a single value may be assigned to that variable for all samples in the pixel, but that value must be interpolated at a location that lies in both the pixel and in the primitive being rendered, including any of the pixel's samples covered by the primitive. Because the location at which the variable is interpolated may be different in neighboring pixels, and derivatives may be computed by computing differences in neighboring pixels, derivatives of centroid-sampled inputs may be less accurate than those for non-centroid interpolated variables. If a fragment shader input is qualified with **sample**, a separate value must be assigned to that variable for each covered sample in the pixel, and that value must be sampled at the location of the individual sample.

4.6 Parameter Qualifiers

Parameters can have these qualifiers:

Qualifier	Meaning
< none: default >	same as in
in	for function parameters passed into a function
out	for function parameters passed back out of a function, but not initialized for use when passed in
inout	for function parameters passed both into and out of a function

Parameter qualifiers are discussed in more detail in section 6.1.1 “Function Calling Conventions”.

4.7 Precision and Precision Qualifiers

4.7.1 Range and Precision

The precision of highp floating-point variables is defined by the IEEE 754 standard for 32-bit floating-point numbers. This includes support for NaNs (Not a Number) and Infs (positive or negative infinities).

4 Variables and Types

The following rules apply to highp operations: Signed infinities and zeros are generated as dictated by IEEE, but subject to the precisions allowed in the following table. Any subnormal (denormalized) value input into a shader or potentially generated by any operation in a shader can be flushed to 0. The rounding mode cannot be set and is undefined. NaNs are not required to be generated. Support for signaling NaNs is not required and exceptions are never raised. Operations including built-in functions that operate on a NaN are not required to return a NaN as the result. However if NaNs are generated, **isnan()** must return the correct value.

Precisions are expressed in terms of maximum relative error in units of ULP (units in the last place), unless otherwise noted.

For single precision operations, precisions are required as follows:

Operation	Precision
$a + b$, $a - b$, $a * b$	Correctly rounded.
$<$, $<=$, $=$, $>$, $>=$	Correct result.
a / b , $1.0 / b$	2.5 ULP for $ b $ in the range $[2^{-126}, 2^{126}]$.
$a * b + c$	Correctly rounded single operation or sequence of two correctly rounded operations.
fma()	Inherited from $a * b + c$
pow (x, y)	Inherited from exp2 ($x * \log_2(y)$).
exp (x), exp2 (x)	$(3 + 2 * x)$ ULP.
log (), log2 ()	3 ULP outside the range $[0.5, 2.0]$. Absolute error $< 2^{-21}$ inside the range $[0.5, 2.0]$.
sqrt ()	Inherited from $1.0 / \text{inversesqrt}()$.
inversesqrt ()	2 ULP.
explicit conversions between types	Correctly rounded.

The rounding mode is not defined but must not affect the result by more than 1 ULP.

Built-in functions defined in the specification with an equation built from the above operations inherit the above errors. These include, for example, the geometric functions, the common functions, and many of the matrix functions. Built-in functions not listed above and not defined as equations of the above have undefined precision. These include, for example, the trigonometric functions and determinant.

Storage requirements are declared through use of *precision qualifiers*. The precision of operations must preserve the storage precisions of the variables involved.

highp floating point values are stored in IEEE 754 single precision floating point format. Mediump and lowp floating point values have minimum range and precision requirements as detailed below and have maximum range and precision as defined by IEEE 754.

4 Variables and Types

All integral types are assumed to be implemented as integers and so may not be emulated by floating point values. Highp signed integers are represented as two's-complement 32-bit signed integers. Highp unsigned integers are represented as unsigned 32-bit integers. Mediump integers (signed and unsigned) must be represented as an integer with between 16 and 32 bits inclusive. Lowp integers (signed and unsigned) must be represented as an integer with between 9 and 32 bits inclusive.

The required ranges and precisions for precision qualifiers are:

Qualifier	Floating Point Range	Floating Point Magnitude Range	Floating Point Precision	Integer Range	
				Signed	Unsigned
highp	Subset of IEEE-754 $(-2^{128}, 2^{128})$	Subset of IEEE-754 $0.0, [2^{-126}, 2^{128})$	Subset of IEEE 754 relative: 2^{-24}	$[-2^{31}, 2^{31}-1]$	$[0, 2^{32}-1]$
mediump (minimum requirements)	$(-2^{14}, 2^{14})$	$(2^{-14}, 2^{14})$	Relative: 2^{-10}	$[-2^{15}, 2^{15}-1]$	$[0, 2^{16}-1]$
lowp (minimum requirements)	$(-2, 2)$	$(2^{-8}, 2)$	Absolute: 2^{-8}	$[-2^8, 2^8-1]$	$[0, 2^9-1]$

Relative precision is defined as the worst case (i.e. largest) ratio of the smallest step in relation to the value for all non-zero values:

$$Precision_{relative} = \left| \frac{|v_1 - v_2|_{min}}{v_1} \right|_{max} \quad v_1, v_2 \neq 0, v_1 \neq v_2$$

It is therefore twice the maximum rounding error when converting from a real number.

In addition, the range and precision of a mediump floating point value must be the same as or greater than the range and precision of a lowp floating point value. The range and precision of a highp floating point value must be the same as or greater than the range and precision of a mediump floating point value.

The range of a mediump integer value must be the same as or greater than the range of a lowp integer value. The range of a highp integer value must be the same as or greater than the range of a mediump integer value.

Within the above specification, an implementation is allowed to vary the representation of numeric values, both within a shader and between different shaders. If necessary, this variance can be controlled using the invariance qualifier.

The actual ranges and precisions provided by an implementation can be queried through the API. See the OpenGL ES 3.2 specification for details on how to do this.

4.7.2 Conversion between precisions

Within the same type, conversion from a lower to a higher precision must be exact. When converting from a higher precision to a lower precision, if the value is representable by the implementation of the target precision, the conversion must also be exact. If the value is not representable, the behavior is dependent on the type:

- For signed and unsigned integers, the value is truncated; bits in positions not present in the target precision are set to zero. (Positions start at zero and the least significant bit is considered to be position zero for this purpose.)
- For floating point values, the value should either clamp to +INF or -INF, or to the maximum or minimum value that the implementation supports. While this behavior is implementation dependent, it should be consistent for a given implementation.

4.7.3 Precision Qualifiers

Any floating point, integer, opaque type declaration can have the type preceded by one of these precision qualifiers:

Qualifier	Meaning
highp	The variable satisfies the minimum requirements for highp described above. Highp variables have the maximum range and precision available but may cause operations to run more slowly on some implementations.
mediump	The variable satisfies the minimum requirements for mediump described above. Mediump variables may typically be used to store high dynamic range colors and low precision geometry.
lowp	The variable satisfies the minimum requirements for lowp described above. Lowp variables may typically be used to store 8-bit color values.

For example:

```
lowp float color;
out mediump vec2 P;
lowp ivec2 foo(lowp mat3);
highp mat4 m;
```

Literal constants do not have precision qualifiers. Neither do Boolean variables. Neither do constructors.

For this paragraph, “operation” includes operators, built-in functions, and constructors, and “operand” includes function arguments and constructor arguments. The precision used to internally evaluate an operation, and the precision qualification subsequently associated with any resulting intermediate values, must be at least as high as the highest precision qualification of the operands consumed by the operation.

In cases where operands do not have a precision qualifier, the precision qualification will come from the other operands. If no operands have a precision qualifier, then the precision qualifications of the operands of the next consuming operation in the expression will be used. This rule can be applied recursively until a precision qualified operand is found. If necessary, it will also include the precision qualification of l-values for assignments, of the declared variable for initializers, of formal parameters for function call arguments, or of function return types for function return values. If the precision cannot be determined by this method e.g. if an entire expression is composed only of operands with no precision qualifier, and the result is not assigned or passed as an argument, then it is evaluated at the default precision of the type or greater. When this occurs in the fragment shader, the default precision must be defined.

For example, consider the statements:

```
uniform highp float h1;
highp float h2 = 2.3 * 4.7; // operation and result are highp precision
mediump float m;
m = 3.7 * h1 * h2;          // all operations are highp precision
h2 = m * h1;               // operation is highp precision
m = h2 - h1;               // operation is highp precision
h2 = m + m;                // addition and result at mediump precision
void f(highp float p);
f(3.3);                    // 3.3 will be passed in at highp precision
```

Precision qualifiers, as with other qualifiers, do not affect the basic type of the variable. In particular, there are no constructors for precision conversions; constructors only convert types. Similarly, precision qualifiers, as with other qualifiers, do not contribute to function overloading based on parameter types. As discussed in section 6.1.1 (“Function Calling Conventions”), function input and output is done through copies, and therefore qualifiers do not have to match.

Precision qualifiers for outputs in one shader matched to inputs in another shader need not match when both shaders are linked into the same program. When both shaders are in separate programs, mismatched precision qualifiers will result in a program interface mismatch that will result in program pipeline validation failures, as described in section 7.4.1 (“Shader Interface Matching”) of the OpenGL ES 3.2 Specification.

The precision of a variable is determined when the variable is declared and cannot be subsequently changed.

Where the precision of a constant integral or constant floating point expression is not specified, evaluation is performed at highp. This rule does not affect the precision qualification of the expression.

The evaluation of constant expressions must be invariant and will usually be performed at compile time.

4.7.4 Default Precision Qualifiers

The precision statement

```
precision precision-qualifier type;
```

can be used to establish a default precision qualifier. The *type* field can be either **int**, **float** or any of the opaque types and the *precision-qualifier* can be **lowp**, **mediump**, or **highp**. Any other types or qualifiers will result in an error. If *type* is **float**, the directive applies to non-precision-qualified floating point type (scalar, vector, and matrix) declarations. If *type* is **int**, the directive applies to all non-precision-qualified integral type (scalar, vector, signed, and unsigned) declarations. This includes global variable declarations, function return declarations, function parameter declarations, and local variable declarations.

Non-precision qualified declarations will use the precision qualifier specified in the most recent **precision** statement that is still in scope. The **precision** statement has the same scoping rules as variable declarations. If it is declared inside a compound statement, its effect stops at the end of the innermost statement it was declared in. Precision statements in nested scopes override precision statements in outer scopes. Multiple precision statements for the same basic type can appear inside the same scope, with later statements overriding earlier statements within that scope.

All languages except for the fragment language have the following predeclared globally scoped default precision statements:

```
precision highp float;
precision highp int;

precision lowp sampler2D;
precision lowp samplerCube;

precision highp atomic_uint;
```

The fragment language has the following predeclared globally scoped default precision statements:

```
precision mediump int;

precision lowp sampler2D;
precision lowp samplerCube;

precision highp atomic_uint;
```

The fragment language has no default precision qualifier for floating point types. Hence for **float**, floating point vector and matrix variable declarations, either the declaration must include a precision qualifier or the default float precision must have been previously declared. Similarly, there is no default precision qualifier for any of the image types, or any of the following sampler types in any of the languages:

```
sampler3D  
samplerCubeShadow  
sampler2DShadow  
sampler2DArray  
sampler2DArrayShadow  
sampler2DMS  
samplerBuffer  
samplerCubeArray  
samplerCubeArrayShadow  
imageCubeArray
```

```
isampler2D  
isampler3D  
isamplerCube  
isampler2DArray  
isampler2DMS  
isamplerBuffer  
isamplerCubeArray  
iimageCubeArray
```

```
usampler2D  
usampler3D  
usamplerCube  
usampler2DArray  
usampler2DMS  
usamplerBuffer  
usamplerCubeArray  
uimageCubeArray
```

```
image2D  
image3D  
imageCube  
image2DArray  
iimage2D  
iimage3D  
iimageCube  
iimage2DArray  
uimage2D  
uimage3D  
uimageCube  
uimage2DArray  
imageBuffer  
iimageBuffer  
uimageBuffer
```

4.7.5 Available Precision Qualifiers

The built-in macro `GL_FRAGMENT_PRECISION_HIGH` is defined to one in GLSL ES 3.2 to indicate that **high** precision is always in the fragment language

```
#define GL_FRAGMENT_PRECISION_HIGH 1
```

This macro is available in all languages except compute.

4.8 Variance and the Invariant Qualifier

In this section, *variance* refers to the possibility of getting different values from the same expression in different programs. For example, consider the situation where two vertex shaders, in different programs, each set *gl_Position* with the same expression, and the input values into that expression are the same when both shaders run. It is possible, due to independent compilation of the two shaders, that the values assigned to *gl_Position* are not exactly the same when the two shaders run. In this example, this can cause problems with alignment of geometry in a multi-pass algorithm.

In general, such variance between shaders is allowed. When such variance does not exist for a particular output variable, that variable is said to be *invariant*.

4.8.1 The Invariant Qualifier

To ensure that a particular output variable is invariant, it is necessary to use the **invariant** qualifier. It can either be used to qualify a previously declared variable as being invariant

```
invariant gl_Position;    // make built-in gl_Position be invariant

out vec3 Color;
invariant Color;          // make existing Color be invariant

invariant Color_2;        // error: Color_2 has not been declared
```

or as part of a declaration when a variable is declared

```
invariant centroid out vec3 Color;
```

Only variables output from a shader can be candidates for invariance. This includes user-defined output variables and the built-in output variables. As only outputs can be declared as invariant, an invariant output from one shader stage will still match an input of a subsequent stage without the input being declared as invariant.

The **invariant** keyword can be followed by a comma separated list of previously declared identifiers. All uses of **invariant** must be at the global scope, and before any use of the variables being declared as invariant.

To guarantee invariance of a particular output variable across two programs, the following must also be true:

- The output variable is declared as invariant in both programs.
- The same values must be input to all shader input variables consumed by expressions and control flow contributing to the value assigned to the output variable.
- The texture formats, texel values, and texture filtering are set the same way for any texture function calls contributing to the value of the output variable.

- All input values are all operated on in the same way. All operations in the consuming expressions and any intermediate expressions must be the same, with the same order of operands and same associativity, to give the same order of evaluation. Intermediate variables and functions must be declared as the same type with the same explicit or implicit precision qualifiers and the same constant qualifiers. Any control flow affecting the output value must be the same, and any expressions consumed to determine this control flow must also follow these invariance rules.
- All the data flow and control flow leading to setting the invariant output variable reside in a single compilation unit.

Essentially, all the data flow and control flow leading to an invariant output must match.

Initially, by default, all output variables are allowed to be variant. To force all output variables to be invariant, use the pragma

```
#pragma STDGL invariant(all)
```

before all declarations in a shader. If this pragma is used after the declaration of any variables or functions, then the set of outputs that behave as invariant is undefined.

Generally, invariance is ensured at the cost of flexibility in optimization, so performance can be degraded by use of invariance. Hence, use of this pragma is intended as a debug aid, to avoid individually declaring all output variables as invariant.

4.8.2 Invariance Within a Shader

When a value is stored in a variable, it is usually assumed it will remain constant unless explicitly changed. However, during the process of optimization, it is possible that the compiler may choose to recompute a value rather than store it in a register. Since the precision of operations is not completely specified (e.g. a low precision operation may be done at medium or high precision), it would be possible for the recomputed value to be different from the original value.

Values are allowed to be variant within a shader. To prevent this, the invariant qualifier or invariant pragma must be used.

Within a shader, there is no invariance for values generated by different non-constant expressions, even if those expressions are identical.

Example 1:

```
precision mediump;
vec4 col;
vec2 a = ...
...
col = texture(tex, a);    // a has a value a1
...
col = texture(tex, a);    // a has a value a2 where possibly a1 ≠ a2
```

To enforce invariance in this example use:

```
#pragma STDGL invariant(all)
```

Example 2:

```
vec2 m = ...;
vec2 n = ...;
vec2 a = m + n;
vec2 b = m + n;    // a and b are not guaranteed to be exactly equal
```

There is no mechanism to enforce invariance between a and b.

4.8.3 Invariance of Constant Expressions

Invariance must be guaranteed for constant expressions. A particular constant expression must evaluate to the same result if it appears again in the same shader or a different shader. This includes the same expression appearing in two shaders of the same language or shaders of two different languages.

Constant expressions must evaluate to the same result when operated on as already described above for invariant variables. Constant expressions are not invariant with respect to equivalent non-constant expressions, even when the invariant qualifier or pragma is used.

4.8.4 Invariance of Undefined Values

Undefined values are not invariant nor can they be made invariant by use of the invariant qualifier or pragma. In some implementations, undefined values may cause unexpected behavior if they are used in control-flow expressions e.g. in the following case, one, both or neither functions may be executed and this may not be consistent over multiple invocations of the shader:

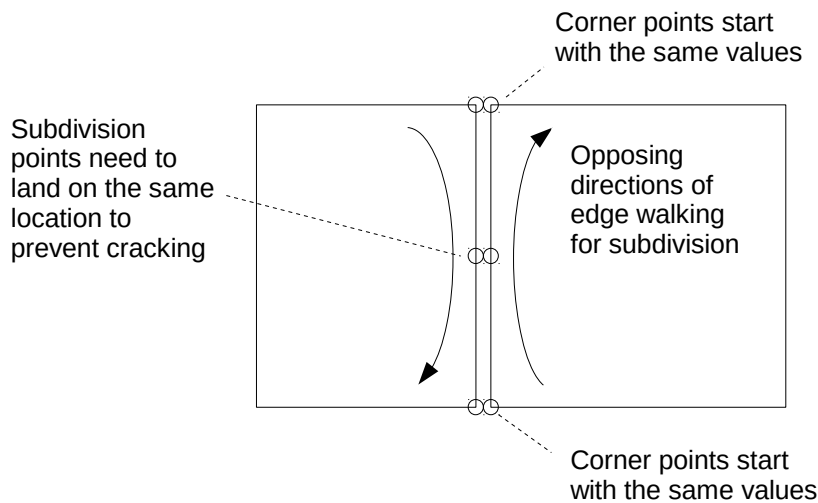
```
int x; // undefined value
if (x == 1)
{
    f(); // Undefined whether f() is executed
}
if (x == 2)
{
    g(); // Undefined whether g() is executed.
}
```

Note that an undefined value is a value that has not been specified. A value that has been specified but has a potentially large error due to, for example, lack of precision in an expression, is not undefined and so can be made invariant.

4.9 The Precise Qualifier

As stated in section 5.11 (“Evaluation of Expressions”), the compiler may transform expressions even if this changes the resulting value. Furthermore, these transforms do not need to be applied consistently so that apparently similar expressions may generate different results. The **invariant** qualifier may be used to guarantee that *identical* expressions produce the same result (see section 4.8) but there are some algorithms which require that non-identical but mathematically equivalent expressions also generate exactly the same result. The **precise** qualifier ensures that operations are (effectively) performed in their stated order and with operator consistency.

The key computation that needs to be made consistent appears when tessellating, where intermediate points for subdivision are synthesized in different directions, yet need to yield the same result, as shown in the diagram below:



The order of evaluation is determined by operator precedence and parenthesis, as described in section 5.1 (“Operators”). Operator consistency means for each particular operator, for example the multiply operator (`*`), its operation is always computed with the same precision. Specifically, values computed by compiler-generated code must adhere to the following identities:

1. $a + b = b + a$
2. $a * b = b * a$
3. $a * b + c * d = b * a + c * d = d * c + b * a = \text{<any other mathematically valid combination>}$

While the following are prevented:

4. $a + (b + c)$ is not allowed to become $(a + b) + c$
5. $a * (b * c)$ is not allowed to become $(a * b) * c$
6. $a * b + c$ is not allowed to become a single operation `fma(a, b, c)`

4 Variables and Types

Where a, b, c, and d, are scalars or vectors, not matrices. (Matrix multiplication generally does not commute.) It is the shader writer's responsibility to express the computation in terms of these rules and the compiler's responsibility to follow these rules. See the description of `gl_TessCoord` for the rules the tessellation stages are responsible for following, which in conjunction with the above allow avoiding cracking when subdividing.

For example,

```
precise out vec4 position;
```

declares that operations used to produce the value of `position` must be performed in exactly the order specified in the source code and with all operators being treated consistently. As with the invariant qualifier (section 4.8.1 “The Invariant Qualifier”), the `precise` qualifier may be used to qualify a built-in or previously declared user-defined variable as being `precise`:

```
out vec3 Color;
precise Color;           // make existing Color be precise
```

This qualifier will affect the evaluation of an r-value in a particular function if and only if the result is eventually consumed in the same function by an l-value qualified as `precise`. Any other expressions within a function are not affected, including return values and output parameters not declared as `precise` but that are eventually consumed outside the function by a variable qualified as `precise`.

Some examples of the use of `precise`:

```
in vec4 a, b, c, d;
precise out vec4 v;

float func(float e, float f, float g, float h)
{
    return (e*f) + (g*h);           // no constraint on order or
                                    // operator consistency
}
```

```
float func2(float e, float f, float g, float h)
{
    precise float result = (e*f) + (g*h); // ensures same precision for
                                           // the two multiplies
    return result;
}

float func3(float i, float j, precise out float k)
{
    k = i * i + j; // precise, due to <k> declaration
}

void main()
{
    vec3 r = vec3(a * b); // precise, used to compute v.xyz
    vec3 s = vec3(c * d); // precise, used to compute v.xyz
    v.xyz = r + s; // precise
    v.w = (a.w * b.w) + (c.w * d.w); // precise
    v.x = func(a.x, b.x, c.x, d.x); // values computed in func()
                                   // are NOT precise
    v.x = func2(a.x, b.x, c.x, d.x); // precise!
    func3(a.x * b.x, c.x * d.x, v.x); // precise!
}
```

For the purposes of determining if an output from one shader stage matches an input of the next stage, the precise qualifier need not match between the input and the output.

All constant expressions are evaluated as if precise was present, whether or not it is present. However, as described in section 4.3.3 (“Constant Expressions”), there is no requirement that a compile-time constant expression evaluates to the same value as a corresponding non-constant expression.

4.10 Memory Access Qualifiers

Shader storage blocks, variables declared within shader storage blocks and variables declared as image types (the basic opaque types with “**image**” in their keyword), can be further qualified with one or more of the following memory qualifiers:

Qualifier	Meaning
coherent	memory variable where reads and writes are coherent with reads and writes from other shader invocations
volatile	memory variable whose underlying value may be changed at any point during shader execution by some source other than the current shader invocation
restrict	memory variable where use of that variable is the only way to read and write the underlying memory in the relevant shader stage
readonly	memory variable that can be used to read the underlying memory, but cannot be used to write the underlying memory
writeonly	memory variable that can be used to write the underlying memory, but cannot be used to read the underlying memory

Memory accesses to image variables declared using the **coherent** qualifier are performed coherently with accesses to the same location from other shader invocations.

As described in section 7.11.1 “Shader Memory Access Ordering” of the OpenGL ES Specification, shader memory reads and writes complete in a largely undefined order. The built-in function **memoryBarrier()** can be used if needed to guarantee the completion and relative ordering of memory accesses performed by a single shader invocation.

When accessing memory using variables not declared as **coherent**, the memory accessed by a shader may be cached by the implementation to service future accesses to the same address. Memory stores may be cached in such a way that the values written may not be visible to other shader invocations accessing the same memory. The implementation may cache the values fetched by memory reads and return the same values to any shader invocation accessing the same memory, even if the underlying memory has been modified since the first memory read. While variables not declared as **coherent** may not be useful for communicating between shader invocations, using non-coherent accesses may result in higher performance.

Memory accesses to image variables declared using the **volatile** storage qualifier must treat the underlying memory as though it could be read or written at any point during shader execution by some source other than the executing shader invocation. When a volatile variable is read, its value must be re-fetched from the underlying memory, even if the shader invocation performing the read had previously fetched its value from the same memory. When a volatile variable is written, its value must be written to the underlying memory, even if the compiler can conclusively determine that its value will be overwritten by a subsequent write. Since the external source reading or writing a **volatile** variable may be another shader invocation, variables declared as **volatile** are automatically treated as coherent.

Memory accesses to image variables declared using the **restrict** storage qualifier may be compiled assuming that the variable used to perform the memory access is the only way to access the underlying memory using the shader stage in question. This allows the compiler to coalesce or reorder loads and stores using **restrict**-qualified image variables in ways that wouldn't be permitted for image variables not so qualified, because the compiler can assume that the underlying image won't be read or written by other code. Applications are responsible for ensuring that image memory referenced by variables qualified with **restrict** will not be referenced using other variables in the same scope; otherwise, accesses to **restrict**-qualified variables will have undefined results.

Memory accesses to image variables declared using the **readonly** qualifier may only read the underlying memory, which is treated as read-only memory and cannot be written to. It is an error to pass an image variable qualified with **readonly** to **imageStore()** or other built-in functions that modify image memory.

Memory accesses to image variables declared using the **writeonly** qualifier may only write the underlying memory; the underlying memory cannot be read. It is an error to pass an image variable qualified with **writeonly** to **imageLoad()** or other built-in functions that read image memory. A variable could be qualified as both **readonly** and **writeonly**, disallowing both read and write, but still be passed to **imageSize()** to have the size queried.

Except for image variables qualified with the format qualifiers **r32f**, **r32i**, and **r32ui**, image variables must specify either memory qualifier **readonly** or the memory qualifier **writeonly**.

The memory qualifiers **coherent**, **volatile**, **restrict**, **readonly**, and **writeonly** may be used in the declaration of buffer variables (i.e., members of shader storage blocks). When a buffer variable is declared with a memory qualifier, the behavior specified for memory accesses involving image variables described above applies identically to memory accesses involving that buffer variable. It is a compile-time error to assign to a buffer variable qualified with **readonly** or to read from a buffer variable qualified with **writeonly**.

Additionally, memory qualifiers may also be used in the declaration of shader storage blocks. When a block declaration is qualified with a memory qualifier, it is as if all of its members were declared with the same memory qualifier. For example, the block declaration

```
coherent buffer Block {
    readonly vec4 member1;
    vec4 member2;
};
```

is equivalent to

```
buffer Block {
    coherent readonly vec4 member1;
    coherent vec4 member2;
};
```

Memory qualifiers are only supported in the declarations of image variables, buffer variables, and shader storage blocks; it is an error to use such qualifiers in any other declarations.

When calling user-defined functions, variables qualified with **coherent**, **volatile**, **readonly**, or **writeonly** may not be passed to functions whose formal parameters lack such qualifiers. (See section 6.1 “Function Definitions” for more detail on function calling.) It is legal to have any additional memory qualifiers on a formal parameter, but only **restrict** can be taken away from a calling argument, by a formal parameter that lacks the restrict qualifier.

When a built-in function is called, the code generated is to be based on the actual qualification of the calling argument, not on the list of memory qualifiers specified on the formal parameter in the prototype. For example, if a calling argument is not coherent but a formal parameter of a built-in function is coherent, the code generated for the built-in function will be for the original non-coherent memory qualification.

```
vec4 funcA(layout(rgba32f) image2D restrict a) { ... }
vec4 funcB(layout(rgba32f) image2D a)          { ... }
layout(rgba32f) uniform image2D img1;
layout(rgba32f) coherent uniform image2D img2;

funcA(img1);           // OK, adding "restrict" is allowed
funcB(img2);           // illegal, stripping "coherent" is not allowed
```

Layout qualifiers cannot be used on formal function parameters, but they are not included in parameter matching.

Note that the use of **const** in an image variable declaration is qualifying the const-ness of variable being declared, not the image it refers to: The qualifier **readonly** qualifies the image memory (as accessed through that variable) while **const** qualifies the variable itself.

4.11 Order of Qualification

When multiple qualifiers are present in a declaration, they may appear in any order, but they must all appear before the type. The **layout** qualifier is the only qualifier that can appear more than once. Further, a declaration can have at most one storage qualifier, at most one auxiliary storage qualifier, and at most one interpolation qualifier. Multiple memory qualifiers can be used. Any violation of these rules will cause a compile-time error.

4.12 Empty Declarations

Empty declarations are allowed. E.g.

```
int;           // No effect
struct S {int x;}; // Defines a struct S
```

The combinations of qualifiers that cause compile-time or link-time errors are the same whether or not the declaration is empty e.g.

```
invariant in float x;      // Error. An input cannot be invariant.
invariant in float;        // Error even though no variable is declared.
```


5 Operators and Expressions

5.1 Operators

The OpenGL ES Shading Language has the following operators.

Precedence	Operator Class	Operators	Associativity
1 (highest)	parenthetical grouping	()	NA
2	array subscript function call and constructor structure field or method selector, swizzler post fix increment and decrement	[] () . ++ --	Left to Right
3	prefix increment and decrement unary	++ -- + - ~ !	Right to Left
4	multiplicative	* / %	Left to Right
5	additive	+ -	Left to Right
6	bit-wise shift	<< >>	Left to Right
7	relational	< > <= >=	Left to Right
8	equality	== !=	Left to Right
9	bit-wise and	&	Left to Right
10	bit-wise exclusive or	^	Left to Right
11	bit-wise inclusive or		Left to Right
12	logical and	&&	Left to Right
13	logical exclusive or	^^	Left to Right
14	logical inclusive or		Left to Right
15	selection	? :	Right to Left
16	Assignment arithmetic assignments	= += -= *= /= %= <<= >>= &= ^= =	Right to Left
17 (lowest)	sequence	,	Left to Right

There is no address-of operator nor a dereference operator. There is no typecast operator; constructors are used instead.

5.2 Array Operations

These are now described in section 5.7 (“Structure and Array Operations”).

5.3 Function Calls

If a function returns a value, then a call to that function may be used as an expression, whose type will be the type that was used to declare or define the function.

Function definitions and calling conventions are discussed in section 6.1 (“Function Definitions”).

5.4 Constructors

Constructors use the function call syntax, where the function name is a type, and the call makes an object of that type. Constructors are used the same way in both initializers and expressions. (See section 9 “Shading Language Grammar” for details.) The parameters are used to initialize the constructed value. Constructors can be used to request a data type conversion to change from one scalar type to another scalar type, or to build larger types out of smaller types, or to reduce a larger type to a smaller type.

In general, constructors are not built-in functions with predetermined prototypes. For arrays and structures, there must be exactly one argument in the constructor for each element or field. For the other types, the arguments must provide a sufficient number of components to perform the initialization, and it is an error to include so many arguments that they cannot all be used. Detailed rules follow. The prototypes actually listed below are merely a subset of examples.

5.4.1 Conversion and Scalar Constructors

Converting between scalar types is done as the following prototypes indicate:

```
int(bool)      // converts a Boolean value to an int
int(float)     // converts a float value to an int
float(bool)    // converts a Boolean value to a float
float(int)     // converts a signed integer value to a float
bool(float)    // converts a float value to a Boolean
bool(int)      // converts a signed integer value to a Boolean
uint(bool)     // converts a Boolean value to an unsigned integer
uint(float)    // converts a float value to an unsigned integer
uint(int)      // converts a signed integer value to an unsigned integer
int(uint)      // converts an unsigned integer to a signed integer
bool(uint)     // converts an unsigned integer value to a Boolean value
float(uint)    // converts an unsigned integer value to a float value
```

When constructors are used to convert a **float** to an **int** or **uint**, the fractional part of the floating-point value is dropped. It is undefined to convert a negative floating point value to an **uint**.

When a constructor is used to convert an **int**, **uint**, or a **float** to a **bool**, 0 and 0.0 are converted to **false**, and non-zero values are converted to **true**. When a constructor is used to convert a **bool** to an **int**, **uint**, or **float**, **false** is converted to 0 or 0.0, and **true** is converted to 1 or 1.0.

The constructor **int(uint)** preserves the bit pattern in the argument, which will change the argument's value if its sign bit is set. The constructor **uint(int)** preserves the bit pattern in the argument, which will change its value if it is negative.

Identity constructors, like **float(float)** are also legal, but of little use.

Scalar constructors with non-scalar parameters can be used to take the first element from a non-scalar. For example, the constructor **float(vec3)** will select the first component of the **vec3** parameter.

5.4.2 Vector and Matrix Constructors

Constructors can be used to create vectors or matrices from a set of scalars, vectors, or matrices. This includes the ability to shorten vectors.

If there is a single scalar parameter to a vector constructor, it is used to initialize all components of the constructed vector to that scalar's value. If there is a single scalar parameter to a matrix constructor, it is used to initialize all the components on the matrix's diagonal, with the remaining components initialized to 0.0.

If a vector is constructed from multiple scalars, one or more vectors, or one or more matrices, or a mixture of these, the vector's components will be constructed in order from the components of the arguments. The arguments will be consumed left to right, and each argument will have all its components consumed, in order, before any components from the next argument are consumed. Similarly for constructing a matrix from multiple scalars or vectors, or a mixture of these. Matrix components will be constructed and consumed in column major order. In these cases, there must be enough components provided in the arguments to provide an initializer for every component in the constructed value. It is an error to provide extra arguments beyond this last used argument.

If a matrix is constructed from a matrix, then each component (column i , row j) in the result that has a corresponding component (column i , row j) in the argument will be initialized from there. All other components will be initialized to the identity matrix. If a matrix argument is given to a matrix constructor, it is an error to have any other arguments.

If the basic type (**bool**, **int**, or **float**) of a parameter to a constructor does not match the basic type of the object being constructed, the scalar construction rules (above) are used to convert the parameters.

5 Operators and Expressions

Some useful vector constructors are as follows:

```
vec3(float)    // initializes each component of the vec3 with the float
vec4(ivec4)    // makes a vec4 with component-wise conversion
vec4(mat2)     // the vec4 is column 0 followed by column 1

vec2(float, float)           // initializes a vec2 with 2 floats
ivec3(int, int, int)         // initializes an ivec3 with 3 ints
bvec4(int, int, float, float) // uses 4 Boolean conversions

vec2(vec3)           // drops the third component of a vec3
vec3(vec4)           // drops the fourth component of a vec4

vec3(vec2, float)    // vec3.x = vec2.x, vec3.y = vec2.y, vec3.z = float
vec3(float, vec2)    // vec3.x = float, vec3.y = vec2.x, vec3.z = vec2.y
vec4(vec3, float)
vec4(float, vec3)
vec4(vec2, vec2)
```

Some examples of these are:

```
vec4 color = vec4(0.0, 1.0, 0.0, 1.0);
vec4 rgba  = vec4(1.0); // sets each component to 1.0
vec3 rgb   = vec3(color); // drop the 4th component
```

To initialize the diagonal of a matrix with all other elements set to zero:

```
mat2(float)
mat3(float)
mat4(float)
```

That is, $result[i][j]$ is set to the float argument for all $i=j$ and set to 0 for all $i \neq j$.

5 Operators and Expressions

To initialize a matrix by specifying vectors or scalars, the components are assigned to the matrix elements in column-major order.

```
mat2(vec2, vec2);           // one column per argument
mat3(vec3, vec3, vec3);     // one column per argument
mat4(vec4, vec4, vec4, vec4); // one column per argument
mat3x2(vec2, vec2, vec2);   // one column per argument

mat2(float, float,         // first column
      float, float);      // second column

mat3(float, float, float,   // first column
      float, float, float,   // second column
      float, float, float); // third column

mat4(float, float, float, float, // first column
      float, float, float, float, // second column
      float, float, float, float, // third column
      float, float, float, float); // fourth column

mat2x3(vec2, float,         // first column
        vec2, float);      // second column
```

A wide range of other possibilities exist, to construct a matrix from vectors and scalars, as long as enough components are present to initialize the matrix. To construct a matrix from a matrix:

```
mat3x3(mat4x4); // takes the upper-left 3x3 of the mat4x4
mat2x3(mat4x2); // takes the upper-left 2x2 of the mat4x2, last row is 0,0
mat4x4(mat3x3); // puts the mat3x3 in the upper-left, sets the lower right
                // component to 1, and the rest to 0
```

5.4.3 Structure Constructors

Once a structure is defined, and its type is given a name, a constructor is available with the same name to construct instances of that structure. For example:

```
struct light {
    float intensity;
    vec3 position;
};

light lightVar = light(3.0, vec3(1.0, 2.0, 3.0));
```

The arguments to the constructor will be used to set the structure's fields, in order, using one argument per field. Each argument must be the same type as the field it sets.

Structure constructors can be used as initializers or in expressions.

5.4.4 Array Constructors

Array types can also be used as constructor names, which can then be used in expressions or initializers. For example,

```
const float c[3] = float[3](5.0, 7.2, 1.1);
const float d[3] = float[](5.0, 7.2, 1.1);

float g;
...
float a[5] = float[5](g, 1, g, 2.3, g);
float b[3];

b = float[3](g, g + 1.0, g + 2.0);

vec4 a[][] = vec4[][](vec4[2](vec4(0.0), vec4(1.0)),
                      vec4[2](vec4(0.0), vec4(1.0)),
                      vec4[2](vec4(0.0), vec4(1.0)));

float a[][] = float[][](float[] (1.0, 2.0), float[] (3.0, 4.0));
float m[2][1];
m = float[][](float[] (1.0), float[] (2.0));
```

There must be exactly the same number of arguments as the size of the array being constructed. The arguments are assigned in order, starting at element 0, to the elements of the constructed array. Each argument must be the same type as the element type of the array.

5.5 Vector Components

The names of the components of a vector are denoted by a single letter. As a notational convenience, several letters are associated with each component based on common usage of position, color or texture coordinate vectors. The individual components of a vector can be selected by following the variable name with period (.) and then the component name.

The component names supported are:

$\{x, y, z, w\}$	Useful when accessing vectors that represent points or normals
$\{r, g, b, a\}$	Useful when accessing vectors that represent colors
$\{s, t, p, q\}$	Useful when accessing vectors that represent texture coordinates

The component names x , r , and s are, for example, synonyms for the same (first) component in a vector.

Note that the third component of the texture coordinate set, r in OpenGL ES, has been renamed p so as to avoid the confusion with r (for red) in a color.

Accessing components beyond those declared for the vector type is an error so, for example:

5 Operators and Expressions

```
vec2 pos;  
pos.x // is legal  
pos.z // is illegal
```

The component selection syntax allows multiple components to be selected by appending their names (from the same name set) after the period (.).

```
vec4 v4;  
v4.rgba; // is a vec4 and the same as just using v4,  
v4.rgb;  // is a vec3,  
v4.b;    // is a float,  
v4.xy;    // is a vec2,  
v4.xgba; // is illegal - the component names do not come from  
          // the same set.
```

No more than 4 components can be selected.

```
vec4 v4;  
v4.xyzw; // is a vec4  
v4.xyzwxy; // is illegal since it has 6 components  
(v4.xyzwxy).xy; // is illegal since the intermediate value has 6 components  
  
vec2 v2;  
v2.xxyy; // is legal. It evaluates to a vec4.
```

The order of the components can be different to swizzle them, or replicated:

```
vec4 pos = vec4(1.0, 2.0, 3.0, 4.0);  
vec4 swiz = pos.wzyx; // swiz = (4.0, 3.0, 2.0, 1.0)  
vec4 dup = pos.xxyy; // dup = (1.0, 1.0, 2.0, 2.0)
```

This notation is more concise than the constructor syntax. To form an r-value, it can be applied to any expression that results in a vector r-value.

The component group notation can occur on the left hand side of an expression.

```
vec4 pos = vec4(1.0, 2.0, 3.0, 4.0);  
pos.xw = vec2(5.0, 6.0); // pos = (5.0, 2.0, 3.0, 6.0)  
pos.wx = vec2(7.0, 8.0); // pos = (8.0, 2.0, 3.0, 7.0)  
pos.xx = vec2(3.0, 4.0); // illegal - 'x' used twice  
pos.xy = vec3(1.0, 2.0, 3.0); // illegal - mismatch between vec2 and vec3
```

To form an l-value, swizzling must be applied to an l-value of vector type, contain no duplicate components, and it results in an l-value of scalar or vector type, depending on number of components specified.

Array subscripting syntax can also be applied to vectors to provide numeric indexing. So in

```
vec4 pos;
```

`pos[2]` refers to the third element of `pos` and is equivalent to `pos.z`. This allows variable indexing into a vector, as well as a generic way of accessing components. Any integer expression can be used as the subscript. The first component is at index zero. Reading from or writing to a vector using a constant integral expression with a value that is negative or greater than or equal to the size of the vector is illegal. When indexing with non-constant expressions, behavior is undefined if the index is negative, or greater than or equal to the size of the vector.

Note that scalars are not considered to be single-component vectors and therefore the use of component selection operators on scalars is illegal.

5.6 Matrix Components

The components of a matrix can be accessed using array subscripting syntax. Applying a single subscript to a matrix treats the matrix as an array of column vectors, and selects a single column, whose type is a vector of the same size as the (column size of the) matrix. The leftmost column is column 0. A second subscript would then operate on the resulting vector, as defined earlier for vectors. Hence, two subscripts select a column and then a row.

```
mat4 m;
m[1] = vec4(2.0);           // sets the second column to all 2.0
m[0][0] = 1.0;              // sets the upper left element to 1.0
m[2][3] = 2.0;              // sets the 4th element of the third column to 2.0
```

Behavior is undefined when accessing a component outside the bounds of a matrix with a non-constant expression. It is an error to access a matrix with a constant expression that is outside the bounds of the matrix.

5.7 Structure and Array Operations

The fields of a structure and the **length** method of an array are selected using the period (`.`).

In total, only the following operators are allowed to operate on arrays and structures as whole entities:

field or method selector	<code>.</code>
equality	<code>== !=</code>
assignment	<code>=</code>
indexing (arrays only)	<code>[]</code>

The equality operators and assignment operator are only allowed if the two operands are same size and type. Array types must be compile-time sized. Structure types must be of the same declared structure. When using the equality operators, two structures are equal if and only if all the fields are component-wise equal, and two arrays are equal if and only if all the elements are element-wise equal.

Array elements are accessed using the array subscript operator (`[]`). An example of accessing an array element is

```
diffuseColor += lightIntensity[3] * NdotL;
```

Array indices start at zero. Array elements are accessed using an expression whose type is **int** or **uint**.

Arrays can also be accessed with the method operator (`.`) and the **length** method to query the size of the array:

```
lightIntensity.length()    // return the size of the array
```

5.8 Assignments

Assignments of values to variable names are done with the assignment operator (`=`):

```
lvalue-expression = rvalue-expression
```

The *lvalue-expression* evaluates to an l-value. The assignment operator stores the value of *rvalue-expression* into the l-value and returns an r-value with the type and precision of *lvalue-expression*. The *lvalue-expression* and *rvalue-expression* must have the same type. Any type-conversions must be specified explicitly via constructors. L-values must be writable. Variables that are built-in types, entire structures or arrays, structure fields, l-values with the field selector (`.`) applied to select components or swizzles without repeated fields, l-values within parentheses, and l-values dereferenced with the array subscript operator (`[]`) are all l-values. Other binary or unary expressions, function names, swizzles with repeated fields, and constants cannot be l-values. The ternary operator (`?:`) is also not allowed as an l-value.

Expressions on the left of an assignment are evaluated before expressions on the right of the assignment.

The other assignment operators are

- add into (`+=`)
- subtract from (`-=`)
- multiply into (`*=`)
- divide into (`/=`)
- modulus into (`%=`)
- left shift by (`<<=`)
- right shift by (`>>=`)
- and into (`&=`)
- inclusive-or into (`|=`)
- exclusive-or into (`^=`)

where the general expression

```
lvalue-expression op= expression
```

is equivalent to

```
lvalue = lvalue op expression
```

where *lvalue* is the value returned by *lvalue-expression*, *op* is as described below, and the *lvalue-expression* and *expression* must satisfy the semantic requirements of both *op* and equals (=).

Reading a variable before writing (or initializing) it is legal, however the value is undefined.

5.9 Expressions

Expressions in the shading language are built from the following:

- Constants of type **bool**, **int**, **uint**, **float**, all vector types, and all matrix types.
- Constructors of all types.
- Variable names of all types.
- An array name with the length method applied.
- Subscripted arrays.
- Function calls that return values.
- Component field selectors and array subscript results.
- Parenthesized expression. Any expression can be parenthesized. Parentheses can be used to group operations. Operations within parentheses are done before operations across parentheses.
- The arithmetic binary operators add (+), subtract (-), multiply (*), and divide (/) operate on integer and floating-point scalars, vectors, and matrices. If the operands are integral types, they must both be signed or both be unsigned. All arithmetic binary operators result in the same fundamental type (signed integer, unsigned integer, or floating-point) as the operands they operate on. The following cases are valid
 - The two operands are scalars. In this case the operation is applied, resulting in a scalar.
 - One operand is a scalar, and the other is a vector or matrix. In this case, the scalar operation is applied independently to each component of the vector or matrix, resulting in the same size vector or matrix.
 - The two operands are vectors of the same size. In this case, the operation is done component-wise resulting in the same size vector.
 - The operator is add (+), subtract (-), or divide (/), and the operands are matrices with the same number of rows and the same number of columns. In this case, the operation is done component-wise resulting in the same size matrix.

- The operator is multiply (*), where both operands are matrices or one operand is a vector and the other a matrix. A right vector operand is treated as a column vector and a left vector operand as a row vector. In all these cases, it is required that the number of columns of the left operand is equal to the number of rows of the right operand. Then, the multiply (*) operation does a linear algebraic multiply, yielding an object that has the same number of rows as the left operand and the same number of columns as the right operand. Section 5.10 (“Vector and Matrix Operations”) explains in more detail how vectors and matrices are operated on.

All other cases are illegal.

Dividing by zero does not cause an exception but does result in an unspecified value. Use the built-in functions **dot**, **cross**, **matrixCompMult**, and **outerProduct**, to get, respectively, vector dot product, vector cross product, matrix component-wise multiplication, and the matrix product of a column vector times a row vector.

- The operator modulus (%) operates on signed or unsigned integers or integer vectors. The operand types must both be signed or both be unsigned. The operands cannot be vectors of differing size. If one operand is a scalar and the other vector, then the scalar is applied component-wise to the vector, resulting in the same type as the vector. If both are vectors of the same size, the result is computed component-wise. The resulting value is undefined for any component computed with a second operand that is zero, while results for other components with non-zero second operands remain defined. If both operands are non-negative, then the remainder is non-negative. Results are undefined if one or both operands are negative. The operator modulus (%) is not defined for any other data types (non-integral types).
- The arithmetic unary operators negate (-), post- and pre-increment and decrement (-- and ++), operate on integer or floating-point values (including vectors and matrices). All unary operators work component-wise on their operands. These result with the same type they operated on. For post- and pre-increment and decrement, the expression must be one that could be assigned to (an l-value). Pre-increment and pre-decrement add or subtract 1 or 1.0 to the contents of the expression they operate on, and the value of the pre-increment or pre-decrement expression is the resulting value of that modification. Post-increment and post-decrement expressions add or subtract 1 or 1.0 to the contents of the expression they operate on, but the resulting expression has the expression’s value before the post-increment or post-decrement was executed.
- The relational operators greater than (>), less than (<), greater than or equal (>=), and less than or equal (<=) operate only on scalar integer and scalar floating-point expressions. The result is scalar Boolean. The types of the operands must match. To do component-wise relational comparisons on vectors, use the built-in functions **lessThan**, **lessThanEqual**, **greaterThan**, and **greaterThanEqual**.
- The equality operators **equal** (==), and not equal (!=) operate on all types except opaque types. They result in a scalar Boolean. The types of the operands must match. For vectors, matrices, structures, and arrays, all components, fields, or elements of one operand must equal the corresponding components, fields, or elements in the other operand for the operands to be considered equal. To get a vector of component-wise equality results for vectors, use the built-in functions **equal** and **notEqual**.

5 Operators and Expressions

- The logical binary operators `and (&&)`, `or (||)`, and `exclusive or (^)` operate only on two Boolean expressions and result in a Boolean expression. `And (&&)` will only evaluate the right hand operand if the left hand operand evaluated to **true**. `Or (||)` will only evaluate the right hand operand if the left hand operand evaluated to **false**. `Exclusive or (^)` will always evaluate both operands.
- The logical unary operator `not (!)`. It operates only on a Boolean expression and results in a Boolean expression. To operate on a vector, use the built-in function **not**.
- The sequence `(,)` operator that operates on expressions by returning the type and value of the right-most expression in a comma separated list of expressions. All expressions are evaluated, in order, from left to right.
- The ternary selection operator `(?:)`. It operates on three expressions (*exp1 ? exp2 : exp3*). This operator evaluates the first expression, which must result in a scalar Boolean. If the result is true, it selects to evaluate the second expression, otherwise it selects to evaluate the third expression. Only one of the second and third expressions is evaluated. The second and third expressions can be any type, as long their types match. This resulting matching type is the type of the entire expression.
- The one's complement operator `(~)`. The operand must be of type signed or unsigned integer or integer vector, and the result is the one's complement of its operand; each bit of each component is complemented, including any sign bits.
- The shift operators `(<<)` and `(>>)`. For both operators, the operands must be signed or unsigned integers or integer vectors. One operand can be signed while the other is unsigned. In all cases, the resulting type will be the same type as the left operand. If the first operand is a scalar, the second operand has to be a scalar as well. If the first operand is a vector, the second operand must be a scalar or a vector with the same size as the first operand, and the result is computed component-wise. The result is undefined if the right operand is negative, or greater than or equal to the number of bits in the left expression's base type. The value of *E1 << E2* is *E1* (interpreted as a bit pattern) left-shifted by *E2* bits. The value of *E1 >> E2* is *E1* right-shifted by *E2* bit positions. If *E1* is a signed integer, the right-shift will extend the sign bit. If *E1* is an unsigned integer, the right-shift will zero-extend.
- The bitwise operators `and (&)`, `exclusive-or (^)`, and `inclusive-or (|)`. The operands must be of type signed or unsigned integers or integer vectors. The operands cannot be vectors of differing size. If one operand is a scalar and the other a vector, the scalar is applied component-wise to the vector, resulting in the same type as the vector. The fundamental types of the operands (signed or unsigned) must match, and will be the resulting fundamental type. For `and (&)`, the result is the bitwise-and function of the operands. For `exclusive-or (^)`, the result is the bitwise exclusive-or function of the operands. For `inclusive-or (|)`, the result is the bitwise inclusive-or function of the operands.

For a complete specification of the syntax of expressions, see section 9 “Shading Language Grammar”.

5.10 Vector and Matrix Operations

With a few exceptions, operations are component-wise. Usually, when an operator operates on a vector or matrix, it is operating independently on each component of the vector or matrix, in a component-wise fashion. For example,

```
vec3 v, u;
float f;

v = u + f;
```

will be equivalent to

```
v.x = u.x + f;
v.y = u.y + f;
v.z = u.z + f;
```

And

```
vec3 v, u, w;
w = v + u;
```

will be equivalent to

```
w.x = v.x + u.x;
w.y = v.y + u.y;
w.z = v.z + u.z;
```

and likewise for most operators and all integer and floating point vector and matrix types. The exceptions are matrix multiplied by vector, vector multiplied by matrix, and matrix multiplied by matrix. These do not operate component-wise, but rather perform the correct linear algebraic multiply.

```
vec3 v, u;
mat3 m;

u = v * m;
```

is equivalent to

```
u.x = dot(v, m[0]); // m[0] is the left column of m
u.y = dot(v, m[1]); // dot(a,b) is the inner (dot) product of a and b
u.z = dot(v, m[2]);
```

And

```
u = m * v;
```

is equivalent to

```
u.x = m[0].x * v.x + m[1].x * v.y + m[2].x * v.z;
u.y = m[0].y * v.x + m[1].y * v.y + m[2].y * v.z;
u.z = m[0].z * v.x + m[1].z * v.y + m[2].z * v.z;
```

And

```
mat3 m, n, r;
```

```
r = m * n;
```

is equivalent to

```
r[0].x = m[0].x * n[0].x + m[1].x * n[0].y + m[2].x * n[0].z;
r[1].x = m[0].x * n[1].x + m[1].x * n[1].y + m[2].x * n[1].z;
r[2].x = m[0].x * n[2].x + m[1].x * n[2].y + m[2].x * n[2].z;
```

```
r[0].y = m[0].y * n[0].x + m[1].y * n[0].y + m[2].y * n[0].z;
r[1].y = m[0].y * n[1].x + m[1].y * n[1].y + m[2].y * n[1].z;
r[2].y = m[0].y * n[2].x + m[1].y * n[2].y + m[2].y * n[2].z;
```

```
r[0].z = m[0].z * n[0].x + m[1].z * n[0].y + m[2].z * n[0].z;
r[1].z = m[0].z * n[1].x + m[1].z * n[1].y + m[2].z * n[1].z;
r[2].z = m[0].z * n[2].x + m[1].z * n[2].y + m[2].z * n[2].z;
```

and similarly for other sizes of vectors and matrices.

5.11 Evaluation of Expressions

The C++ standard requires that expressions must be evaluated in the order specified by the precedence of operations and may only be regrouped if the result is the same or where the result is undefined. No other transforms may be applied that affect the result of an operation. GLSL ES relaxes these requirements for scalar operations in the following ways:

- Addition and multiplication are assumed to be associative.
- Multiplication is assumed to be distributive over addition. Therefore expressions may be expanded and re-factored.
- Floating point division may be replaced by reciprocal and multiplication.
- Multiplication may be replaced by repeated addition.
- Within the constraints of invariance (where applicable), the precision used may vary.

These rules also apply to the built-in functions.

6 Statements and Structure

The fundamental building blocks of the OpenGL ES Shading Language are:

- statements and declarations
- function definitions
- selection (**if-else** and **switch-case-default**)
- iteration (**for**, **while**, and **do-while**)
- jumps (**discard**, **return**, **break**, and **continue**)

The overall structure of a shader is as follows

translation-unit:
 global-declaration
 translation-unit global-declaration

global-declaration:
 function-definition
 declaration

That is, a shader is a sequence of declarations and function bodies. Function bodies are defined as

function-definition:
 function-prototype { statement-list }

statement-list:
 statement
 statement-list statement

statement:
 compound-statement
 simple-statement

Curly braces are used to group sequences of statements into compound statements.

compound-statement:
 { statement-list }

simple-statement:
declaration-statement
expression-statement
selection-statement
iteration-statement
jump-statement

Simple declaration, expression, and jump statements end in a semi-colon.

This above is slightly simplified, and the complete grammar specified in section 9 “Shading Language Grammar” should be used as the definitive specification.

Declarations and expressions have already been discussed.

6.1 Function Definitions

As indicated by the grammar above, a valid shader is a sequence of global declarations and function definitions. A function is declared as the following example shows:

```
// prototype
returnType functionName (type0 arg0, type1 arg1, ..., typen argn);
```

and a function is defined like

```
// definition
returnType functionName (type0 arg0, type1 arg1, ..., typen argn)
{
    // do some computation
    return returnValue;
}
```

where *returnType* must be present and cannot be void, or:

```
void functionName (type0 arg0, type1 arg1, ..., typen argn)
{
    // do some computation
    return; // optional
}
```

Each of the *typeN* must include a type and can optionally include a parameter qualifier and/or **const**.

A function is called by using its name followed by a list of arguments in parentheses.

Arrays are allowed as arguments and as the return type. In both cases, the array must be compile-time sized. An array is passed or returned by using just its name, without brackets, and the size of the array must match the size specified in the function's declaration.

Structures are also allowed as argument types. The return type can also be a structure.

See section 9 “Shading Language Grammar” for the definitive reference on the syntax to declare and define functions.

All functions must be either declared with a prototype or defined with a body before they are called. For example:

6 Statements and Structure

```
float myfunc (float f,          // f is an input parameter
             out float g);    // g is an output parameter
```

Functions that return no value must be declared as **void**. A void function can only use return without a return argument, even if the return argument has void type. Return statements only accept values:

```
void func1() { }
void func2() { return func1(); } // illegal return statement
```

Only a precision qualifier is allowed on the return type of a function. Formal parameters can have parameter, precision and memory qualifiers, but no other qualifiers.

Functions that accept no input arguments need not use **void** in the argument list because prototypes (or definitions) are required and therefore there is no ambiguity when an empty argument list “()” is declared. The idiom “(void)” as a parameter list is provided for convenience.

Function names can be overloaded. The same function name can be used for multiple functions, as long as the parameter types differ. If a function name is declared twice with the same parameter types, then the return types and all qualifiers must also match, and it is the same function being declared. When function calls are resolved, an exact type match for all the arguments is required.

For example,

```
vec4 f(in vec4 x, out vec4 y);
vec4 f(in vec4 x, out ivec4 y); // allowed, different argument type
int  f(in vec4 x, out ivec4 y); // error, only return type differs
vec4 f(in vec4 x, in ivec4 y); // error, only qualifier differs
int  f(const in vec4 x, out ivec4 y); // error, only qualifier differs
```

Calling the first two functions above with the following argument types yields

```
f(vec4, vec4)    // exact match of vec4 f(in vec4 x, out vec4 y)
f(vec4, ivec4)   // exact match of vec4 f(in vec4 x, out ivec4 y)
f(ivec4, vec4)   // error, no exact match.
f(ivec4, ivec4)  // error, no exact match.
```

User-defined functions can have multiple declarations, but only one definition.

A shader cannot redefine or overload built-in functions.

The function *main* is used as the entry point to a shader executable. All shaders must define a function named *main*. This function takes no arguments, returns no value, and must be declared as type **void**:

```
void main()
{
    ...
}
```

The function *main* can contain uses of **return**. See section 6.4 (“Jumps”) for more details.

It is an error to declare or define a function **main** with any other parameters or return type.

6.1.1 Function Calling Conventions

Functions are called by value-return. This means input arguments are copied into the function at call time, and output arguments are copied back to the caller before function exit. Because the function works with local copies of parameters, there are no issues regarding aliasing of variables within a function. To control what parameters are copied in and/or out through a function definition or declaration:

- The keyword **in** is used as a qualifier to denote a parameter is to be copied in, but not copied out.
- The keyword **out** is used as a qualifier to denote a parameter is to be copied out, but not copied in. This should be used whenever possible to avoid unnecessarily copying parameters in.
- The keyword **inout** is used as a qualifier to denote the parameter is to be both copied in and copied out.
- A function parameter declared with no such qualifier means the same thing as specifying **in**.

All arguments are evaluated at call time, exactly once, in order, from left to right. Evaluation of an **in** parameter results in a value that is copied to the formal parameter. Evaluation of an **out** parameter results in an l-value that is used to copy out a value when the function returns. Evaluation of an **inout** parameter results in both a value and an l-value; the value is copied to the formal parameter at call time and the l-value is used to copy out a value when the function returns.

The order in which output parameters are copied back to the caller is undefined.

In a function, writing to an input-only parameter is allowed. Only the function's copy is modified. This can be prevented by declaring a parameter with the **const** qualifier.

When calling a function, expressions that do not evaluate to l-values cannot be passed to parameters declared as **out** or **inout**.

The return type can be any basic type, array type, structure name, or *structure definition*. Only precision qualifiers are allowed on the return type of a function.

The syntax for function prototypes can be informally expressed as:

function_prototype:

[*type_qualifier*] *type_specifier IDENTIFIER*
LEFT_PAREN parameter_declaration, parameter_declaration, ... , RIGHT_PAREN

parameter_declaration:

[*type_qualifier*] *type_specifier* [*IDENTIFIER* [*array_specifier*]]

type_qualifier:

single_type_qualifier, single_type_qualifier, ...

The qualifiers allowed on formal parameters are:

empty

in

out**inout****const***memory-qualifier**precision-qualifier*

However, the **const** qualifier cannot be used with **out** or **inout**. The above is used for function declarations (i.e., prototypes) and for function definitions. Hence, function definitions can have unnamed arguments.

Static, and hence dynamic recursion, are not allowed. Static recursion is present if the static function call graph of the program contains cycles. Dynamic recursion occurs if at any time control flow has entered but not exited a single function more than once.

6.2 Selection

Conditional control flow in the shading language is done by either **if**, **if-else**, or **switch** statements:

selection-statement:

```
if ( bool-expression ) statement
if ( bool-expression ) statement else statement
switch ( init-expression ) { switch-statement-listopt }
```

Where *switch-statement-list* is a nested scope containing a list of zero or more *switch-statement* and other statements defined by the language, where *switch-statement* adds some forms of labels. That is

switch-statement-list:

```
switch-statement
switch-statement-list switch-statement
```

switch-statement:

```
case constant-expression :
default :
statement
```

If an **if**-expression evaluates to **true**, then the first *statement* is executed. If it evaluates to **false** and there is an **else** part then the second *statement* is executed.

Any expression whose type evaluates to a Boolean can be used as the conditional expression *bool-expression*. Vector types are not accepted as the expression to **if**.

Conditionals can be nested.

The type of *init-expression* in a **switch** statement must be a scalar integer. If a **case** label has a *constant-expression* of equal value, then execution will continue after that label. Otherwise, if there is a **default** label, execution will continue after that label. Otherwise, execution skips the rest of the switch statement. It is an error to have more than one **default** or a replicated *constant-expression*. A **break** statement not nested in a loop or other switch statement (either not nested or nested only in **if** or **if-else** statements) will also skip the rest of the switch statement. Fall through labels are allowed. No statements are allowed in a switch statement before the first **case** statement.

The type of *init-expression* must match the type of the **case** labels within each **switch** statement. Either signed integers or unsigned integers are allowed but there is no implicit type conversion between the two.

No **case** or **default** labels can be nested inside other statements or compound statements within their corresponding **switch**.

6.3 Iteration

For, while, and do loops are allowed as follows:

```
for (init-expression; condition-expression; loop-expression)
    sub-statement

while (condition-expression)
    sub-statement

do
    statement
while (condition-expression)
```

See section 9 (“Shading Language Grammar”) for the definitive specification of loops.

The **for** loop first evaluates the *init-expression*, then the *condition-expression*. If the *condition-expression* evaluates to **true**, then the body of the loop is executed. After the body is executed, a **for** loop will then evaluate the *loop-expression*, and then loop back to evaluate the *condition-expression*, repeating until the *condition-expression* evaluates to **false**. The loop is then exited, skipping its body and skipping its *loop-expression*. Variables modified by the *loop-expression* maintain their value after the loop is exited, provided they are still in scope. Variables declared in *init-expression* or *condition-expression* are only in scope until the end of the sub-statement of the **for** loop.

The **while** loop first evaluates the *condition-expression*. If **true**, then the body is executed. This is then repeated, until the *condition-expression* evaluates to **false**, exiting the loop and skipping its body. Variables declared in the *condition-expression* are only in scope until the end of the sub-statement of the while loop.

For both **for** and **while** loops, the sub-statement does not introduce a new scope for variable names, so the following has a redeclaration error:

```
for (int i = 0; i < 10; i++)
{
    int i; // redeclaration error
}
```

The **do-while** loop first executes the body, then executes the *condition-expression*. This is repeated until *condition-expression* evaluates to **false**, and then the loop is exited.

Expressions for *condition-expression* must evaluate to a Boolean.

Both the *condition-expression* and the *init-expression* can declare and initialize a variable, except in the **do-while** loop, which cannot declare a variable in its *condition-expression*. The variable's scope lasts only until the end of the sub-statement that forms the body of the loop.

Loops can be nested.

Non-terminating loops are allowed. The consequences of very long or non-terminating loops are platform dependent.

6.4 Jumps

These are the jumps:

```
jump_statement:
    continue;
    break;
    return;
    return expression;
    discard; // in the fragment shader language only
```

There is no “goto” or other non-structured flow of control.

The **continue** jump is used only in loops. It skips the remainder of the body of the innermost loop of which it is inside. For **while** and **do-while** loops, this jump is to the next evaluation of the loop *condition-expression* from which the loop continues as previously defined. For **for** loops, the jump is to the *loop-expression*, followed by the *condition-expression*.

The **break** jump can also be used only in loops and switch statements. It is simply an immediate exit of the inner-most loop or switch statements containing the **break**. No further execution of *condition-expression*, *loop-expression*, or *switch-statement* is done.

The **discard** keyword is only allowed within fragment shaders. It can be used within a fragment shader to abandon the operation on the current fragment. This keyword causes the fragment to be discarded and no updates to the framebuffer will occur. Any prior writes to other buffers such as shader storage buffers are unaffected. Control flow exits the shader, and subsequent implicit or explicit derivatives are undefined when this control flow is non-uniform (meaning different fragments within the primitive take different control paths). It would typically be used within a conditional statement, for example:

```
if (intensity < 0.0)
    discard;
```

A fragment shader may test a fragment's alpha value and discard the fragment based on that test. However, it should be noted that coverage testing occurs after the fragment shader runs, and the coverage test can change the alpha value.

The **return** jump causes immediate exit of the current function. If it has *expression* then that is the return value for the function.

The function *main* can use **return**. This simply causes *main* to exit in the same way as when the end of the function had been reached. It does not imply a use of **discard** in a fragment shader. Using **return** in *main* before defining outputs will have the same behavior as reaching the end of *main* before defining outputs.

7 Built-in Variables

7.1 Built-in Language Variables

Some OpenGL ES operations occur in fixed functionality and need to provide values to or receive values from shader executables. Shaders communicate with fixed-function OpenGL ES pipeline stages, and optionally with other shader executables, through the use of built-in input and output variables.

7.1.1 Vertex Shader Special Variables

The built-in vertex shader variables are intrinsically declared as follows:

```
in highp int gl_VertexID;
in highp int gl_InstanceID;

out gl_PerVertex
{
    out highp vec4   gl_Position;
    out highp float  gl_PointSize;
};
```

Unless otherwise noted elsewhere, these variables are only available in the vertex language as declared above.

The variable *gl_Position* is intended for writing the homogeneous vertex position. It can be written at any time during shader execution. This value will be used by primitive assembly, clipping, culling, and other fixed functionality operations, if present, that operate on primitives after vertex processing has occurred. Its value is undefined after the vertex processing stage if the vertex shader executable does not write *gl_Position*.

The variable *gl_PointSize* is intended for a shader to write the size of the point to be rasterized. It is measured in pixels. If *gl_PointSize* is not written to, its value is undefined in subsequent pipe stages.

The variable *gl_VertexID* is a vertex shader input variable that holds an integer index for the vertex, as defined under “Shader Inputs” in section 11.1.3 “Shader Execution” in the OpenGL ES Graphics System Specification. While the variable *gl_VertexID* is always present, its value is not always defined.

The variable *gl_InstanceID* is a vertex shader input variable that holds the instance number of the current primitive in an instanced draw call (see “Shader Inputs” in section 11.1.3 “Shader Execution” in the OpenGL ES 3.2 Graphics System Specification). If the current primitive does not come from an instanced draw call, the value of *gl_InstanceID* is zero.

7.1.2 Tessellation Control Shader Special Variables

In the tessellation control shader, built-in variables are intrinsically declared as:

```
in gl_PerVertex
{
    highp vec4 gl_Position;
} gl_in[gl_MaxPatchVertices];

in highp int gl_PatchVerticesIn;
in highp int gl_PrimitiveID;
in highp int gl_InvocationID;

out gl_PerVertex
{
    highp vec4 gl_Position;
} gl_out[];

patch out highp float gl_TessLevelOuter[4];
patch out highp float gl_TessLevelInner[2];
patch out highp vec4 gl_BoundingBox[2];
```

7.1.2.1 Tessellation Control Input Variables

gl_Position contains the output written in the previous shader stage to *gl_Position*.

gl_PatchVerticesIn contains the number of vertices in the input patch being processed by the shader. A single shader can read patches of differing sizes, so the value of *gl_PatchVerticesIn* may differ between patches.

gl_PrimitiveID contains the number of primitives processed by the shader since the current set of rendering primitives was started.

gl_InvocationID contains the number of the output patch vertex assigned to the tessellation control shader invocation. It is assigned integer values in the range [0, N-1], where N is the number of output patch vertices per primitive.

7.1.2.2 Tessellation Control Output Variables

gl_Position is used in the same fashion as the corresponding output variable in the vertex shader.

The values written to *gl_TessLevelOuter* and *gl_TessLevelInner* are assigned to the corresponding outer and inner tessellation levels of the output patch. They are used by the tessellation primitive generator to control primitive tessellation, and may be read by tessellation evaluation shaders.

7.1.3 Tessellation Evaluation Shader Special Variables

In the tessellation evaluation language, built-in variables are intrinsically declared as:


```

in gl_PerVertex
{
    highp vec4 gl_Position;
} gl_in[gl_MaxPatchVertices];

in highp int gl_PatchVerticesIn;
in highp int gl_PrimitiveID;
in highp vec3 gl_TessCoord;
patch in highp float gl_TessLevelOuter[4];
patch in highp float gl_TessLevelInner[2];

out gl_PerVertex
{
    highp vec4 gl_Position;
};

```

7.1.3.1 Tessellation Evaluation Input Variables

gl_Position contains the output written in the previous shader stage to *gl_Position*.

gl_PatchVerticesIn and *gl_PrimitiveID* are defined in the same fashion as the corresponding input variables in the tessellation control shader.

gl_TessCoord specifies a three-component (u,v,w) vector identifying the position of the vertex being processed by the shader relative to the primitive being tessellated. Its values will obey the properties

```

gl_TessCoord.x == 1.0 - (1.0 - gl_TessCoord.x) // two operations performed
gl_TessCoord.y == 1.0 - (1.0 - gl_TessCoord.y) // two operations performed
gl_TessCoord.z == 1.0 - (1.0 - gl_TessCoord.z) // two operations performed

```

gl_TessLevelOuter and *gl_TessLevelInner* are filled with the corresponding output variables written by the active tessellation control shader.

7.1.3.2 Tessellation Evaluation Output Variables

gl_Position is used in the same fashion as the corresponding output variable in the vertex shader.

7.1.4 Geometry Shader Special Variables

In the geometry language, the built-in variables are intrinsically declared as:

```

in gl_PerVertex
{
    highp vec4 gl_Position;
} gl_in[];

in highp int gl_PrimitiveIDIn;
in highp int gl_InvocationID;

out gl_PerVertex
{
    highp vec4 gl_Position;
};

out highp int gl_PrimitiveID;
out highp int gl_Layer;

```

7.1.4.1 Geometry Shader Input Variables

gl_Position contains the output written in the previous shader stage to *gl_Position*.

gl_PrimitiveIDIn contains the number of primitives processed by the shader since the current set of rendering primitives was started.

gl_InvocationID contains the invocation number assigned to the geometry shader invocation. It is assigned integer values in the range [0, N-1], where N is the number of geometry shader invocations per primitive.

7.1.4.2 Geometry Shader Output Variables

gl_Position is used in the same fashion as the corresponding output variable in the vertex shader. Its value is undefined after geometry processing if the shader calls **EmitVertex()** without having written *gl_Position* since the last **EmitVertex()**, or does not write it at all.

gl_PrimitiveID is filled with a single integer that serves as a primitive identifier to the fragment shader. This is then available to fragment shaders, which will select the written primitive ID from the provoking vertex of the primitive being shaded. If a fragment shader using *gl_PrimitiveID* is active and a geometry shader is also active, the geometry shader must write to *gl_PrimitiveID* or the fragment shader input *gl_PrimitiveID* is undefined. See section 11.3.4.4 “Geometry Shader Outputs” of the OpenGL ES Specification for more information.

gl_Layer is used to select a specific layer (or face and layer of a cube map) of a multi-layer framebuffer attachment. The actual layer used will come from one of the vertices in the primitive being shaded. Which vertex the layer comes from is determined as discussed in section 11.3.4.4 of the OpenGL ES Specification but may be undefined, so it is best to write the same layer value for all vertices of a primitive. If a shader statically assigns a value to *gl_Layer*, layered rendering mode is enabled. See section 11.3.4.4 “Geometry Shader Outputs” and section 9.8 “Layered Framebuffers” of the OpenGL ES Specification for more information. If a shader statically assigns a value to *gl_Layer*, and there is an execution path through the shader that does not set *gl_Layer*, then the value of *gl_Layer* is undefined for executions of the shader that take that path.

`gl_Layer` takes on a special value when used with an array of cube map textures. Instead of only referring to the layer, it is used to select a cube map face and a layer. Setting `gl_Layer` to the value $(\text{layer} * 6 + \text{face})$ will render to the face *face* of the cube defined in layer *layer*. The face values are defined in table 8.25 of the OpenGL ES Specification.

For example, to render to the positive *y* cube map face located in the 5th layer of the cube map array, `gl_Layer` should be set to $5 * 6 + 2$.

Face Value	Resulting Target
0	TEXTURE_CUBE_MAP_POSITIVE_X
1	TEXTURE_CUBE_MAP_NEGATIVE_X
2	TEXTURE_CUBE_MAP_POSITIVE_Y
3	TEXTURE_CUBE_MAP_NEGATIVE_Y
4	TEXTURE_CUBE_MAP_POSITIVE_Z
5	TEXTURE_CUBE_MAP_NEGATIVE_Z

7.1.5 Fragment Shader Special Variables

The built-in special variables that are accessible from a fragment shader are intrinsically declared as follows:

```

in  highp   vec4    gl_FragCoord;
in           bool    gl_FrontFacing;
out highp   float    gl_FragDepth;
in  mediump vec2    gl_PointCoord;
in           bool    gl_HelperInvocation;
in  highp   int      gl_PrimitiveID;
in  highp   int      gl_Layer;
in  lowp    int      gl_SampleID;
in  mediump vec2    gl_SamplePosition;
in  highp   int      gl_SampleMaskIn[(gl_MaxSamples+31)/32];
out highp   int      gl_SampleMask[(gl_MaxSamples+31)/32];

```

Except as noted below, they behave as other input and output variables.

The output of the fragment shader executable is processed by the fixed function operations at the back end of the OpenGL ES pipeline.

The fixed functionality computed depth for a fragment may be obtained by reading `gl_FragCoord.z`, described below.

Writing to *gl_FragDepth* will establish the depth value for the fragment being processed. If depth buffering is enabled, and no shader writes *gl_FragDepth*, then the fixed function value for depth will be used as the fragment's depth value. If a shader statically assigns a value to *gl_FragDepth*, and there is an execution path through the shader that does not set *gl_FragDepth*, then the value of the fragment's depth may be undefined for executions of the shader that take that path. That is, if the set of linked fragment shaders statically contain a write to *gl_FragDepth*, then it is responsible for always writing it.

If a shader executes the **discard** keyword, the fragment is discarded, and the values of any user-defined fragment outputs, become irrelevant.

The variable *gl_FragCoord* is available as an input variable from within fragment shaders and it holds the window relative coordinates (*x*, *y*, *z*, *1/w*) values for the fragment. If multi-sampling, this value can be for any location within the pixel, or one of the fragment samples. The use of **centroid** does not further restrict this value to be inside the current primitive. This value is the result of the fixed functionality that interpolates primitives after vertex processing to generate fragments. The *z* component is the depth value that would be used for the fragment's depth if no shader contained any writes to *gl_FragDepth*. This is useful for invariance if a shader conditionally computes *gl_FragDepth* but otherwise wants the fixed functionality fragment depth.

Fragment shaders have access to the input built-in variable *gl_FrontFacing*, whose value is **true** if the fragment belongs to a front-facing primitive. One use of this is to emulate two-sided lighting by selecting one of two colors calculated by a vertex or geometry shader.

The values in *gl_PointCoord* are two-dimensional coordinates indicating where within a point primitive the current fragment is located, when point sprites are enabled. They range from 0.0 to 1.0 across the point. If the current primitive is not a point, or if point sprites are not enabled, then the values read from *gl_PointCoord* are undefined.

The value *gl_HelperInvocation* is **true** if the fragment shader invocation is considered a "helper" invocation and is **false** otherwise. A helper invocation is a fragment shader invocation that is created solely for the purposes of evaluating derivatives for the built-in functions *texture()* (section 8.9 "Texture Functions"), *dFdx()*, *dFdy()*, and *fwidth()* for other non-helper fragment shader invocations.

Fragment shader helper invocations execute the same shader code as non-helper invocations, but will not have side effects that modify the framebuffer or other shader-accessible memory. In particular:

- Fragments corresponding to helper invocations are discarded when shader execution is complete, without updating the framebuffer.
- Stores to image and buffer variables performed by helper invocations have no effect on the underlying image or buffer memory.
- Atomic operations to image, buffer, or atomic counter variables performed by helper invocations have no effect on the underlying image or buffer memory. The values returned by such atomic operations are undefined.

Helper invocations may be generated for pixels not covered by a primitive being rendered. While fragment shader inputs qualified with **centroid** are normally required to be sampled in the intersection of the pixel and the primitive, the requirement is ignored for such pixels since there is no intersection between the pixel and primitive.

Helper invocations may also be generated for fragments that are covered by a primitive being rendered when the fragment is killed by early fragment tests (using the **early_fragment_tests** qualifier) or where the implementation is able to determine that executing the fragment shader would have no effect other than assisting in computing derivatives for other fragment shader invocations.

The set of helper invocations generated when processing any set of primitives is implementation-dependent.

The input variable *gl_PrimitiveID* is filled with the value written to the *gl_PrimitiveID* geometry shader output, if a geometry shader is present. Otherwise, it is filled with the number of primitives processed by the shader since the current set of rendering primitives was started.

The input variable *gl_Layer* is filled with the value written to the *gl_Layer* geometry shader output, if a geometry shader is present. If the geometry stage does not dynamically assign a value to *gl_Layer*, the value of *gl_Layer* in the fragment stage will be undefined. If the geometry stage makes no static assignment to *gl_Layer*, the input value in the fragment stage will be zero. Otherwise, the fragment stage will read the same value written by the geometry stage, even if that value is out of range. If a fragment shader contains a static access to *gl_Layer*, it will count against the implementation defined limit for the maximum number of inputs to the fragment stage.

The input variable *gl_SampleID* is filled with the sample number of the sample currently being processed. This variable is in the range 0 to *gl_NumSamples*-1, where *gl_NumSamples* is the total number of samples in the framebuffer, or one if rendering to a non-multisample framebuffer. Any static use of *gl_SampleID* in a fragment shader causes the entire shader to be executed per-sample.

The input variable *gl_SamplePosition* contains the position of the current sample within the multi-sample draw buffer. The x and y components of *gl_SamplePosition* contain the sub-pixel coordinate of the current sample and will have values in the range 0.0 to 1.0. Any static use of this variable in a fragment shader causes the entire shader to be executed per-sample.

For the both the input array *gl_SampleMaskIn[]* and the output array *gl_SampleMask[]*, bit B of mask M (*gl_SampleMaskIn[M]* or *gl_SampleMask[M]*) corresponds to sample $32 \cdot M + B$. These arrays have $\text{ceil}(\text{gl_MaxSamples}/32)$ elements, where *gl_MaxSamples* is the maximum number of color samples supported by the implementation.

The input variable *gl_SampleMaskIn* indicates the set of samples covered by the primitive generating the fragment during multisample rasterization. It has a sample bit set if and only if the sample is considered covered for this fragment shader invocation. The output array *gl_SampleMask[]* sets the sample mask for the fragment being processed. Coverage for the current fragment will be the logical AND of the coverage mask and the output *gl_SampleMask*. If the fragment shader statically assigns a value to *gl_SampleMask*, the sample mask will be undefined for any array elements of any fragment shader invocations that fails to assign a value. If a shader does not statically assign a value to *gl_SampleMask*, the sample mask has no effect on the processing of a fragment.

7.1.6 Compute Shader Special Variables

In the compute language, the built-in variables are declared as follows:

```
// work group dimensions
in    uvec3 gl_NumWorkGroups;
const uvec3 gl_WorkGroupSize;

// work group and invocation IDs
in    uvec3 gl_WorkGroupID;
in    uvec3 gl_LocalInvocationID;

// derived variables
in    uvec3 gl_GlobalInvocationID;
in    uint  gl_LocalInvocationIndex;
```

The built-in variable *gl_NumWorkGroups* is a compute-shader input variable containing the total number of global work items in each dimension of the work group that will execute the compute shader. Its content is equal to the values specified in the *num_groups_x*, *num_groups_y*, and *num_groups_z* parameters passed to the *DispatchCompute* API entry point.

The built-in constant *gl_WorkGroupSize* is a compute-shader constant containing the local work-group size of the shader. The size of the work group in the *X*, *Y*, and *Z* dimensions is stored in the *x*, *y*, and *z* components. The constants values in *gl_WorkGroupSize* will match those specified in the required **local_size_x**, **local_size_y**, and **local_size_z** layout qualifiers for the current shader. This is a constant so that it can be used to size arrays of memory that can be shared within the local work group. It is a compile-time error to use *gl_WorkGroupSize* in a shader that does not declare a fixed local group size, or before that shader has declared a fixed local group size, using **local_size_x**, **local_size_y**, and **local_size_z**. When a size is given for some of these identifiers, but not all, the corresponding *gl_WorkGroupSize* will have a size of 1.

The built-in variable *gl_WorkGroupID* is a compute-shader input variable containing the three-dimensional index of the global work group that the current invocation is executing in. The possible values range across the parameters passed into *DispatchCompute*, i.e., from (0, 0, 0) to (*gl_NumWorkGroups.x* - 1, *gl_NumWorkGroups.y* - 1, *gl_NumWorkGroups.z* - 1).

The built-in variable *gl_LocalInvocationID* is a compute-shader input variable containing the three-dimensional index of the local work group within the global work group that the current invocation is executing in. The possible values for this variable range across the local work group size, i.e., (0,0,0) to (*gl_WorkGroupSize.x* - 1, *gl_WorkGroupSize.y* - 1, *gl_WorkGroupSize.z* - 1).

The built-in variable *gl_GlobalInvocationID* is a compute shader input variable containing the global index of the current work item. This value uniquely identifies this invocation from all other invocations across all local and global work groups initiated by the current *DispatchCompute* call. This is computed as:

```
gl_GlobalInvocationID =
    gl_WorkGroupID * gl_WorkGroupSize + gl_LocalInvocationID;
```

The built-in variable *gl_LocalInvocationIndex* is a compute shader input variable that contains the one-dimensional representation of the *gl_LocalInvocationID*. This is useful for uniquely identifying a unique region of shared memory within the local work group for this invocation to use. This is computed as:

```
gl_LocalInvocationIndex =  
    gl_LocalInvocationID.z * gl_WorkGroupSize.x * gl_WorkGroupSize.y +  
    gl_LocalInvocationID.y * gl_WorkGroupSize.x +  
    gl_LocalInvocationID.x;
```

7.2 Built-In Constants

The following built-in constants are provided to all shaders. The actual values used are implementation dependent, but must be at least the value shown.

```
//  
// Implementation dependent constants. The example values below  
// are the minimum values allowed for these maximums.  
//  
  
const mediump int gl_MaxVertexAttribs = 16;  
const mediump int gl_MaxVertexUniformVectors = 256;
```

7 Built-in Variables

```
const mediump int gl_MaxVertexOutputVectors = 16;

const mediump int gl_MaxTessControlInputComponents = 64;
const mediump int gl_MaxTessControlOutputComponents = 64;
const mediump int gl_MaxTessControlTextureImageUnits = 16;
const mediump int gl_MaxTessControlUniformComponents = 1024;
const mediump int gl_MaxTessControlTotalOutputComponents = 4096;

const mediump int gl_MaxTessControlImageUniforms = 0;
const mediump int gl_MaxTessEvaluationImageUniforms = 0;
const mediump int gl_MaxTessControlAtomicCounters = 0;
const mediump int gl_MaxTessEvaluationAtomicCounters = 0;
const mediump int gl_MaxTessControlAtomicCounterBuffers = 0;
const mediump int gl_MaxTessEvaluationAtomicCounterBuffers = 0;

const mediump int gl_MaxTessEvaluationInputComponents = 64;
const mediump int gl_MaxTessEvaluationOutputComponents = 64;
const mediump int gl_MaxTessEvaluationTextureImageUnits = 16;
const mediump int gl_MaxTessEvaluationUniformComponents = 1024;

const mediump int gl_MaxTessPatchComponents = 120;

const mediump int gl_MaxPatchVertices = 32;
const mediump int gl_MaxTessGenLevel = 64;

const mediump int gl_MaxGeometryInputComponents = 64;
const mediump int gl_MaxGeometryOutputComponents = 64;
const mediump int gl_MaxGeometryImageUniforms = 0;
const mediump int gl_MaxGeometryTextureImageUnits = 16;
const mediump int gl_MaxGeometryOutputVertices = 256;
const mediump int gl_MaxGeometryTotalOutputComponents = 1024;
const mediump int gl_MaxGeometryUniformComponents = 1024;
const mediump int gl_MaxGeometryAtomicCounters = 0;
const mediump int gl_MaxGeometryAtomicCounterBuffers = 0;

const mediump int gl_MaxFragmentInputVectors = 15;
const mediump int gl_MaxFragmentUniformVectors = 224;
const mediump int gl_MaxDrawBuffers = 4;

const mediump int gl_MaxVertexTextureImageUnits = 16;
const mediump int gl_MaxCombinedTextureImageUnits = 96;
const mediump int gl_MaxTextureImageUnits = 16;
```



```
const mediump int gl_MinProgramTexelOffset = -8;
const mediump int gl_MaxProgramTexelOffset = 7;

const mediump int gl_MaxImageUnits = 4;
const mediump int gl_MaxVertexImageUniforms = 0;
const mediump int gl_MaxFragmentImageUniforms = 0;

const mediump int gl_MaxSamples = 4;

const mediump int gl_MaxComputeImageUniforms = 4;
const mediump int gl_MaxCombinedImageUniforms = 4;

const mediump int gl_MaxCombinedShaderOutputResources = 4;

const highp ivec3 gl_MaxComputeWorkGroupCount = ivec3(65535, 65535, 65535);
const highp ivec3 gl_MaxComputeWorkGroupSize = ivec3(128, 128, 64);
const mediump int gl_MaxComputeUniformComponents = 512;
const mediump int gl_MaxComputeTextureImageUnits = 16;

const mediump int gl_MaxComputeAtomicCounters = 8;
const mediump int gl_MaxComputeAtomicCounterBuffers = 1;

const mediump int gl_MaxVertexAtomicCounters = 0;
const mediump int gl_MaxFragmentAtomicCounters = 0;
const mediump int gl_MaxCombinedAtomicCounters = 8;
const mediump int gl_MaxAtomicCounterBindings = 1;

const mediump int gl_MaxFragmentAtomicCounterBuffers = 0;
const mediump int gl_MaxVertexAtomicCounterBuffers = 0;
const mediump int gl_MaxCombinedAtomicCounterBuffers = 1;
const mediump int gl_MaxAtomicCounterBufferSize = 32;
```

7.3 Built-In Uniform State

As an aid to accessing OpenGL ES processing state, the following uniform variables are built into the OpenGL ES Shading Language.

```
//
// Depth range in window coordinates,
// section 12.5.1 "Controlling the Viewport" in the
// OpenGL ES Graphics System Specification.
//
struct gl_DepthRangeParameters {
    highp float near;          // n
    highp float far;           // f
    highp float diff;          // f - n
};
uniform gl_DepthRangeParameters gl_DepthRange;

uniform lowp int gl_NumSamples;
```

7.4 Redefining Built-in Blocks

The `gl_PerVertex` block can be redeclared in a shader to explicitly indicate what subset of the members will be used. This is necessary to establish the interface between multiple programs. If the `gl_PerVertex` block is not redefined in a given program, the intrinsically declared definition of that block is used for the program interface.

For example:

```
out gl_PerVertex
{
    highp vec4 gl_Position;    // will use gl_Position
    highp vec4 t;              // error, only gl_PerVertex members allowed
}; // no other members of gl_PerVertex will be used
```

This establishes the output interface the shader will use with the subsequent pipeline stage. It must be a subset of the built-in members of `gl_PerVertex`. Such a redeclaration can also add the invariant qualifier and interpolation qualifiers.

Other layout qualifiers, like location, cannot be added to such a redeclaration, unless specifically stated.

If a built-in interface block is redeclared, it must appear in the shader before any use of any member included in the built-in declaration, or a compile-time error will result. It is also a compile-time error to redeclare the block more than once or to redeclare a built-in block and then use a member from that built-in block that was not included in the redeclaration. Also, if a built-in interface block is redeclared, no member of the built-in declaration can be redeclared outside the block redeclaration. If multiple shaders using members of a built-in block belonging to the same interface are linked together in the same program, they must all redeclare the built-in block in the same way, as described in section 4.3.9 ("Interface Blocks") for interface block matching, or a link-time error will result. It will also be a link-time error if some shaders in a program redeclare a specific built-in interface block while another shader in that program does not redeclare that interface block yet still uses a member of that interface block. If a built-in block interface is formed across shaders in different programs, the shaders must all redeclare the built-in block in the same way (as described for a single program), or the values passed along the interface are undefined.

8 Built-in Functions

The OpenGL ES Shading Language defines an assortment of built-in convenience functions for scalar and vector operations. Many of these built-in functions can be used in more than one type of shader, but some are intended to provide a direct mapping to hardware and so are available only for a specific type of shader.

The built-in functions basically fall into three categories:

- They expose some necessary hardware functionality in a convenient way such as accessing a texture map. There is no way in the language for these functions to be emulated by a shader.
- They represent a trivial operation (clamp, mix, etc.) that is very simple for the user to write, but they are very common and may have direct hardware support. It is a very hard problem for the compiler to map expressions to complex assembler instructions.
- They represent an operation graphics hardware is likely to accelerate at some point. The trigonometry functions fall into this category.

Many of the functions are similar to the same named ones in common C libraries, but they support vector input as well as the more traditional scalar input.

Applications should be encouraged to use the built-in functions rather than do the equivalent computations in their own shader code since the built-in functions are assumed to be optimal (e.g., perhaps supported directly in hardware).

When the built-in functions are specified below, where the input arguments (and corresponding output) can be **float**, **vec2**, **vec3**, or **vec4**, *genFType* is used as the argument. Where the input arguments (and corresponding output) can be **int**, **ivec2**, **ivec3**, or **ivec4**, *genIType* is used as the argument. Where the input arguments (and corresponding output) can be **uint**, **uvec2**, **uvec3**, or **uvec4**, *genUType* is used as the argument. Where the input arguments (or corresponding output) can be **bool**, **bvec2**, **bvec3**, or **bvec4**, *genBType* is used as the argument. For any specific use of a function, the actual types substituted for *genFType*, *genIType*, *genUType*, or *genBType* have to have the same number of components for all arguments and for the return type. Similarly for *mat*, which can be any matrix basic type.

The precision of built-in functions is dependent on the function and arguments. There are three categories:

- Some functions have predefined precisions. The precision is specified by the function signature e.g.
 - Floating-point pack and unpack functions
 - imageSize** and **textureSize** Functions
 - Atomic operation functions
- For the texture sampling, image load and image store functions, the precision of the return type matches the precision of the sampler type:

```

uniform lowp sampler2D sampler;
highp vec2 coord;
...
lowp vec4 col = texture (sampler, coord); // texture() returns lowp

```

- For other built-in functions, a call will return a precision qualification matching the highest precision qualification of the call's input arguments. See Section 4.7.3 “Precision Qualifiers” for more detail.

The built-in functions are assumed to be implemented according to the equations specified in the following sections. The precision at which the calculations are performed follows the general rules for precision of operations as specified in section 4.7.3 “Precision Qualifiers”.

Example:

$$\textit{normalize} \begin{pmatrix} x \\ y \\ z \end{pmatrix} = \frac{1}{\sqrt{x^2 + y^2 + z^2}} \begin{pmatrix} x \\ y \\ z \end{pmatrix}$$

If the input vector is lowp, the entire calculation is performed at lowp. For some inputs, this will cause the calculation to overflow, even when the correct result is within the range of lowp.

8.1 Angle and Trigonometry Functions

Function parameters specified as *angle* are assumed to be in units of radians. In no case will any of these functions result in a divide by zero error. If the divisor of a ratio is 0, then results will be undefined.

These all operate component-wise. The description is per component.

Syntax	Description
genFType radians (genFType <i>degrees</i>)	Converts <i>degrees</i> to radians, i.e., $\frac{\pi}{180} \cdot \text{degrees}$
genFType degrees (genFType <i>radians</i>)	Converts <i>radians</i> to degrees, i.e., $\frac{180}{\pi} \cdot \text{radians}$
genFType sin (genFType <i>angle</i>)	The standard trigonometric sine function.
genFType cos (genFType <i>angle</i>)	The standard trigonometric cosine function.
genFType tan (genFType <i>angle</i>)	The standard trigonometric tangent.
genFType asin (genFType <i>x</i>)	Arc sine. Returns an angle whose sine is <i>x</i> . The range of values returned by this function is $\left[-\frac{\pi}{2}, \frac{\pi}{2}\right]$. Results are undefined if $ x > 1$.
genFType acos (genFType <i>x</i>)	Arc cosine. Returns an angle whose cosine is <i>x</i> . The range of values returned by this function is $[0, \pi]$. Results are undefined if $ x > 1$.
genFType atan (genFType <i>y</i> , genFType <i>x</i>)	Arc tangent. Returns an angle whose tangent is <i>y/x</i> . The signs of <i>x</i> and <i>y</i> are used to determine what quadrant the angle is in. The range of values returned by this function is $[-\pi, \pi]$. Results are undefined if <i>x</i> and <i>y</i> are both 0.
genFType atan (genFType <i>y_over_x</i>)	Arc tangent. Returns an angle whose tangent is <i>y_over_x</i> . The range of values returned by this function is $\left[-\frac{\pi}{2}, \frac{\pi}{2}\right]$.

Syntax	Description
genFType sinh (genFType x)	Returns the hyperbolic sine function $\frac{e^x - e^{-x}}{2}$
genFType cosh (genFType x)	Returns the hyperbolic cosine function $\frac{e^x + e^{-x}}{2}$
genFType tanh (genFType x)	Returns the hyperbolic tangent function $\frac{\sinh(x)}{\cosh(x)}$
genFType asinh (genFType x)	Arc hyperbolic sine; returns the inverse of sinh .
genFType acosh (genFType x)	Arc hyperbolic cosine; returns the non-negative inverse of cosh . Results are undefined if $x < 1$.
genFType atanh (genFType x)	Arc hyperbolic tangent; returns the inverse of tanh . Results are undefined if $ x \geq 1$.

8.2 Exponential Functions

These all operate component-wise. The description is per component.

Syntax	Description
genFType pow (genFType x , genFType y)	Returns x raised to the y power, i.e., x^y Results are undefined if $x < 0$. Results are undefined if $x = 0$ and $y \leq 0$.
genFType exp (genFType x)	Returns the natural exponentiation of x , i.e., e^x .
genFType log (genFType x)	Returns the natural logarithm of x , i.e., returns the value y which satisfies the equation $x = e^y$. Results are undefined if $x \leq 0$.
genFType exp2 (genFType x)	Returns 2 raised to the x power, i.e., 2^x .
genFType log2 (genFType x)	Returns the base 2 logarithm of x , i.e., returns the value y which satisfies the equation $x = 2^y$. Results are undefined if $x \leq 0$.

Syntax	Description
genFType sqrt (genFType x)	Returns \sqrt{x} Results are undefined if $x < 0$.
genFType inversesqrt (genFType x)	Returns $\frac{1}{\sqrt{x}}$ Results are undefined if $x \leq 0$.

8.3 Common Functions

These all operate component-wise. The description is per component.

Syntax	Description
genFType abs (genFType x) genIType abs (genIType x)	Returns x if $x \geq 0$, otherwise it returns $-x$.
genFType sign (genFType x) genIType sign (genIType x)	Returns 1.0 if $x > 0$, 0.0 if $x = 0$ or -1.0 if $x < 0$.
genFType floor (genFType x)	Returns a value equal to the nearest integer that is less than or equal to x .
genFType trunc (genFType x)	Returns a value equal to the nearest integer to x whose absolute value is not larger than the absolute value of x .
genFType round (genFType x)	Returns a value equal to the nearest integer to x . The fraction 0.5 will round in a direction chosen by the implementation, presumably the direction that is fastest. This includes the possibility that round (x) returns the same value as roundEven (x) for all values of x .
genFType roundEven (genFType x)	Returns a value equal to the nearest integer to x . A fractional part of 0.5 will round toward the nearest even integer. (Both 3.5 and 4.5 for x will return 4.0.)
genFType ceil (genFType x)	Returns a value equal to the nearest integer that is greater than or equal to x .
genFType fract (genFType x)	Returns $x - \mathbf{floor}(x)$.

Syntax	Description
<code>genFType mod (genFType <i>x</i>, float <i>y</i>)</code> <code>genFType mod (genFType <i>x</i>, genFType <i>y</i>)</code>	Modulus. Returns $x - y * \text{floor}(x/y)$.
<code>genFType modf (genFType <i>x</i>, out genFType <i>i</i>)</code>	Returns the fractional part of <i>x</i> and sets <i>i</i> to the integer part (as a whole number floating point value). Both the return value and the output parameter will have the same sign as <i>x</i> .
<code>genFType min (genFType <i>x</i>, genFType <i>y</i>)</code> <code>genFType min (genFType <i>x</i>, float <i>y</i>)</code> <code>genIType min (genIType <i>x</i>, genIType <i>y</i>)</code> <code>genIType min (genIType <i>x</i>, int <i>y</i>)</code> <code>genUType min (genUType <i>x</i>, genUType <i>y</i>)</code> <code>genUType min (genUType <i>x</i>, uint <i>y</i>)</code>	Returns <i>y</i> if $y < x$, otherwise it returns <i>x</i> .
<code>genFType max (genFType <i>x</i>, genFType <i>y</i>)</code> <code>genFType max (genFType <i>x</i>, float <i>y</i>)</code> <code>genIType max (genIType <i>x</i>, genIType <i>y</i>)</code> <code>genIType max (genIType <i>x</i>, int <i>y</i>)</code> <code>genUType max (genUType <i>x</i>, genUType <i>y</i>)</code> <code>genUType max (genUType <i>x</i>, uint <i>y</i>)</code>	Returns <i>y</i> if $x < y$, otherwise it returns <i>x</i> .
<code>genFType clamp (genFType <i>x</i>, genFType <i>minVal</i>, genFType <i>maxVal</i>)</code> <code>genFType clamp (genFType <i>x</i>, float <i>minVal</i>, float <i>maxVal</i>)</code> <code>genIType clamp (genIType <i>x</i>, genIType <i>minVal</i>, genIType <i>maxVal</i>)</code> <code>genIType clamp (genIType <i>x</i>, int <i>minVal</i>, int <i>maxVal</i>)</code> <code>genUType clamp (genUType <i>x</i>, genUType <i>minVal</i>, genUType <i>maxVal</i>)</code> <code>genUType clamp (genUType <i>x</i>, uint <i>minVal</i>, uint <i>maxVal</i>)</code>	Returns min (max (<i>x</i> , <i>minVal</i>), <i>maxVal</i>). Results are undefined if $\text{minVal} > \text{maxVal}$.

Syntax	Description
<code>genFType mix (genFType x, genFType y, genFType a)</code> <code>genFType mix (genFType x, genFType y, float a)</code>	Returns the linear blend of x and y , i.e., $x \cdot (1 - a) + y \cdot a$
<code>genFType mix (genFType x, genFType y, genBType a)</code> <code>genIType mix (genIType x, genIType y, genBType a)</code> <code>genUType mix (genUType x, genUType y, genBType a)</code> <code>genBType mix (genBType x, genBType y, genBType a)</code>	<p>Selects which vector each returned component comes from. For a component of a that is false, the corresponding component of x is returned. For a component of a that is true, the corresponding component of y is returned. Components of x and y that are not selected are allowed to be invalid floating point values and will have no effect on the results. Thus, this provides different functionality than</p> <p style="text-align: center;"><code>genFType mix(genFType x, genFType y, genFType(a))</code></p> <p>where a is a Boolean vector.</p>
<code>genFType step (genFType edge, genFType x)</code> <code>genFType step (float edge, genFType x)</code>	Returns 0.0 if $x < edge$, otherwise it returns 1.0.
<code>genFType smoothstep (genFType edge0, genFType edge1, genFType x)</code> <code>genFType smoothstep (float edge0, float edge1, genFType x)</code>	<p>Returns 0.0 if $x \leq edge0$ and 1.0 if $x \geq edge1$ and performs smooth Hermite interpolation between 0 and 1 when $edge0 < x < edge1$. This is useful in cases where you would want a threshold function with a smooth transition. This is equivalent to:</p> <pre>genFType t; t = clamp ((x - edge0) / (edge1 - edge0), 0, 1); return t * t * (3 - 2 * t);</pre> <p>Results are undefined if $edge0 \geq edge1$.</p>
<code>genBType isnan (genFType x)</code>	Returns true if x holds a NaN. Returns false otherwise.
<code>genBType isinf (genFType x)</code>	Returns true if x holds a positive infinity or negative infinity. Returns false otherwise.

Syntax	Description
genIType floatBitsToInt (genFType <i>value</i>) genUType floatBitsToUint (genFType <i>value</i>)	Returns a signed or unsigned highp integer value representing the encoding of a floating-point value. For highp floating point, the value's bit level representation is preserved. For mediump and lowp, the value is first converted to highp floating point and the encoding of that value is returned.
genFType intBitsToFloat (genIType <i>value</i>) genFType uintBitsToFloat (genUType <i>value</i>)	Returns a highp floating-point value corresponding to a signed or unsigned integer encoding of a floating-point value. If an inf or NaN is passed in, it will not signal, and the resulting floating point value is unspecified. Otherwise, the bit-level representation is preserved. For lowp and mediump, the value is first converted to the corresponding signed or unsigned highp integer and then reinterpreted as a highp floating point value as before.
GenType fma (genType a, genType b, genType c)	<p>Computes and returns $a * b + c$</p> <p>In uses where the return value is eventually consumed by a variable declared as precise:</p> <ul style="list-style-type: none"> • fma() is considered a single operation, whereas the expression $a * b + c$ consumed by a variable declared precise is considered two operations. • The precision of fma() can differ from the precision of the expression $a * b + c$. • fma() will be computed with the same precision as any other fma() consumed by a precise variable, giving invariant results for the same input values of a, b, and c. <p>Otherwise, in the absence of precise consumption, there are no special constraints on the number of operations or difference in precision between fma() and the expression $a * b + c$.</p>

Syntax	Description
<p>highp genFType frexp(highp genFType <i>x</i>, out highp genIType <i>exp</i>);</p>	<p>The function <code>frexp()</code> splits each single-precision floating-point number in <i>x</i> into a binary significand, a floating-point number in the range [0.5, 1.0), and an integral exponent of two, such that:</p> $x = \text{significand} \cdot 2^{\text{exponent}}$ <p>The significand is returned by the function; the exponent is returned in the parameter <i>exp</i>. For a floating-point value of zero, the significand and exponent are both zero. If an implementation supports signed zero, an input value of minus zero should return a significand of minus zero. For a floating-point value that is an infinity or is not a number, the results of <code>frexp()</code> are undefined. If the input <i>x</i> is a vector, this operation is performed in a component-wise manner; the value returned by the function and the value written to <i>exp</i> are vectors with the same number of components as <i>x</i>.</p>
<p>highp genFType ldexp(highp genFType <i>x</i>, in highp genIType <i>exp</i>);</p>	<p>The function <code>ldexp()</code> builds a single-precision floating-point number from each significand component in <i>x</i> and the corresponding integral exponent of two in <i>exp</i>, returning:</p> $x = \text{significand} \cdot 2^{\text{exponent}}$ <p>If <i>exponent</i> is greater than +128, the value returned is undefined. If <i>exponent</i> is less than -126, the value returned may be flushed to zero. Additionally, splitting the value into a significand and exponent using frexp() and then reconstructing a floating-point value using ldexp() should yield the original input for zero and all finite non-subnormal values. If the input <i>x</i> is a vector, this operation is performed in a component-wise manner; the value passed in <i>exp</i> and returned by the function are vectors with the same number of components as <i>x</i>.</p>

8.4 Floating-Point Pack and Unpack Functions

These functions do not operate component-wise, rather as described in each case.

Syntax	Description
highp uint packSnorm2x16 (vec2 v)	<p>First, converts each component of the normalized floating-point value v into 16-bit integer values. Then, the results are packed into the returned 32-bit unsigned integer.</p> <p>The conversion for component c of v to fixed point is done as follows:</p> $\text{fixed point value} = \text{round}(\text{clamp}(c, -1, +1) * 32767.0)$ <p>The first component of the vector will be written to the least significant bits of the output; the last component will be written to the most significant bits.</p>
highp vec2 unpackSnorm2x16 (highp uint p)	<p>First, unpacks a single 32-bit unsigned integer p into a pair of 16-bit signed integers. Then, each component is converted to a normalized floating-point value to generate the returned two-component vector.</p> <p>The conversion for unpacked fixed-point value f to floating point is done as follows:</p> $\text{floating point value} = \text{clamp}(f / 32767.0, -1, +1)$ <p>The first component of the returned vector will be extracted from the least significant bits of the input; the last component will be extracted from the most significant bits.</p>
highp uint packUnorm2x16 (vec2 v)	<p>First, converts each component of the normalized floating-point value v into 16-bit integer values. Then, the results are packed into the returned 32-bit unsigned integer.</p> <p>The conversion for component c of v to fixed point is done as follows:</p> $\text{fixed point value} = \text{round}(\text{clamp}(c, 0, +1) * 65535.0)$ <p>The first component of the vector will be written to the least significant bits of the output; the last component will be written to the most significant bits.</p>

Syntax	Description
highp vec2 unpackUnorm2x16 (highp uint p)	<p>First, unpacks a single 32-bit unsigned integer p into a pair of 16-bit unsigned integers. Then, each component is converted to a normalized floating-point value to generate the returned two-component vector.</p> <p>The conversion for unpacked fixed-point value f to floating point is done as follows:</p> $\text{floating point value} = f / 65535.0$ <p>The first component of the returned vector will be extracted from the least significant bits of the input; the last component will be extracted from the most significant bits.</p>
highp uint packHalf2x16 (mediump vec2 v)	<p>Returns an unsigned integer obtained by converting the components of a two-component floating-point vector to the 16-bit floating-point representation found in the OpenGL ES Specification, and then packing these two 16-bit integers into a 32-bit unsigned integer.</p> <p>The first vector component specifies the 16 least-significant bits of the result; the second component specifies the 16 most-significant bits.</p>
mediump vec2 unpackHalf2x16 (highp uint v)	<p>Returns a two-component floating-point vector with components obtained by unpacking a 32-bit unsigned integer into a pair of 16-bit values, interpreting those values as 16-bit floating-point numbers according to the OpenGL ES Specification, and converting them to 32-bit floating-point values.</p> <p>The first component of the vector is obtained from the 16 least-significant bits of v; the second component is obtained from the 16 most-significant bits of v.</p>

Syntax	Description
highp uint packUnorm4x8 (mediump vec4 v)	<p>First convert each component of four-component vector of normalized floating-point values into 8-bit unsigned integer values.</p> <p>fixed-value = round(clamp(floating-value, 0, 1)*255.0</p> <p>Then, the results are packed into the returned 32-bit unsigned integer. The first component of the vector will be written to the least significant bits of the output; the last component will be written to the most significant bits.</p>
highp uint packSnorm4x8 (mediump vec4 v);	<p>First convert each component of four-component vector of normalized floating-point values into signed 8-bit integer values.</p> <p>fixed-value = round(clamp(floating-value, -1, 1)*127.0</p> <p>Then, the results are packed into the returned 32-bit unsigned integer. The first component of the vector will be written to the least significant bits of the output; the last component will be written to the most significant bits.</p>
mediump vec4 unpackUnorm4x8 (highp uint v)	<p>First unpack a single 32-bit unsigned integer into four 8-bit unsigned integers. Then, each component is converted to a normalized floating-point value to generate the returned four-component vector.</p> <p>float-value = fixed-value / 255.0</p> <p>The first component of the vector will be extracted from the least significant bits of the input; the last component will be extracted from the most significant bits.</p>

Syntax	Description
mediump vec4 unpackSnorm4x8 (highp uint v);	<p>First unpack a single 32-bit unsigned integer into four 8-bit signed integers. Then, each component is converted to a normalized floating-point value to generate the returned four-component vector.</p> <p>float-value = clamp(fixed-value / 127.0, -1, 1)</p> <p>The first component of the vector will be extracted from the least significant bits of the input; the last component will be extracted from the most significant bits.</p>

8.5 Geometric Functions

These operate on vectors as vectors, not component-wise.

Syntax	Description
float length (genFType <i>x</i>)	Returns the length of vector <i>x</i> , i.e., $\sqrt{x[0]^2 + x[1]^2 + \dots}$
float distance (genFType <i>p0</i> , genFType <i>p1</i>)	Returns the distance between <i>p0</i> and <i>p1</i> , i.e., length (<i>p0</i> - <i>p1</i>)
float dot (genFType <i>x</i> , genFType <i>y</i>)	Returns the dot product of <i>x</i> and <i>y</i> , i.e., $x[0] \cdot y[0] + x[1] \cdot y[1] + \dots$
vec3 cross (vec3 <i>x</i> , vec3 <i>y</i>)	Returns the cross product of <i>x</i> and <i>y</i> , i.e., $\begin{bmatrix} x[1] \cdot y[2] - y[1] \cdot x[2] \\ x[2] \cdot y[0] - y[2] \cdot x[0] \\ x[0] \cdot y[1] - y[0] \cdot x[1] \end{bmatrix}$
genFType normalize (genFType <i>x</i>)	Returns a vector in the same direction as <i>x</i> but with a length of 1 i.e. $\frac{x}{length(x)}$
genFType faceforward (genFType <i>N</i> , genFType <i>I</i> , genFType <i>Nref</i>)	If dot (<i>Nref</i> , <i>I</i>) < 0 return <i>N</i> , otherwise return - <i>N</i> .

Syntax	Description
genFType reflect (genFType <i>I</i> , genFType <i>N</i>)	<p>For the incident vector <i>I</i> and surface orientation <i>N</i>, returns the reflection direction:</p> $I - 2 * \mathbf{dot}(N, I) * N$ <p><i>N</i> must already be normalized in order to achieve the desired result.</p>
genFType refract (genFType <i>I</i> , genFType <i>N</i> , float <i>eta</i>)	<p>For the incident vector <i>I</i> and surface normal <i>N</i>, and the ratio of indices of refraction <i>eta</i>, return the refraction vector. The result is computed by</p> $k = 1.0 - eta * eta * (1.0 - \mathbf{dot}(N, I) * \mathbf{dot}(N, I))$ <p>if (<i>k</i> < 0.0) return genFType(0.0) else return $eta * I - (eta * \mathbf{dot}(N, I) + \mathbf{sqrt}(k)) * N$</p> <p>The input parameters for the incident vector <i>I</i> and the surface normal <i>N</i> must already be normalized to get the desired results.</p>

8.6 Matrix Functions

Syntax	Description
mat matrixCompMult (mat <i>x</i> , mat <i>y</i>)	Multiply matrix <i>x</i> by matrix <i>y</i> component-wise, i.e., result[i][j] is the scalar product of <i>x</i> [i][j] and <i>y</i> [i][j]. Note: to get linear algebraic matrix multiplication, use the multiply operator (*).
mat2 outerProduct (vec2 <i>c</i> , vec2 <i>r</i>) mat3 outerProduct (vec3 <i>c</i> , vec3 <i>r</i>) mat4 outerProduct (vec4 <i>c</i> , vec4 <i>r</i>) mat2x3 outerProduct (vec3 <i>c</i> , vec2 <i>r</i>) mat3x2 outerProduct (vec2 <i>c</i> , vec3 <i>r</i>) mat2x4 outerProduct (vec4 <i>c</i> , vec2 <i>r</i>) mat4x2 outerProduct (vec2 <i>c</i> , vec4 <i>r</i>) mat3x4 outerProduct (vec4 <i>c</i> , vec3 <i>r</i>) mat4x3 outerProduct (vec3 <i>c</i> , vec4 <i>r</i>)	Treats the first parameter <i>c</i> as a column vector (matrix with one column) and the second parameter <i>r</i> as a row vector (matrix with one row) and does a linear algebraic matrix multiply <i>c</i> * <i>r</i> , yielding a matrix whose number of rows is the number of components in <i>c</i> and whose number of columns is the number of components in <i>r</i> .
mat2 transpose (mat2 <i>m</i>) mat3 transpose (mat3 <i>m</i>) mat4 transpose (mat4 <i>m</i>) mat2x3 transpose (mat3x2 <i>m</i>) mat3x2 transpose (mat2x3 <i>m</i>) mat2x4 transpose (mat4x2 <i>m</i>) mat4x2 transpose (mat2x4 <i>m</i>) mat3x4 transpose (mat4x3 <i>m</i>) mat4x3 transpose (mat3x4 <i>m</i>)	Returns a matrix that is the transpose of <i>m</i> . The input matrix <i>m</i> is not modified.
float determinant (mat2 <i>m</i>) float determinant (mat3 <i>m</i>) float determinant (mat4 <i>m</i>)	Returns the determinant of <i>m</i> .
mat2 inverse (mat2 <i>m</i>) mat3 inverse (mat3 <i>m</i>) mat4 inverse (mat4 <i>m</i>)	Returns a matrix that is the inverse of <i>m</i> . The input matrix <i>m</i> is not modified. The values in the returned matrix are undefined if <i>m</i> is singular or poorly-conditioned (nearly singular).

8.7 Vector Relational Functions

Relational and equality operators ($<$, $<=$, $>$, $>=$, $==$, $!=$) are defined to produce scalar Boolean results. For vector results, use the following built-in functions. Below, “bvec” is a placeholder for one of **bvec2**, **bvec3**, or **bvec4**, “ivec” is a placeholder for one of **ivec2**, **ivec3**, or **ivec4**, “uvec” is a placeholder for **uvec2**, **uvec3**, or **uvec4**, and “vec” is a placeholder for **vec2**, **vec3**, or **vec4**. In all cases, the sizes of the input and return vectors for any particular call must match.

Syntax	Description
bvec lessThan (vec x , vec y) bvec lessThan (ivec x , ivec y) bvec lessThan (uvec x , uvec y)	Returns the component-wise compare of $x < y$.
bvec lessThanEqual (vec x , vec y) bvec lessThanEqual (ivec x , ivec y) bvec lessThanEqual (uvec x , uvec y)	Returns the component-wise compare of $x <= y$.
bvec greaterThan (vec x , vec y) bvec greaterThan (ivec x , ivec y) bvec greaterThan (uvec x , uvec y)	Returns the component-wise compare of $x > y$.
bvec greaterThanEqual (vec x , vec y) bvec greaterThanEqual (ivec x , ivec y) bvec greaterThanEqual (uvec x , uvec y)	Returns the component-wise compare of $x >= y$.
bvec equal (vec x , vec y) bvec equal (ivec x , ivec y) bvec equal (uvec x , uvec y) bvec equal (bvec x , bvec y)	Returns the component-wise compare of $x == y$.
bvec notEqual (vec x , vec y) bvec notEqual (ivec x , ivec y) bvec notEqual (uvec x , uvec y) bvec notEqual (bvec x , bvec y)	Returns the component-wise compare of $x != y$.
bool any (bvec x)	Returns true if any component of x is true .
bool all (bvec x)	Returns true only if all components of x are true .
bvec not (bvec x)	Returns the component-wise logical complement of x .

8.8 Integer Functions

Syntax	Description
<p>genIType bitfieldExtract(genIType <i>value</i>, int <i>offset</i>, int <i>bits</i>);</p> <p>genUType bitfieldExtract(genUType <i>value</i>, int <i>offset</i>, int <i>bits</i>);</p>	<p>Extracts bits <i>offset</i> through <i>offset+bits-1</i> from each component in <i>value</i>, returning them in the least significant bits of corresponding component of the result. For unsigned data types, the most significant bits of the result will be set to zero. For signed data types, the most significant bits will be set to the value of bit <i>offset+base-1</i>. If <i>bits</i> is zero, the result will be zero. The result will be undefined if <i>offset</i> or <i>bits</i> is negative, or if the sum of <i>offset</i> and <i>bits</i> is greater than the number of bits used to store the operand. Note that for vector versions of bitfieldExtract(), a single pair of <i>offset</i> and <i>bits</i> values is shared for all components.</p> <p>The precision qualification of the value returned from bitfieldExtract() matches the precision qualification of the call's input argument “<i>value</i>”.</p>
<p>genIType bitfieldInsert(genIType <i>base</i>, genIType <i>insert</i>, int <i>offset</i>, int <i>bits</i>);</p> <p>genUType bitfieldInsert(genUType <i>base</i>, genUType <i>insert</i>, int <i>offset</i>, int <i>bits</i>);</p>	<p>Inserts the <i>bits</i> least significant bits of each component of <i>insert</i> into the corresponding component of <i>base</i>. The result will have bits numbered <i>offset</i> through <i>offset+bits-1</i> taken from bits 0 through <i>bits-1</i> of <i>insert</i>, and all other bits taken directly from the corresponding bits of <i>base</i>. If <i>bits</i> is zero, the result will simply be <i>base</i>. The result will be undefined if <i>offset</i> or <i>bits</i> is negative, or if the sum of <i>offset</i> and <i>bits</i> is greater than the number of bits used to store the operand. Note that for vector versions of bitfieldInsert(), a single pair of <i>offset</i> and <i>bits</i> values is shared for all components.</p> <p>The precision qualification of the value returned from bitfieldInsert matches the highest precision qualification of the call's input arguments “<i>base</i>” and “<i>insert</i>”.</p>
<p>highp genIType bitfieldReverse(highp genIType <i>value</i>);</p> <p>highp genUType bitfieldReverse(highp genUType <i>value</i>);</p>	<p>Reverses the bits of <i>value</i>. The bit numbered <i>n</i> of the result will be taken from bit $(bits-1)-n$ of <i>value</i>, where <i>bits</i> is the total number of bits used to represent <i>value</i>.</p>
<p>lowp genIType bitCount(genIType <i>value</i>);</p> <p>lowp genIType bitCount(genUType <i>value</i>);</p>	<p>Returns the number of one bits in the binary representation of <i>value</i>.</p>

Syntax	Description
lowp genIType findLSB (genIType <i>value</i>); lowp genIType findLSB (genUType <i>value</i>);	Returns the bit number of the least significant one bit in the binary representation of <i>value</i> . If <i>value</i> is zero, -1 will be returned.
lowp genIType findMSB (highp genIType <i>value</i>); lowp genIType findMSB (highp genUType <i>value</i>);	Returns the bit number of the most significant bit in the binary representation of <i>value</i> . For positive integers, the result will be the bit number of the most significant one bit. For negative integers, the result will be the bit number of the most significant zero bit. For a <i>value</i> of zero or negative one, -1 will be returned.
highp genUType uaddCarry (highp genUType <i>x</i> , highp genUType <i>y</i> , out lowp genUType <i>carry</i>);	Adds 32-bit unsigned integers or vectors <i>x</i> and <i>y</i> , returning the sum modulo 2^{32} . The value <i>carry</i> is set to zero if the sum was less than 2^{32} , or one otherwise.
highp genUType usubBorrow (highp genUType <i>x</i> , highp genUType <i>y</i> , out lowp genUType <i>borrow</i>);	Subtracts the 32-bit unsigned integer or vector <i>y</i> from <i>x</i> , returning the difference if non-negative or 2^{32} plus the difference, otherwise. The value <i>borrow</i> is set to zero if $x \geq y$, or one otherwise.
void umulExtended (highp genUType <i>x</i> , highp genUType <i>y</i> , out highp genUType <i>msb</i> , out highp genUType <i>lsb</i>); void imulExtended (highp genIType <i>x</i> , highp genIType <i>y</i> , out highp genIType <i>msb</i> , out highp genIType <i>lsb</i>);	Multiply 32-bit unsigned or signed integers or vectors <i>x</i> and <i>y</i> , producing a 64-bit result. The 32 least significant bits are returned in <i>lsb</i> ; the 32 most significant bits are returned in <i>msb</i> .

8.9 Texture Functions

Texture lookup functions are available in all shading stages. However, level of detail is implicitly computed only for fragment shaders. Other shaders operate as though the base level of detail were computed as zero. The functions in the table below provide access to textures through samplers, as set up through the OpenGL ES API. Texture properties such as size, pixel format, number of dimensions, filtering method, number of mip-map levels, depth comparison, and so on are also defined by OpenGL ES API calls. Such properties are taken into account as the texture is accessed via the built-in functions defined below.

Texture data can be stored by the GL as floating point, unsigned normalized integer, unsigned integer or signed integer data. This is determined by the type of the internal format of the texture. Texture lookups on unsigned normalized integer and floating point data return floating point values in the range [0, 1].

Texture lookup functions are provided that can return their result as floating point, unsigned integer or signed integer, depending on the sampler type passed to the lookup function. Care must be taken to use the right sampler type for texture access. The following table lists the supported combinations of sampler types and texture internal formats. Blank entries are unsupported. Doing a texture lookup will return undefined values for unsupported combinations.

Internal Texture Format	Floating Point Sampler Types	Signed Integer Sampler Types	Unsigned Integer Sampler Types
Floating point	Supported		
Normalized Integer	Supported		
Signed Integer		Supported	
Unsigned Integer			Supported

If an integer sampler type is used, the result of a texture lookup is an **ivec4**. If an unsigned integer sampler type is used, the result of a texture lookup is a **uvec4**. If a floating point sampler type is used, the result of a texture lookup is a **vec4**.

In the prototypes below, the “g” in the return type “*gvec4*” is used as a placeholder for nothing, “*i*”, or “*u*” making a return type of **vec4**, **ivec4**, or **uvec4**. In these cases, the sampler argument type also starts with “g”, indicating the same substitution done on the return type; it is either a floating point, signed integer, or unsigned integer sampler, matching the basic type of the return type, as described above.

For shadow forms (the sampler parameter is a shadow-type), a depth comparison lookup on the depth texture bound to *sampler* is done as described in section 8.19 “Texture Comparison Modes” of the OpenGL ES Graphics System Specification. See the table below for which component specifies D_{ref} . The texture bound to *sampler* must be a depth texture, or results are undefined. If a non-shadow texture call is made to a sampler that represents a depth texture with depth comparisons turned on, then results are undefined. If a shadow texture call is made to a sampler that represents a depth texture with depth comparisons turned off, then results are undefined. If a shadow texture call is made to a sampler that does not represent a depth texture, then results are undefined.

In all functions below, the *bias* parameter is optional for fragment shaders. The *bias* parameter is not accepted in any other stage. For a fragment shader, if *bias* is present, it is added to the implicit level of detail prior to performing the texture access operation. No *bias* or *lod* parameters for multi-sample textures or texture buffers are supported because mipmaps are not allowed for these types of textures.

The implicit level of detail is selected as follows: For a texture that is not mip-mapped, the texture is used directly. If it is mip-mapped and running in a fragment shader, the LOD computed by the implementation is used to do the texture lookup. If it is mip-mapped and running in a non fragment shader, then the base texture is used.

Some texture functions (non-“**Lod**” and non-“**Grad**” versions) may require implicit derivatives. Implicit derivatives are undefined within non-uniform control flow and for vertex texture fetches.

For **Cube** forms, the direction of *P* is used to select which face to do a 2-dimensional texture lookup in, as described in section 8.12 “Cube Map Texture Selection” in the OpenGL ES Graphics System Specification.

For **Array** forms, the array layer used will be

$$\max(0, \min(d - 1, \text{floor}(\text{layer} + 0.5)))$$

where *d* is the depth of the texture array and *layer* comes from the component indicated in the tables below.

8.9.1 Texture Query Functions

Syntax	Description
highp ivec2 textureSize (gsampler2D <i>sampler</i> , int <i>lod</i>)	Returns the dimensions of level <i>lod</i> for the texture bound to <i>sampler</i> , as described in section 11.1.3.4 “Texture Queries” of the OpenGL ES 3.2 Graphics System Specification. The components in the return value are filled in, in order, with the width, height, depth of the texture. For the array forms, the last component of the return value is the number of layers in the texture array.
highp ivec3 textureSize (gsampler3D <i>sampler</i> , int <i>lod</i>)	
highp ivec2 textureSize (gsamplerCube <i>sampler</i> , int <i>lod</i>)	
highp ivec2 textureSize (gsampler2DMS <i>sampler</i>)	
highp ivec3 textureSize (gsampler2DArray <i>sampler</i> , int <i>lod</i>)	
highp ivec2 textureSize (samplerCubeShadow <i>sampler</i> , int <i>lod</i>)	
highp ivec2 textureSize (sampler2DShadow <i>sampler</i> , int <i>lod</i>)	
highp ivec3 textureSize (sampler2DArrayShadow <i>sampler</i> , int <i>lod</i>)	
highp int textureSize (gsamplerBuffer <i>sampler</i>)	
highp ivec3 textureSize (gsamplerCubeArray <i>sampler</i> , int <i>lod</i>)	
highp ivec3 textureSize (gsamplerCubeArrayShadow <i>sampler</i> , int <i>lod</i>)	
highp ivec3 textureSize (gsampler2DMSArray <i>sampler</i>)	

8.9.2 Texel Lookup Functions

Syntax	Description
<p>gvec4 texture (gsampler2D <i>sampler</i>, vec2 <i>P</i> [, float <i>bias</i>])</p> <p>gvec4 texture (gsampler3D <i>sampler</i>, vec3 <i>P</i> [, float <i>bias</i>])</p> <p>gvec4 texture (gsamplerCube <i>sampler</i>, vec3 <i>P</i> [, float <i>bias</i>])</p> <p>gvec4 texture (gsampler2DArray <i>sampler</i>, vec3 <i>P</i> [, float <i>bias</i>])</p> <p>float texture (sampler2DShadow <i>sampler</i>, vec3 <i>P</i> [, float <i>bias</i>])</p> <p>float texture (samplerCubeShadow <i>sampler</i>, vec4 <i>P</i> [, float <i>bias</i>])</p> <p>float texture (sampler2DArrayShadow <i>sampler</i>, vec4 <i>P</i>)</p> <p>gvec4 texture (gsamplerCubeArray <i>sampler</i>, vec4 <i>P</i> [, float <i>bias</i>])</p> <p>float texture (samplerCubeArrayShadow <i>sampler</i>, vec4 <i>P</i>, float <i>compare</i>)</p>	<p>Use the texture coordinate <i>P</i> to do a texture lookup in the texture currently bound to <i>sampler</i>.</p> <p>For shadow forms: When <i>compare</i> is present, it is used as D_{ref} and the array layer comes from the last component of <i>P</i>. When <i>compare</i> is not present, the last component of <i>P</i> is used as D_{ref} and the array layer comes from the second to last component of <i>P</i>.</p> <p>For array forms, the array layer comes from the last component of <i>P</i> in the non-shadow forms, and the second to last component of <i>P</i> in the shadow forms.</p>
<p>gvec4 textureProj (gsampler2D <i>sampler</i>, vec3 <i>P</i> [, float <i>bias</i>])</p> <p>gvec4 textureProj (gsampler2D <i>sampler</i>, vec4 <i>P</i> [, float <i>bias</i>])</p> <p>gvec4 textureProj (gsampler3D <i>sampler</i>, vec4 <i>P</i> [, float <i>bias</i>])</p> <p>float textureProj (sampler2DShadow <i>sampler</i>, vec4 <i>P</i> [, float <i>bias</i>])</p>	<p>Do a texture lookup with projection. The texture coordinates consumed from <i>P</i>, not including the last component of <i>P</i>, are divided by the last component of <i>P</i> to form projected coordinates P'. The resulting third component of P' in the shadow forms is used as D_{ref}. The third component of <i>P</i> is ignored when <i>sampler</i> has type gsampler2D and <i>P</i> has type vec4. After these values are computed, texture lookup proceeds as in texture.</p>

Syntax	Description
<p>gvec4 textureLod (gsampler2D <i>sampler</i>, vec2 <i>P</i>, float <i>lod</i>)</p> <p>gvec4 textureLod (gsampler3D <i>sampler</i>, vec3 <i>P</i>, float <i>lod</i>)</p> <p>gvec4 textureLod (gsamplerCube <i>sampler</i>, vec3 <i>P</i>, float <i>lod</i>)</p> <p>gvec4 textureLod (gsampler2DArray <i>sampler</i>, vec3 <i>P</i>, float <i>lod</i>)</p> <p>float textureLod (sampler2DShadow <i>sampler</i>, vec3 <i>P</i>, float <i>lod</i>)</p> <p>gvec4 textureLod (gsamplerCubeArray <i>sampler</i>, vec4 <i>P</i>, float <i>lod</i>)</p>	<p>Do a texture lookup as in texture but with explicit LOD; <i>lod</i> specifies λ_{base} and sets the partial derivatives as follows. (See section 8.13 “Texture Minification” in the OpenGL ES 3.2 Graphics System Specification.)</p> $\frac{\partial u}{\partial x} = 0 \quad \frac{\partial v}{\partial x} = 0 \quad \frac{\partial w}{\partial x} = 0$ $\frac{\partial u}{\partial y} = 0 \quad \frac{\partial v}{\partial y} = 0 \quad \frac{\partial w}{\partial y} = 0$
<p>gvec4 textureOffset (gsampler2D <i>sampler</i>, vec2 <i>P</i>, ivec2 <i>offset</i> [, float <i>bias</i>])</p> <p>gvec4 textureOffset (gsampler3D <i>sampler</i>, vec3 <i>P</i>, ivec3 <i>offset</i> [, float <i>bias</i>])</p> <p>float textureOffset (sampler2DShadow <i>sampler</i>, vec3 <i>P</i>, ivec2 <i>offset</i> [, float <i>bias</i>])</p> <p>gvec4 textureOffset (gsampler2DArray <i>sampler</i>, vec3 <i>P</i>, ivec2 <i>offset</i> [, float <i>bias</i>])</p>	<p>Do a texture lookup as in texture but with <i>offset</i> added to the (<i>u,v,w</i>) texel coordinates before looking up each texel. The offset value must be a constant expression. A limited range of offset values are supported; the minimum and maximum offset values are implementation-dependent and given by MIN_PROGRAM_TEXEL_OFFSET and MAX_PROGRAM_TEXEL_OFFSET, respectively.</p> <p>Note that <i>offset</i> does not apply to the layer coordinate for texture arrays. This is explained in detail in section 8.13 “Texture Minification” of the OpenGL ES Graphics System Specification, where <i>offset</i> is $(\delta_u, \delta_v, \delta_w)$. Note that texel offsets are also not supported for cube maps.</p>

Syntax	Description
gvec4 texelFetch (gsampler2D <i>sampler</i> , ivec2 <i>P</i> , int <i>lod</i>) gvec4 texelFetch (gsampler3D <i>sampler</i> , ivec3 <i>P</i> , int <i>lod</i>) gvec4 texelFetch (gsampler2DArray <i>sampler</i> , ivec3 <i>P</i> , int <i>lod</i>) gvec4 texelFetch (gsampler2DMS <i>sampler</i> , ivec2 <i>P</i> , int <i>sample</i>) gvec4 texelFetch (gsamplerBuffer <i>sampler</i> , int <i>P</i>)	Use integer texture coordinate <i>P</i> to lookup a single texel from <i>sampler</i> . The array layer comes from the last component of <i>P</i> for the array forms. The level-of-detail <i>lod</i> is as described in section 11.1.3.2 “Texel Fetches” of the OpenGL ES 3.2 Graphics System Specification.
gvec4 texelFetch (gsampler2DMSArray <i>sampler</i> , ivec3 <i>P</i> , int <i>sample</i>)	Use integer texture coordinate <i>P</i> to lookup a single sample <i>sample</i> on the texture bound to <i>sampler</i> as described in section 2.11.9.3 of the OpenGL ES specification “Multisample Texel Fetches”.
gvec4 texelFetchOffset (gsampler2D <i>sampler</i> , ivec2 <i>P</i> , int <i>lod</i> , ivec2 <i>offset</i>) gvec4 texelFetchOffset (gsampler3D <i>sampler</i> , ivec3 <i>P</i> , int <i>lod</i> , ivec3 <i>offset</i>) gvec4 texelFetchOffset (gsampler2DArray <i>sampler</i> , ivec3 <i>P</i> , int <i>lod</i> , ivec2 <i>offset</i>)	Fetch a single texel as in texelFetch offset by <i>offset</i> as described in textureOffset .
gvec4 textureProjOffset (gsampler2D <i>sampler</i> , vec3 <i>P</i> , ivec2 <i>offset</i> [, float <i>bias</i>]) gvec4 textureProjOffset (gsampler2D <i>sampler</i> , vec4 <i>P</i> , ivec2 <i>offset</i> [, float <i>bias</i>]) gvec4 textureProjOffset (gsampler3D <i>sampler</i> , vec4 <i>P</i> , ivec3 <i>offset</i> [, float <i>bias</i>]) float textureProjOffset (sampler2DShadow <i>sampler</i> , vec4 <i>P</i> , ivec2 <i>offset</i> [, float <i>bias</i>])	Do a projective texture lookup as described in textureProj offset by <i>offset</i> as described in textureOffset .

Syntax	Description
<p><code>gvec4 textureLodOffset</code> (<code>gsampler2D sampler</code>, <code>vec2 P</code>, <code>float lod</code>, <code>ivec2 offset</code>)</p> <p><code>gvec4 textureLodOffset</code> (<code>gsampler3D sampler</code>, <code>vec3 P</code>, <code>float lod</code>, <code>ivec3 offset</code>)</p> <p><code>gvec4 textureLodOffset</code> (<code>gsampler2DArray sampler</code>, <code>vec3 P</code>, <code>float lod</code>, <code>ivec2 offset</code>)</p> <p><code>float textureLodOffset</code> (<code>sampler2DShadow sampler</code>, <code>vec3 P</code>, <code>float lod</code>, <code>ivec2 offset</code>)</p>	Do an offset texture lookup with explicit LOD. See textureLod and textureOffset .
<p><code>gvec4 textureProjLod</code> (<code>gsampler2D sampler</code>, <code>vec3 P</code>, <code>float lod</code>)</p> <p><code>gvec4 textureProjLod</code> (<code>gsampler2D sampler</code>, <code>vec4 P</code>, <code>float lod</code>)</p> <p><code>gvec4 textureProjLod</code> (<code>gsampler3D sampler</code>, <code>vec4 P</code>, <code>float lod</code>)</p> <p><code>float textureProjLod</code> (<code>sampler2DShadow sampler</code>, <code>vec4 P</code>, <code>float lod</code>)</p>	Do a projective texture lookup with explicit LOD. See textureProj and textureLod .
<p><code>gvec4 textureProjLodOffset</code> (<code>gsampler2D sampler</code>, <code>vec3 P</code>, <code>float lod</code>, <code>ivec2 offset</code>)</p> <p><code>gvec4 textureProjLodOffset</code> (<code>gsampler2D sampler</code>, <code>vec4 P</code>, <code>float lod</code>, <code>ivec2 offset</code>)</p> <p><code>gvec4 textureProjLodOffset</code> (<code>gsampler3D sampler</code>, <code>vec4 P</code>, <code>float lod</code>, <code>ivec3 offset</code>)</p> <p><code>float textureProjLodOffset</code> (<code>sampler2DShadow sampler</code>, <code>vec4 P</code>, <code>float lod</code>, <code>ivec2 offset</code>)</p>	Do an offset projective texture lookup with explicit LOD. See textureProj , textureLod , and textureOffset .

Syntax	Description
<p>gvec4 textureGrad (gsampler2D <i>sampler</i>, vec2 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>)</p> <p>gvec4 textureGrad (gsampler3D <i>sampler</i>, vec3 <i>P</i>, vec3 <i>dPdx</i>, vec3 <i>dPdy</i>)</p> <p>gvec4 textureGrad (gsamplerCube <i>sampler</i>, vec3 <i>P</i>, vec3 <i>dPdx</i>, vec3 <i>dPdy</i>)</p> <p>gvec4 textureGrad (gsampler2DArray <i>sampler</i>, vec3 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>)</p> <p>float textureGrad (sampler2DShadow <i>sampler</i>, vec3 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>)</p> <p>float textureGrad (samplerCubeShadow <i>sampler</i>, vec4 <i>P</i>, vec3 <i>dPdx</i>, vec3 <i>dPdy</i>)</p> <p>float textureGrad (sampler2DArrayShadow <i>sampler</i>, vec4 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>)</p> <p>gvec4 textureGrad (gsamplerCubeArray <i>sampler</i>, vec4 <i>P</i>, vec3 <i>dPdx</i>, vec3 <i>dPdy</i>);</p>	<p>Do a texture lookup as in texture but with explicit gradients. The partial derivatives of <i>P</i> are with respect to window x and window y. Set</p> $\frac{\partial s}{\partial x} = \frac{\partial P.s}{\partial x}$ $\frac{\partial s}{\partial y} = \frac{\partial P.s}{\partial y}$ $\frac{\partial t}{\partial x} = \frac{\partial P.t}{\partial x}$ $\frac{\partial t}{\partial y} = \frac{\partial P.t}{\partial y}$ $\frac{\partial r}{\partial x} = \frac{\partial P.p}{\partial x} \text{ (cube)}$ $\frac{\partial r}{\partial y} = \frac{\partial P.p}{\partial y} \text{ (cube)}$ <p>For the cube version, the partial derivatives of <i>P</i> are assumed to be in the coordinate system used before texture coordinates are projected onto the appropriate cube face.</p>
<p>gvec4 textureGradOffset (gsampler2D <i>sampler</i>, vec2 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>, ivec2 <i>offset</i>)</p> <p>gvec4 textureGradOffset (gsampler3D <i>sampler</i>, vec3 <i>P</i>, vec3 <i>dPdx</i>, vec3 <i>dPdy</i>, ivec3 <i>offset</i>)</p> <p>gvec4 textureGradOffset (gsampler2DArray <i>sampler</i>, vec3 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>, ivec2 <i>offset</i>)</p> <p>float textureGradOffset (sampler2DShadow <i>sampler</i>, vec3 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>, ivec2 <i>offset</i>)</p> <p>float textureGradOffset (sampler2DArrayShadow <i>sampler</i>, vec4 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>, ivec2 <i>offset</i>)</p>	<p>Do a texture lookup with both explicit gradient and offset, as described in textureGrad and textureOffset.</p>

Syntax	Description
<p>gvec4 textureProjGrad (gsampler2D <i>sampler</i>, vec3 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>)</p> <p>gvec4 textureProjGrad (gsampler2D <i>sampler</i>, vec4 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>)</p> <p>gvec4 textureProjGrad (gsampler3D <i>sampler</i>, vec4 <i>P</i>, vec3 <i>dPdx</i>, vec3 <i>dPdy</i>)</p> <p>float textureProjGrad (sampler2DShadow <i>sampler</i>, vec4 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>)</p>	Do a texture lookup both projectively, as described in textureProj , and with explicit gradient as described in textureGrad . The partial derivatives <i>dPdx</i> and <i>dPdy</i> are assumed to be already projected.
<p>gvec4 textureProjGradOffset (gsampler2D <i>sampler</i>, vec3 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>, ivec2 <i>offset</i>)</p> <p>gvec4 textureProjGradOffset (gsampler2D <i>sampler</i>, vec4 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>, ivec2 <i>offset</i>)</p> <p>gvec4 textureProjGradOffset (gsampler3D <i>sampler</i>, vec4 <i>P</i>, vec3 <i>dPdx</i>, vec3 <i>dPdy</i>, ivec3 <i>offset</i>)</p> <p>float textureProjGradOffset (sampler2DShadow <i>sampler</i>, vec4 <i>P</i>, vec2 <i>dPdx</i>, vec2 <i>dPdy</i>, ivec2 <i>offset</i>)</p>	Do a texture lookup projectively and with explicit gradient as described in textureProjGrad , as well as with offset, as described in textureOffset .

8.9.3 Texture Gather Functions

The texture gather functions take components of a single floating-point vector operand as a texture coordinate, determine a set of four texels to sample from the base level of detail of the specified texture image, and return one component from each texel in a four-component result vector.

When performing a texture gather operation, the minification and magnification filters are ignored, and the rules for LINEAR filtering in the OpenGL ES Specification are applied to the base level of the texture image to identify the four texels i_{0j1} , i_{1j1} , i_{1j0} , and i_{0j0} . The texels are then converted to texture base colors (Rs, Gs, Bs, As) according to table 8.11 in section 8.5 “Texture Image Specification” of the OpenGL ES specification, followed by application of the texture swizzle as described section 14.2.1 “Texture Access” of the OpenGL ES Graphics System Specification. A four-component vector is assembled by taking the selected component from each of the post-swizzled texture source colors in the order (i_{0j1} , i_{1j1} , i_{1j0} , i_{0j0}).

The selected component is identified by the optional *comp* argument, where the values zero, one, two, and three identify the Rs, Gs, Bs, or As component, respectively. If *comp* is omitted, it is treated as identifying the Rs component.

Incomplete textures (section 8.16 “Texture Completeness” of the OpenGL ES specification) return a texture source color of (0,0,0,1) for all four source texels.

For texture gather functions using a shadow sampler type, each of the four texel lookups perform a depth comparison against the depth reference value passed in (*refZ*), and returns the result of that comparison in the appropriate component of the result vector.

As with other texture lookup functions, the results of a texture gather are undefined for shadow samplers if the texture referenced is not a depth texture or has depth comparisons disabled; or for non-shadow samplers if the texture referenced is a depth texture with depth comparisons enabled.

The **textureGatherOffset** built-in functions from the OpenGL ES Shading Language return a vector derived from sampling four texels in the image array of level *level_base*. For each of the four texel offsets specified by the *offsets* argument, the rules for the LINEAR minification filter are applied to identify a 2x2 texel footprint, from which the single texel T_{i0j0} is selected. A four-component vector is then assembled by taking a single component from each of the four T_{i0j0} texels in the same manner as for the **textureGather** function.

Syntax	Description
gvec4 textureGather (gsampler2D <i>sampler</i> , vec2 <i>P</i> [, int <i>comp</i>])	Returns the value
gvec4 textureGather (gsampler2DArray <i>sampler</i> , vec3 <i>P</i> [, int <i>comp</i>])	vec4(Sample_i0_j1(<i>P</i> , base). <i>comp</i> , Sample_i1_j1(<i>P</i> , base). <i>comp</i> , Sample_i1_j0(<i>P</i> , base). <i>comp</i> , Sample_i0_j0(<i>P</i> , base). <i>comp</i>)
gvec4 textureGather (gsamplerCube <i>sampler</i> , vec3 <i>P</i> [, int <i>comp</i>])	
vec4 textureGather (sampler2DShadow <i>sampler</i> , vec2 <i>P</i> , float <i>refZ</i>)	If specified, the value of <i>comp</i> must be a constant integer expression with a value of 0, 1, 2, or 3, identifying the <i>x</i> , <i>y</i> , <i>z</i> , or <i>w</i> post-swizzled component of the four-component vector lookup result for each texel, respectively. If <i>comp</i> is not specified, it is treated as 0, selecting the <i>x</i> component of each texel to generate the result.
vec4 textureGather (sampler2DArrayShadow <i>sampler</i> , vec3 <i>P</i> , float <i>refZ</i>)	
vec4 textureGather (samplerCubeShadow <i>sampler</i> , vec3 <i>P</i> , float <i>refZ</i>)	
gvec4 textureGather (gsamplerCubeArray <i>sampler</i> , vec4 <i>P</i> [, int <i>comp</i>])	
vec4 textureGather (samplerCubeArrayShadow <i>sampler</i> , vec4 <i>P</i> , float <i>refZ</i>)	

Syntax	Description
<pre>gvec4 textureGatherOffset (gsampler2D <i>sampler</i>, vec2 <i>P</i>, ivec2 <i>offset</i> [, int <i>comp</i>])</pre> <pre>gvec4 textureGatherOffset (gsampler2DArray <i>sampler</i>, vec3 <i>P</i>, ivec2 <i>offset</i> [, int <i>comp</i>])</pre> <pre>vec4 textureGatherOffset (sampler2DShadow <i>sampler</i>, vec2 <i>P</i>, float <i>refZ</i>, ivec2 <i>offset</i>)</pre> <pre>vec4 textureGatherOffset (sampler2DArrayShadow <i>sampler</i>, vec3 <i>P</i>, float <i>refZ</i>, ivec2 <i>offset</i>)</pre>	<p>Perform a texture gather operation as in textureGather offset by <i>offset</i> as described in textureOffset except that the <i>offset</i> can be variable (non-constant) and the implementation-dependent minimum and maximum offset values are given by MIN_PROGRAM_TEXTURE_GATHER_OFFSET and MAX_PROGRAM_TEXTURE_GATHER_OFFSET respectively.</p>
<pre>gvec4 textureGatherOffsets (gsampler2D <i>sampler</i>, vec2 <i>P</i>, ivec2 <i>offsets</i>[4] [, int <i>comp</i>])</pre> <pre>gvec4 textureGatherOffsets (gsampler2DArray <i>sampler</i>, vec3 <i>P</i>, ivec2 <i>offsets</i>[4] [, int <i>comp</i>])</pre> <pre>vec4 textureGatherOffsets (sampler2DShadow <i>sampler</i>, vec2 <i>P</i>, float <i>refZ</i>, ivec2 <i>offsets</i>[4])</pre> <pre>vec4 textureGatherOffsets (sampler2DArrayShadow <i>sampler</i>, vec3 <i>P</i>, float <i>refZ</i>, ivec2 <i>offsets</i>[4])</pre>	

Operate identically to **textureGatherOffset** except that *offsets* is used to determine the location of the four texels to sample. Each of the four texels is obtained by applying the corresponding offset in *offsets* as a (u,v) coordinate offset to *P*, identifying the four-texel linear footprint, and then selecting texel (i_0, j_0) of that footprint. The specified values in *offsets* must be constant integral expressions.

8.10 Atomic-Counter Functions

The atomic-counter operations in this section operate atomically with respect to each other. They are atomic for any single counter, meaning any of these operations on a specific counter in one shader instantiation will be indivisible by any of these operations on the same counter from another shader instantiation. There is no guarantee that these operations are atomic with respect to other forms of access to the counter or that they are serialized when applied to separate counters. Such cases would require additional use of fences, barriers, or other forms of synchronization, if atomicity or serialization is desired.

The value returned by an atomic-counter function is the value of an atomic counter, which may be

- returned and incremented in an atomic operation, or
- decremented and returned in an atomic operation, or
- simply returned.

The underlying counter is a 32-bit unsigned integer. Increments and decrements at the limit of the range will wrap to $[0, 2^{32}-1]$.

Syntax	Description
uint atomicCounterIncrement (atomic_uint <i>c</i>)	Atomically <ol style="list-style-type: none"> 1. increments the counter for <i>c</i>, and 2. returns its value prior to the increment operation. These two steps are done atomically with respect to the atomic counter functions in this table.
uint atomicCounterDecrement (atomic_uint <i>c</i>)	Atomically <ol style="list-style-type: none"> 1. decrements the counter for <i>c</i>, and 2. returns the value resulting from the decrement operation. These two steps are done atomically with respect to the atomic counter functions in this table.
uint atomicCounter (atomic_uint <i>c</i>)	Returns the counter value for <i>c</i> .

8.11 Atomic Memory Functions

Atomic memory functions perform atomic operations on an individual signed or unsigned integer stored in buffer-object or shared-variable storage. All of the atomic memory operations read a value from memory, compute a new value using one of the operations described below, write the new value to memory, and return the original value read. The contents of the memory being updated by the atomic operation are guaranteed not to be modified by any other assignment or atomic memory function in any shader invocation between the time the original value is read and the time the new value is written.

Atomic memory functions are supported only for a limited set of variables. A shader will fail to compile if the value passed to the *mem* argument of an atomic memory function does not correspond to a buffer or shared variable. It is acceptable to pass an element of an array or a single component of a vector to the *mem* argument of an atomic memory function, as long as the underlying array or vector is a buffer or shared variable.

All the built-in functions in this section accept arguments with combinations of **restrict**, **coherent**, and **volatile** memory qualification, despite not having them listed in the prototypes. The atomic operation will operate as required by the calling argument's memory qualification, not by the built-in function's formal parameter memory qualification.

Syntax	Description
uint atomicAdd (inout uint <i>mem</i> , uint <i>data</i>) int atomicAdd (inout int <i>mem</i> , int <i>data</i>)	Computes a new value by adding the value of <i>data</i> to the contents <i>mem</i> .
uint atomicMin (inout uint <i>mem</i> , uint <i>data</i>) int atomicMin (inout int <i>mem</i> , int <i>data</i>)	Computes a new value by taking the minimum of the value of <i>data</i> and the contents of <i>mem</i> .
uint atomicMax (inout uint <i>mem</i> , uint <i>data</i>) int atomicMax (inout int <i>mem</i> , int <i>data</i>)	Computes a new value by taking the maximum of the value of <i>data</i> and the contents of <i>mem</i> .
uint atomicAnd (inout uint <i>mem</i> , uint <i>data</i>) int atomicAnd (inout int <i>mem</i> , int <i>data</i>)	Computes a new value by performing a bit-wise AND of the value of <i>data</i> and the contents of <i>mem</i> .
uint atomicOr (inout uint <i>mem</i> , uint <i>data</i>) int atomicOr (inout int <i>mem</i> , int <i>data</i>)	Computes a new value by performing a bit-wise OR of the value of <i>data</i> and the contents of <i>mem</i> .
uint atomicXor (inout uint <i>mem</i> , uint <i>data</i>) int atomicXor (inout int <i>mem</i> , int <i>data</i>)	Computes a new value by performing a bit-wise EXCLUSIVE OR of the value of <i>data</i> and the contents of <i>mem</i> .
uint atomicExchange (inout uint <i>mem</i> , uint <i>data</i>) int atomicExchange (inout int <i>mem</i> , int <i>data</i>)	Computes a new value by simply copying the value of <i>data</i> .
uint atomicCompSwap (inout uint <i>mem</i> , uint <i>compare</i> , uint <i>data</i>) int atomicCompSwap (inout int <i>mem</i> , int <i>compare</i> , int <i>data</i>)	Compares the value of <i>compare</i> and the contents of <i>mem</i> . If the values are equal, the new value is given by <i>data</i> ; otherwise, it is taken from the original contents of <i>mem</i> .

8.12 Image Functions

Variables using one of the image basic types may be used by the built-in shader image memory functions defined in this section to read and write individual texels of a texture. Each image variable references an image unit, which has a texture image attached.

When image memory functions below access memory, an individual texel in the image is identified using an (i) , (i, j) , or (i, j, k) coordinate corresponding to the values of P . The coordinates are used to select an individual texel in the manner described in section 8.22 “Texture Image Loads and Stores” of the OpenGL ES specification.

Loads and stores support float, integer, and unsigned integer types. The data types below starting “*gimage*” serve as placeholders meaning types starting either “**image**”, “**iimage**”, or “**uimage**” in the same way as *gvec* or *gsampler* in earlier sections.

The *IMAGE_PARAMS* in the prototypes below is a placeholder representing 18 separate functions, each for a different type of image variable. The *IMAGE_PARAMS* placeholder is replaced by one of the following parameter lists:

gimage2D image, ivec2 P

gimage3D image, ivec3 P

gimageCube image, ivec3 P

gimage2DArray image, ivec3 P

gimageBuffer image, int P

gimageCubeArray image, ivec3 P

where each of the lines represents one of three different image variable types, and *image*, *P* specify the individual texel to operate on. The method for identifying the individual texel operated on from *image*, *P*, and the method for reading and writing the texel are specified in section 8.22 “Texture Image Loads and Stores” of the OpenGL ES specification.

The atomic functions perform operations on individual texels or samples of an image variable. Atomic memory operations read a value from the selected texel, compute a new value using one of the operations described below, write the new value to the selected texel, and return the original value read. The contents of the texel being updated by the atomic operation are guaranteed not to be updated by any other image store or atomic function between the time the original value is read and the time the new value is written.

Atomic memory operations are supported on only a subset of all image variable types; *image* must be either:

- a signed integer image variable (type starts “**iimage**”) and a format qualifier of **r32i**, used with a *data* argument of type **int**, or
- an unsigned integer image variable (type starts “**uimage**”) and a format qualifier of **r32ui**, used with a *data* argument of type **uint**, or
- a float image variable (type starts “**image**”) and a format qualifier of **r32f**, used with a *data* argument of type **float** (imageAtomicExchange only).

All the built-in functions in this section accept arguments with combinations of **restrict**, **coherent**, and **volatile** memory qualification, despite not having them listed in the prototypes. The image operation will operate as required by the calling argument's memory qualification, not by the built-in function's formal parameter memory qualification.

Syntax	Description
highp ivec2 imageSize (readonly writeonly gimage2D image) highp ivec3 imageSize (readonly writeonly gimage3D image) highp ivec2 imageSize (readonly writeonly gimageCube image) highp ivec3 imageSize (readonly writeonly gimage2DArray image) highp int imageSize (readonly writeonly gimageBuffer image) highp ivec3 imageSize (readonly writeonly gimageCubeArray image)	<p>Returns the dimensions of the image or images bound to <i>image</i>. For arrayed images, the last component of the return value will hold the size of the array. Cube images only return the dimensions of one face.</p> <p>Note: The qualification readonly writeonly accepts a variable qualified with readonly, writeonly, both, or neither. It means the formal argument will be used for neither reading nor writing to the underlying memory.</p>
highp gvec4 imageLoad (readonly <i>IMAGE_PARAMS</i>)	<p>Loads the texel at the coordinate <i>P</i> from the image unit <i>image</i> (in <i>IMAGE_PARAMS</i>). When <i>image</i>, <i>P</i> identify a valid texel, the bits used to represent the selected texel in memory are converted to a vec4, ivec4, or uvec4 in the manner described in section 8.22 “Texture Image Loads and Stores” of the OpenGL ES Specification and returned.</p>
void imageStore (writeonly <i>IMAGE_PARAMS</i> , gvec4 <i>data</i>)	<p>Stores <i>data</i> into the texel at the coordinate <i>P</i> from the image specified by <i>image</i>. When <i>image</i> and <i>P</i> identify a valid texel, the bits used to represent <i>data</i> are converted to the format of the image unit in the manner described in section 8.22 “Texture Image Loads and Stores” of the OpenGL ES Specification and stored to the specified texel.</p>
highp uint imageAtomicAdd (<i>IMAGE_PARAMS</i> , uint <i>data</i>); highp int imageAtomicAdd (<i>IMAGE_PARAMS</i> , int <i>data</i>);	<p>Computes a new value by adding the value of <i>data</i> to the contents of the selected texel.</p>
highp uint imageAtomicMin (<i>IMAGE_PARAMS</i> , uint <i>data</i>); highp int imageAtomicMin (<i>IMAGE_PARAMS</i> , int <i>data</i>);	<p>Computes a new value by taking the minimum of the value of <i>data</i> and the contents of the selected texel.</p>

Syntax	Description
highp uint imageAtomicMax (IMAGE_PARAMS, uint data); highp int imageAtomicMax (IMAGE_PARAMS, int data);	Computes a new value by taking the maximum of the value of <i>data</i> and the contents of the selected texel.
highp uint imageAtomicAnd (IMAGE_PARAMS, uint data); highp int imageAtomicAnd (IMAGE_PARAMS, int data);	Computes a new value by performing a bitwise AND of the value of <i>data</i> and the contents of the selected texel
highp uint imageAtomicOr (IMAGE_PARAMS, uint data); highp int imageAtomicOr (IMAGE_PARAMS, int data);	Computes a new value by performing a bitwise OR of the value of <i>data</i> and the contents of the selected texel
highp uint imageAtomicXor (IMAGE_PARAMS, uint data); highp int imageAtomicXor (IMAGE_PARAMS, int data);	Computes a new value by performing a bitwise EXCLUSIVE OR of the value of <i>data</i> and the contents of the selected texel
highp uint imageAtomicExchange (IMAGE_PARAMS, uint data); highp int imageAtomicExchange (IMAGE_PARAMS, int data); highp float imageAtomicExchange (IMAGE_PARAMS, int float);	Computes a new value by simply copying the value of <i>data</i> .
highp uint imageAtomicCompSwap (IMAGE_PARAMS, uint data); highp int imageAtomicCompSwap (IMAGE_PARAMS, int data);	Compares the value of <i>compare</i> and the contents of the selected texel. If the values are equal, the new value is given by <i>data</i> ; otherwise, it is taken from the original value loaded from the texel.

8.13 Geometry Shader Functions

These functions are only available in geometry shaders. They are described in more depth following the table.

Syntax	Description
void EmitVertex ()	Emits the current values of output variables to the current output primitive. On return from this call, the values of output variables are undefined.
void EndPrimitive ()	Completes the current output primitive and starts a new one. No vertex is emitted.

The function `EmitVertex()` specifies that a vertex is completed. A vertex is added to the current output primitive using the current values of all built-in and user-defined output variables. The values of all output variables are undefined after a call to `EmitVertex()`. If a geometry shader invocation has emitted more vertices than permitted by the output layout qualifier `max_vertices`, the results of calling `EmitVertex()` are undefined.

The function `EndPrimitive()` specifies that the current output primitive is completed and a new output primitive (of the same type) will be started by any subsequent `EmitVertex()`. This function does not emit a vertex. If the output layout is declared to be **points**, calling `EndPrimitive()` is optional.

A geometry shader starts with an output primitive containing no vertices. When a geometry shader terminates, the current output primitive is automatically completed. It is not necessary to call `EndPrimitive()` if the geometry shader writes only a single primitive.

8.14 Fragment Processing Functions

Fragment processing functions are only available in fragment shaders.

Derivatives may be computationally expensive and/or numerically unstable. Therefore, an OpenGL ES implementation may approximate the true derivatives by using a fast but not entirely accurate derivative computation. Derivatives are undefined within non-uniform control flow.

The expected behavior of a derivative is specified using forward/backward differencing.

Forward differencing:

$$F(x+dx) - F(x) \sim dFdx(x) \cdot dx \quad 1a$$

$$dFdx(x) \sim \frac{F(x+dx) - F(x)}{dx} \quad 1b$$

Backward differencing:

$$F(x-dx) - F(x) \sim -dFdx(x) \cdot dx \quad 2a$$

$$dFdx(x) \sim \frac{F(x) - F(x-dx)}{dx} \quad 2b$$

With single-sample rasterization, $dx \leq 1.0$ in equations 1b and 2b. For multi-sample rasterization, $dx < 2.0$ in equations 1b and 2b.

dFdy is approximated similarly, with y replacing x .

An OpenGL ES implementation may use the above or other methods to perform the calculation, subject to the following conditions:

1. The method may use piecewise linear approximations. Such linear approximations imply that higher order derivatives, **dFdx(dFdx(x))** and above, are undefined.
2. The method may assume that the function evaluated is continuous. Therefore derivatives within the body of a non-uniform conditional are undefined.
3. The method may differ per fragment, subject to the constraint that the method may vary by window coordinates, not screen coordinates. The invariance requirement described in section 13.2 “Invariance” of the OpenGL ES Graphics System Specification, is relaxed for derivative calculations, because the method may be a function of fragment location.

Other properties that are desirable, but not required, are:

1. Functions should be evaluated within the interior of a primitive (interpolated, not extrapolated).
2. Functions for **dFdx** should be evaluated while holding y constant. Functions for **dFdy** should be evaluated while holding x constant. However, mixed higher order derivatives, like **dFdx(dFdy(y))** and **dFdy(dFdx(x))** are undefined.
3. Derivatives of constant arguments should be 0.

In some implementations, varying degrees of derivative accuracy may be obtained by providing GL hints (section 18.1 “Hints” of the OpenGL ES 3.0 Graphics System Specification), allowing a user to make an image quality versus speed trade off.

Syntax	Description
genFType dFdx (genFType <i>p</i>)	Returns the derivative in x using local differencing for the input argument <i>p</i> .
genFType dFdy (genFType <i>p</i>)	<p>Returns the derivative in y using local differencing for the input argument <i>p</i>.</p> <p>These two functions are commonly used to estimate the filter width used to anti-alias procedural textures. We are assuming that the expression is being evaluated in parallel on a SIMD array so that at any given point in time the value of the function is known at the grid points represented by the SIMD array. Local differencing between SIMD array elements can therefore be used to derive dFdx, dFdy, etc.</p>
genFType fwidth (genFType <i>p</i>)	Returns the sum of the absolute derivative in x and y using local differencing for the input argument <i>p</i> , i.e., abs (dFdx (<i>p</i>)) + abs (dFdy (<i>p</i>)) ;

8.14.1 Interpolation Functions

Built-in interpolation functions are available to compute an interpolated value of a fragment shader input variable at a shader-specified (x,y) location. A separate (x,y) location may be used for each invocation of the built-in function, and those locations may differ from the default (x,y) location used to produce the default value of the input. For the interpolateAt* functions, the call will return a precision qualification matching the precision of the *interpolant* argument to the function call.

For all of the interpolation functions, *interpolant* must be an input variable or an element of an input variable declared as an array. Component selection operators (e.g., **.xy**) may not be used when specifying *interpolant*. If *interpolant* is declared with a **flat** qualifier, the interpolated value will have the same value everywhere for a single primitive, so the location used for the interpolation has no effect and the functions just return that same value. If *interpolant* is declared with the **centroid** qualifier, the value returned by interpolateAtSample() and interpolateAtOffset() will be evaluated at the specified location, ignoring the location normally used with the **centroid** qualifier.

Syntax	Description
float interpolateAtCentroid (float interpolant) vec2 interpolateAtCentroid (vec2 interpolant) vec3 interpolateAtCentroid (vec3 interpolant) vec4 interpolateAtCentroid (vec4 interpolant)	Returns the value of the input <i>interpolant</i> sampled at a location inside the both the pixel and the primitive being processed. The value obtained would be the same value assigned to the input variable if declared with the centroid qualifier.
float interpolateAtSample (float interpolant, int sample) vec2 interpolateAtSample (vec2 interpolant, int sample) vec3 interpolateAtSample (vec3 interpolant, int sample) vec4 interpolateAtSample (vec4 interpolant, int sample)	Returns the value of the input <i>interpolant</i> variable at the location of the sample numbered <i>sample</i> . If multisample buffers are not available, the input varying will be evaluated at the center of the pixel. If sample <i>sample</i> does not exist, the position used to interpolate the input varying is undefined.
float interpolateAtOffset (float interpolant, vec2 offset) vec2 interpolateAtOffset (vec2 interpolant, vec2 offset) vec3 interpolateAtOffset (vec3 interpolant, vec2 offset) vec4 interpolateAtOffset (vec4 interpolant, vec2 offset);	Returns the value of the input <i>interpolant</i> variable sampled at an offset from the center of the pixel specified by <i>offset</i> . The two floating-point components of <i>offset</i> give the offset in pixels in the x and y directions, respectively. An offset of (0,0) identifies the center of the pixel. The range and granularity of offsets supported by this function is implementation-dependent.

8.15 Shader Invocation Control Functions

The shader invocation control function is only available in tessellation control and compute shaders. It is used to control the relative execution order of multiple shader invocations used to process a local work group (in the case of compute shaders), which are otherwise executed with an undefined order.

Syntax	Description
void barrier ()	For any given static instance of barrier() , all tessellation control shader invocations for a single input patch, or all compute shader invocations for a single work group must enter it before any will continue beyond it.

The function **barrier()** provides a partially defined order of execution between shader invocations. For tessellation control shaders but not compute shaders, it also defines an order of memory operations.

For tessellation control shaders, the `barrier` ensures that values written by one invocation prior to a given static instance of `barrier()` can be safely read by other invocations after their call to the same static instance `barrier()`. Because invocations may execute in an undefined order between these barrier calls, the values of a per-vertex or per-patch output variable for tessellation control shaders, or the values of shared variables for compute shaders will be undefined in a number of cases enumerated in section 4.3.6 (“Output Variables”) for tessellation control shaders and section 4.3.8 (“Shared Variables”) for compute shaders.

For tessellation control shaders, the `barrier()` function may only be placed inside the function `main()` of the shader and may not be called within any control flow. Barriers are also disallowed after a return statement in the function `main()`. Any such misplaced barriers result in a compile-time error.

For compute shaders, a barrier only affects control flow and does not by itself synchronize memory accesses. In particular, it does not ensure that values written by one invocation prior to a given static instance of **`barrier()`** can be safely read by other invocations after their call to the same static instance of **`barrier()`**. To achieve this requires the use of both **`barrier()`** and a memory barrier. Because invocations may execute in an undefined order between these barrier calls, the values of shared variables for compute shaders will be undefined in a number of cases enumerated in Section 4.3.8 “Shared Variables” (for compute shaders).

For compute shaders, the **`barrier()`** function may be placed within control flow, but that control flow must be uniform control flow. That is, all the controlling expressions that lead to execution of the barrier must be dynamically uniform expressions. This ensures that if any shader invocation enters a conditional statement, then all invocations will enter it. While compilers are encouraged to give warnings if they can detect this might not happen, compilers cannot completely determine this. Hence, it is the author's responsibility to ensure **`barrier()`** only exists inside uniform control flow. Otherwise, some shader invocations will stall indefinitely, waiting for a barrier that is never reached by other invocations.

8.16 Shader Memory Control Functions

Within a single shader invocation, the visibility and order of writes made by that invocation are well-defined. However, the relative order of reads and writes to a single shared memory address from multiple separate shader invocations is largely undefined. Additionally, the order of accesses to multiple memory addresses performed by a single shader invocation, as observed by other shader invocations, is also undefined.

The following built-in functions can be used to control the ordering of reads and writes:

Syntax	Description
void memoryBarrier ()	Control the ordering of all memory transactions issued by a single shader invocation.
void memoryBarrierAtomicCounter ()	Control the ordering of accesses to atomic counter variables issued by a single shader invocation.
void memoryBarrierBuffer ()	Control the ordering of memory transactions to buffer variables issued within a single shader invocation.
void memoryBarrierImage ()	Control the ordering of memory transactions to images issued within a single shader invocation.
void memoryBarrierShared ()	Control the ordering of memory transactions to shared variables issued within a single shader invocation, as viewed by other invocations in the same work group. Only available in compute shaders.
void groupMemoryBarrier ()	Control the ordering of all memory transactions issued within a single shader invocation, as viewed by other invocations in the same work group. Only available in compute shaders.

The memory barrier built-in functions can be used to order reads and writes to variables stored in memory accessible to other shader invocations. When called, these functions will wait until prior memory transactions reach a point where they cannot be passed by subsequent transactions, as observed by other shader invocations and then return with no other effect. The built-in functions **memoryBarrierAtomicCounter()**, **memoryBarrierBuffer()**, **memoryBarrierImage()**, and **memoryBarrierShared()** wait for the completion of accesses to atomic counter, buffer, image, and shared variables, respectively. The built-in functions **memoryBarrier()** and **groupMemoryBarrier()** wait for the completion of accesses to all of the above variable types. The functions **memoryBarrierShared()** and **groupMemoryBarrier()** are available only in compute shaders; the other functions are available in all shader types.

When these functions return, any memory stores performed using coherent variables prior to the call will be visible to any future coherent access to the same memory performed by any other shader invocation. In particular, the values written this way in one shader stage are guaranteed to be visible to coherent memory accesses performed by shader invocations in subsequent stages when those invocations were triggered by the execution of the original shader invocation (e.g., fragment shader invocations for a primitive resulting from a particular vertex or geometry shader invocation).

Additionally, memory barrier functions order stores performed by the calling invocation, as observed by other shader invocations. Without memory barriers, if one shader invocation performs two stores to coherent variables, a second shader invocation might see the values written by the second store prior to seeing those written by the first. However, if the first shader invocation calls a memory barrier function between the two stores, selected other shader invocations will never see the results of the second store before seeing those of the first. When using the function **groupMemoryBarrier()**, this ordering guarantee applies only to other shader invocations in the same compute shader work group; all other memory barrier functions provide the guarantee to all other shader invocations. No memory barrier is required to guarantee the order of memory stores as observed by the invocation performing the stores; an invocation reading from a variable that it previously wrote will always see the most recently written value unless another shader invocation also wrote to the same memory.

9 Shader Interface Matching

As described in chapter 7 of the OpenGL ES specification, shaders may be linked together to form a *program object* before being bound to the pipeline or may be linked and bound individually as *separable program objects*¹.

Within a *program object* or a *separable program object*, qualifiers for matching variables must themselves match according to the rules specified in this section. There are also matching rules for qualifiers of matching variables between separable program objects but only for variables across an input/output boundary between shader stages. For other shader interface variables such as uniforms, each program object or separable program object has its own name space and so the same name can refer to multiple independent variables. Consequently, there are no matching rules for qualifiers in these cases.

9.1 Input Output Matching by Name in Linked Programs

When linking shaders, the type of declared vertex outputs and fragment inputs with the same name must match, otherwise the link command will fail. Only those fragment inputs statically used (i.e. read) in the fragment shader must be declared as outputs in the vertex shader; declaring superfluous vertex shader outputs is permissible.

The following table summarizes the rules for matching shader outputs to shader inputs in consecutive stages when shaders are linked together.

Treatment of Mismatched Input Variables		Consuming Shader (input variables)		
		No declaration	Declared but no static use	Declared and static use
Generating Shader (output variables)	No declaration	Allowed	Allowed	error
	Declares; no static use	Allowed	Allowed	Allowed (values are undefined)
	Declares and static use	Allowed	Allowed	Allowed (values are potentially undefined)

See section 3.9.1 for the definition of *static use*.

¹ These were previously known as separate shader objects (SSOs) but the mechanism has been extended to support future versions of the specification that have more than two shader stages. It allows a subset of the shaders to be linked together.

The precision of a vertex output does not need to match the precision of the corresponding fragment input. The minimum precision at which vertex outputs are interpolated is the minimum of the vertex output precision and the fragment input precision, with the exception that for highp, implementations do not have to support full IEEE 754 precision.

The precision of values exported to a transform feedback buffer is the precision of the outputs of the vertex shader. However, they are converted to highp format before being written.

9.2 Matching of Qualifiers

The following tables summarize the requirements for matching of qualifiers. It applies whenever there are two or more matching variables in a shader interface.

Notes:

1. *Yes* means the qualifiers must match.
2. *No* means the qualifiers do not need to match.
3. *Consistent* means qualifiers may be missing from a subset of declarations but they cannot conflict
4. The rules apply to all declared variables, irrespective of whether they are statically used, with the exception of inputs and outputs when shaders are linked (see section 9.1 “Input Output Matching by Name in Linked Programs”)
5. Errors are generated for any conflicts.

9.2.1 Linked Shaders

Qualifier Class	Qualifier	in/out	Default Uniforms	uniform Block	buffer Block
Storage ¹	in out uniform	N/A	N/A	N/A	N/A
Auxiliary	centroid sample	No	N/A	N/A	N/A
	patch	Yes	N/A	N/A	N/A
Layout	location	Yes ²	Consistent	N/A	N/A
	Block layout	N/A	N/A	Yes	Yes
	binding	N/A	Yes	Yes	Yes
	offset	N/A	Yes	N/A	N/A
	format	N/A	Yes	N/A	N/A
Interpolation	smooth flat	Yes	N/A	N/A	N/A
Precision	lowp mediump highp	No	Yes	Yes	Yes
Variance	invariant precise	No	N/A	N/A	N/A
Memory	all	N/A	Yes	Yes	Yes

-
- Storage qualifiers determine *when* variables match rather than being *required* to match for matching variables. Note also that each shader interface has a separate name space so for example, it is possible to use the same name for a vertex output and fragment uniform.
 - If present, the location qualifier determines the matching of inputs and outputs. See section 7.4.1. (“Shader interface Matching”) of the OpenGL ES 3.1 specification for details.

9.2.2 Separable Programs

Qualifier Class	Qualifier	in/out
Storage	in out uniform	N/A
Auxiliary	centroid sample	No
	patch	Yes
Layout	location	Yes
	Block layout	N/A
	binding	N/A
	offset	N/A
	format	N/A
Interpolation	smooth flat	Yes
Precision	lowp mediump highp	Yes
Variance	invariant precise	No
Memory	all	N/A

10 Shading Language Grammar

The grammar is fed from the output of lexical analysis. The tokens returned from lexical analysis are

```

VOID
BOOL FLOAT INT UINT
BVEC2 BVEC3 BVEC4 IVEC2 IVEC3 IVEC4 UVEC2 UVEC3 UVEC4 VEC2 VEC3 VEC4
MAT2 MAT3 MAT4
MAT2X2 MAT2X3 MAT2X4
MAT3X2 MAT3X3 MAT3X4
MAT4X2 MAT4X3 MAT4X4
STRUCT

ATOMIC_UINT
SAMPLER2D SAMPLER3D SAMPLERCUBE SAMPLER2DARRAY
SAMPLER2DSHADOW SAMPLERCUBESHADOW SAMPLER2DARRAYSHADOW
ISAMPLER2D ISAMPLER3D ISAMPLERCUBE ISAMPLER2DARRAY
USAMPLER2D USAMPLER3D USAMPLERCUBE USAMPLER2DARRAY

SAMPLER2DMS ISAMPLER2DMS USAMPLER2DMS
SAMPLER2DMSARRAY ISAMPLER2DMSARRAY USAMPLER2DMSARRAY
SAMPLERBUFFER ISAMPLERBUFFER USAMPLERBUFFER
SAMPLER2DMSARRAY ISAMPLER2DMSARRAY USAMPLER2DMSARRAY
SAMPLERCUBEARRAY ISAMPLERCUBEARRAY USAMPLERCUBEARRAY
SAMPLERCUBEARRAYSHADOW

IMAGE2D IIMAGE2D UIMAGE2D
IMAGE3D IIMAGE3D UIMAGE3D
IMAGECUBE IIMAGECUBE UIMAGECUBE
IMAGE2DARRAY IIMAGE2DARRAY UIMAGE2DARRAY
IMAGEBUFFER IIMAGEBUFFER UIMAGEBUFFER
UIMAGECUBEARRAY IMAGECUBEARRAY IIMAGECUBEARRAY

INVARIANT
HIGH_PRECISION MEDIUM_PRECISION LOW_PRECISION PRECISION
IN OUT INOUT
CONST UNIFORM BUFFER SHARED
COHERENT VOLATILE RESTRICT READONLY WRITEONLY
FLAT SMOOTH CENTROID

LAYOUT

PATCH    SAMPLE

WHILE BREAK CONTINUE DO ELSE FOR IF DISCARD RETURN SWITCH CASE DEFAULT

```


10 Shading Language Grammar

```
IDENTIFIER TYPE_NAME FLOATCONSTANT INTCONSTANT UINTCONSTANT BOOLCONSTANT  
FIELD_SELECTION  
LEFT_OP RIGHT_OP  
INC_OP DEC_OP LE_OP GE_OP EQ_OP NE_OP  
AND_OP OR_OP XOR_OP MUL_ASSIGN DIV_ASSIGN ADD_ASSIGN  
MOD_ASSIGN LEFT_ASSIGN RIGHT_ASSIGN AND_ASSIGN XOR_ASSIGN OR_ASSIGN  
SUB_ASSIGN
```

```
LEFT_PAREN RIGHT_PAREN LEFT_BRACKET RIGHT_BRACKET LEFT_BRACE RIGHT_BRACE DOT  
COMMA COLON EQUAL SEMICOLON BANG DASH TILDE PLUS STAR SLASH PERCENT  
LEFT_ANGLE RIGHT_ANGLE VERTICAL_BAR CARET AMPERSAND QUESTION
```

The following describes the grammar for the OpenGL ES Shading Language in terms of the above tokens.

```

variable_identifier:
    IDENTIFIER

primary_expression:
    variable_identifier
    INTCONSTANT
    UINTCONSTANT
    FLOATCONSTANT
    BOOLCONSTANT
    LEFT_PAREN expression RIGHT_PAREN

postfix_expression:
    primary_expression
    postfix_expression LEFT_BRACKET integer_expression RIGHT_BRACKET
    function_call
    postfix_expression DOT FIELD_SELECTION
    postfix_expression INC_OP
    postfix_expression DEC_OP

// FIELD_SELECTION includes fields in structures, component selection for vectors
// and the 'length' identifier for the length() method

integer_expression:
    expression

function_call:
    function_call_or_method

function_call_or_method:
    function_call_generic

function_call_generic:
    function_call_header_with_parameters RIGHT_PAREN
    function_call_header_no_parameters RIGHT_PAREN

function_call_header_no_parameters:
    function_call_header VOID
    function_call_header

function_call_header_with_parameters:
    function_call_header assignment_expression
    function_call_header_with_parameters COMMA assignment_expression

function_call_header:
    function_identifier LEFT_PAREN

```

// Grammar Note: Constructors look like functions, but lexical analysis recognized most of them as
 // keywords. They are now recognized through “type_specifier”.

// Methods (.length) and identifiers are recognized through postfix_expression.

function_identifier:
 type_specifier
 postfix_expression

unary_expression:
 postfix_expression
 INC_OP *unary_expression*
 DEC_OP *unary_expression*
 unary_operator unary_expression

// Grammar Note: No traditional style type casts.

unary_operator:
 PLUS
 DASH
 BANG
 TILDE

// Grammar Note: No '*' or '&' unary ops. Pointers are not supported.

multiplicative_expression:
 unary_expression
 multiplicative_expression STAR *unary_expression*
 multiplicative_expression SLASH *unary_expression*
 multiplicative_expression PERCENT *unary_expression*

additive_expression:
 multiplicative_expression
 additive_expression PLUS *multiplicative_expression*
 additive_expression DASH *multiplicative_expression*

shift_expression:
 additive_expression
 shift_expression LEFT_OP *additive_expression*
 shift_expression RIGHT_OP *additive_expression*

relational_expression:
 shift_expression
 relational_expression LEFT_ANGLE *shift_expression*
 relational_expression RIGHT_ANGLE *shift_expression*
 relational_expression LE_OP *shift_expression*
 relational_expression GE_OP *shift_expression*

equality_expression:

relational_expression

equality_expression EQ_OP *relational_expression*

equality_expression NE_OP *relational_expression*

and_expression:

equality_expression

and_expression AMPERSAND *equality_expression*

exclusive_or_expression:

and_expression

exclusive_or_expression CARET *and_expression*

inclusive_or_expression:

exclusive_or_expression

inclusive_or_expression VERTICAL_BAR *exclusive_or_expression*

logical_and_expression:

inclusive_or_expression

logical_and_expression AND_OP *inclusive_or_expression*

logical_xor_expression:

logical_and_expression

logical_xor_expression XOR_OP *logical_and_expression*

logical_or_expression:

logical_xor_expression

logical_or_expression OR_OP *logical_xor_expression*

conditional_expression:

logical_or_expression

logical_or_expression QUESTION *expression* COLON *assignment_expression*

assignment_expression:

conditional_expression

unary_expression *assignment_operator* *assignment_expression*

assignment_operator:

EQUAL

MUL_ASSIGN

DIV_ASSIGN

MOD_ASSIGN

ADD_ASSIGN

SUB_ASSIGN

LEFT_ASSIGN

RIGHT_ASSIGN

AND_ASSIGN

XOR_ASSIGN

OR_ASSIGN

expression:

$$assignment_expression$$
$$\text{expression COMMA assignment_expression}$$

constant_expression:

$$conditional_expression$$

declaration:

function_prototype SEMICOLON

```
init_declarator_list SEMICOLON
```

```
PRECISION precision_qualifier type_specifier SEMICOLON
```

```

type_qualifier IDENTIFIER LEFT_BRACE struct_declaration_list RIGHT_BRACE
                                     SEMICOLON

```

```

type_qualifier IDENTIFIER LEFT_BRACE struct_declaration_list RIGHT_BRACE
                                IDENTIFIER SEMICOLON

```

```
type_qualifier IDENTIFIER LEFT_BRACE struct_declaration_list RIGHT_BRACE  
IDENTIFIER array_specifier SEMICOLON
```

type_qualifier SEMICOLON

type_qualifier IDENTIFIER SEMICOLON // e.g. to qualify an existing variable as invariant

$$type_qualifier\ IDENTIFIER\ identifier_list\ SEMICOLON$$

identifier_list:

COMMA IDENTIFIER

identifiser_list COMMA IDENTIFIER

function prototype:

```
function_declarator RIGHT_PAREN
```

function_declarator:

function_header

function_header_with_parameters

function_header_with_parameters:

function_header parameter_declaration

function header with parameters COMMA *parameter declaration*

function_header:

fully_specified_type IDENTIFIER LEFT_PAREN

parameter_declarator:

type_specifier IDENTIFIER

type_specifier IDENTIFIER *array_specifier*

parameter_declaration:

type_qualifier *parameter_declarator*

parameter_declarator

type_qualifier *parameter_type_specifier*

parameter_type_specifier

parameter type specifier:

type_specifier

```

init_declarator_list:
    single_declaration
    init_declarator_list COMMA IDENTIFIER
    init_declarator_list COMMA IDENTIFIER array_specifier
    init_declarator_list COMMA IDENTIFIER array_specifier EQUAL initializer
    init_declarator_list COMMA IDENTIFIER EQUAL initializer

single_declaration:
    fully_specified_type
    fully_specified_type IDENTIFIER
    fully_specified_type IDENTIFIER array_specifier
    fully_specified_type IDENTIFIER array_specifier EQUAL initializer
    fully_specified_type IDENTIFIER EQUAL initializer

// Grammar Note: No 'enum', or 'typedef'.

fully_specified_type:
    type_specifier
    type_qualifier type_specifier

invariant_qualifier:
    INVARIANT

interpolation_qualifier:
    SMOOTH
    FLAT

layout_qualifier:
    LAYOUT LEFT_PAREN layout_qualifier_id_list RIGHT_PAREN

layout_qualifier_id_list:
    layout_qualifier_id
    layout_qualifier_id_list COMMA layout_qualifier_id

layout_qualifier_id:
    IDENTIFIER
    IDENTIFIER EQUAL INTCONSTANT
    IDENTIFIER EQUAL UINTCONSTANT
    SHARED

precise_qualifier:
    PRECISE

type_qualifier:
    single_type_qualifier
    type_qualifier single_type_qualifier

```

single_type_qualifier:

storage_qualifier
layout_qualifier
precision_qualifier
interpolation_qualifier
invariant_qualifier
precise_qualifier

storage_qualifier:

CONST
 IN
 OUT
 INOUT
 CENTROID
 PATCH
 SAMPLE
 UNIFORM
 BUFFER
 SHARED
 COHERENT
 VOLATILE
 RESTRICT
 READONLY
 WRITEONLY

type_specifier:

type_specifier_nonarray
type_specifier_nonarray array_specifier

array_specifier:

LEFT_BRACKET RIGHT_BRACKET
 LEFT_BRACKET *constant_expression* RIGHT_BRACKET
array_specifier LEFT_BRACKET RIGHT_BRACKET
array_specifier LEFT_BRACKET *constant_expression* RIGHT_BRACKET

type_specifier_nonarray:

VOID
FLOAT
INT
UINT
BOOL
VEC2
VEC3
VEC4
BVEC2
BVEC3
BVEC4
IVEC2
IVEC3
IVEC4
UVEC2
UVEC3
UVEC4
MAT2
MAT3
MAT4
MAT2X2
MAT2X3
MAT2X4
MAT3X2
MAT3X3
MAT3X4
MAT4X2
MAT4X3
MAT4X4
ATOMIC_UINT
SAMPLER2D
SAMPLER3D
SAMPLERCUBE
SAMPLER2DSHADOW
SAMPLERCUBESHADOW
SAMPLER2DARRAY
SAMPLER2DARRAYSHADOW
SAMPLERBUFFER
SAMPLER2DMSARRAY
SAMPLERCUBEARRAY
SAMPLERCUBEARRAYSHADOW
ISAMPLER2D
ISAMPLER3D
ISAMPLERCUBE
ISAMPLER2DARRAY
ISAMPLERBUFFER

ISAMPLER2DMSARRAY
 ISAMPLERCUBEARRAY
 USAMPLER2D
 USAMPLER3D
 USAMPLERCUBE
 USAMPLER2DARRAY
 SAMPLER2DMS
 ISAMPLER2DMS
 USAMPLER2DMS
 SAMPLER2DMSARRAY
 ISAMPLER2DMSARRAY
 USAMPLER2DMSARRAY
 IMAGE2D
 IIMAGE2D
 UIMAGE2D
 IMAGE3D
 IIMAGE3D
 UIMAGE3D
 IMAGECUBE
 IIMAGECUBE
 UIMAGECUBE
 IMAGE2DARRAY
 IIMAGE2DARRAY
 UIMAGE2DARRAY
 IMAGEBUFFER
 IIMAGEBUFFER
 UIMAGEBUFFER
 UIMAGECUBEARRAY
 IMAGECUBEARRAY
 IIMAGECUBEARRAY
 struct_specifier
 TYPE_NAME

precision_qualifier:

HIGH_PRECISION
 MEDIUM_PRECISION
 LOW_PRECISION

struct_specifier:

STRUCT IDENTIFIER LEFT_BRACE *struct_declaration_list* RIGHT_BRACE
 STRUCT LEFT_BRACE *struct_declaration_list* RIGHT_BRACE

struct_declaration_list:

struct_declaration
struct_declaration_list *struct_declaration*

struct_declaration:

type_specifier *struct_declarator_list* SEMICOLON
type_qualifier *type_specifier* *struct_declarator_list* SEMICOLON

```

struct_declarator_list:
    struct_declarator
    struct_declarator_list COMMA struct_declarator

struct_declarator:
    IDENTIFIER
    IDENTIFIER array_specifier

initializer:
    assignment_expression

declaration_statement:
    declaration

statement:
    compound_statement_with_scope
    simple_statement

statement_no_new_scope:
    compound_statement_no_new_scope
    simple_statement

statement_with_scope:
    compound_statement_no_new_scope
    simple_statement

// Grammar Note: labeled statements for SWITCH only; 'goto' is not supported.

simple_statement:
    declaration_statement
    expression_statement
    selection_statement
    switch_statement
    case_label
    iteration_statement
    jump_statement

compound_statement_with_scope:
    LEFT_BRACE RIGHT_BRACE
    LEFT_BRACE statement_list RIGHT_BRACE

compound_statement_no_new_scope:
    LEFT_BRACE RIGHT_BRACE
    LEFT_BRACE statement_list RIGHT_BRACE

statement_list:
    statement
    statement_list statement

```

expression_statement:
 SEMICOLON
expression SEMICOLON

selection_statement:
 IF LEFT_PAREN *expression* RIGHT_PAREN *selection_rest_statement*

selection_rest_statement:
statement_with_scope ELSE *statement_with_scope*
statement_with_scope

condition:
expression
fully_specified_type IDENTIFIER EQUAL *initializer*

switch_statement:
 SWITCH LEFT_PAREN *expression* RIGHT_PAREN LEFT_BRACE *switch_statement_list*
 RIGHT_BRACE

switch_statement_list:
 /* nothing */
statement_list

case_label:
 CASE *expression* COLON
 DEFAULT COLON

iteration_statement:
 WHILE LEFT_PAREN *condition* RIGHT_PAREN *statement_no_new_scope*
 DO *statement_with_scope* WHILE LEFT_PAREN *expression* RIGHT_PAREN SEMICOLON
 FOR LEFT_PAREN *for_init_statement* *for_rest_statement* RIGHT_PAREN
statement_no_new_scope

for_init_statement:
expression_statement
declaration_statement

conditionopt:
condition
 /* empty */

for_rest_statement:
conditionopt SEMICOLON
conditionopt SEMICOLON *expression*

jump_statement:
 CONTINUE SEMICOLON
 BREAK SEMICOLON
 RETURN SEMICOLON
 RETURN *expression* SEMICOLON
 DISCARD SEMICOLON // Fragment shader only.

// Grammar Note: No 'goto'. Gotos are not supported.

```

translation_unit:
    external_declaration
    translation_unit external_declaration

external_declaration:
    function_definition
    declaration

function_definition:
    function_prototype compound_statement_no_new_scope

```

In general the above grammar describes a super set of the GLSL ES language. Certain constructs that are valid purely in terms of the grammar are disallowed by statements elsewhere in this specification.

Rules specifying the scoping are present only to assist the understanding of scoping and they do not affect the language accepted by the grammar. If required, the grammar can be simplified by making the following substitutions:

- Replace *compound_statement_with_scope* and *compound_statement_no_new_scope* with a new rule *compound_statement*
- Replace *statement_with_scope* and *statement_no_new_scope* with the existing rule *statement*.

11 Errors

This section lists errors that must be detected by the compiler or linker.

Lexical errors, including all those found during pre-processing, all grammatical errors and all semantic errors must be reported at compile-time. Errors due to mismatches between stages must be reported at link-time. All other errors, including exceeding resource limits must be reported either at compile-time or link-time.

The error string returned is implementation-dependent.

11.1 Preprocessor Errors

P0001: Preprocessor syntax error

P0002: `#error`

P0003: `#extension` if a required extension `extension_name` is not supported, or if all is specified.

P0005: Invalid `#version` construct

P0006: `#line` has wrong parameters

P0007: Language version not supported

P0008: Use of undefined macro

P0009: Macro name too long

11.2 Lexer/Parser Errors

Grammatical errors occurs whenever the grammar rules are not followed. They are not listed individually here.

G0001: Syntax error

The parser also detects the following errors:

G0002: Undefined identifier.

G0003: Use of reserved keywords

G0004: Identifier too long

G0005: Integer constant too long

11.3 Semantic Errors

S0001: Type mismatch in expression e.g. `1 + 1.0`

S0002: Array parameter must be an integer

- S0003: Conditional jump parameter (**if**, **for**, **while**, **do-while**) must be a boolean
- S0004: Operator not supported for operand types (e.g. `mat4 * vec3`)
- S0005: **?:** parameter must be a boolean
- S0006: 2nd and 3rd parameters of **?:** must have the same type
- S0007: Wrong arguments for constructor
- S0008: Argument unused in constructor
- S0009: Too few arguments for constructor

- S0011: Arguments in wrong order for structure constructor
- S0012: Expression must be a constant expression
- S0013: Initializer for constant variable must be a constant expression

- S0015: Expression must be a constant integral expression

- S0017: Array size must be greater than zero
- S0018: Array size not defined (except for the last element of a shader storage buffer object)

- S0020: Indexing an array with a constant integral expression greater than its declared size
- S0021: Indexing an array with a negative constant integral expression
- S0022: Redefinition of variable in same scope
- S0023: Redefinition of function in same scope
- S0024: Redefinition of name in same scope (e.g. declaring a function with the same name as a struct)
- S0025: Field selectors must be from the same set (cannot mix `xyzw` with `rgba`)
- S0026: Illegal field selector (e.g. using `.z` with a **vec2**)
- S0027: Target of assignment is not an l-value
- S0028: Precision used with type other than **int**, **float** or sampler type
- S0029: Declaring a main function with the wrong signature or return type

- S0031: **const** variable does not have initializer
- S0032: Use of **a type** without a precision qualifier where the default precision is not defined
- S0033: Expression that does not have an intrinsic precision where the default precision is not defined

S0034: Variable cannot be declared invariant

S0035: All uses of invariant must be at the global scope

S0037: L-value contains duplicate components (e.g. `v.xx = q;`)

S0038: Function declared with a return value but return statement has no argument

S0039: Function declared void but return statement has an argument

S0040: Function declared with a return value but not all paths return a value

S0042: Return type of function definition must match return type of function declaration.

S0043: Parameter qualifiers of function definition must match parameter qualifiers of function declaration.

S0045: Declaring an input inside a function

S0046: Declaring a uniform inside a function

S0047: Declaring an output inside a function

S0048: Illegal data type for vertex output or fragment input

S0049: Illegal data type for vertex input (can only use **float**, floating-point vectors, matrices, signed and unsigned integers and integer vectors)

S0050: Initializer for input

S0051: Initializer for output

S0052: Initializer for uniform

S0053: Static recursion present

S0054: Overloading built-in functions not allowed.

S0055: Vertex output with integral type must be declared as flat

S0056: Fragment input with integral type must be declared as flat

S0057: *init-expression* in switch statement must be a scalar integer

S0058: Illegal data type for fragment output

S0059: Invalid layout qualifier

S0060: Invalid use of layout qualifier (e.g. use of *binding* on a non-opaque type)

11.4 Linker

L0001: Global variables must have the same type (including the same names for structure and field names and the same size for arrays) and precision.

L0003: Too many vertex input values

L0004: Too many vertex output values

L0005: Too many uniform values

L0006: Too many fragment output values

L0007: Fragment shader uses an input where there is no corresponding vertex output

L0008: Type mismatch between vertex output and fragment input

L0009: Missing main function for shader

12 Counting of Inputs and Outputs

This section applies to vertex shader outputs and fragment shader inputs.

GLSL ES 3.2 specifies the storage available for vertex shader outputs and fragment shader inputs in terms of an array of 4-vectors. The assumption is that variables will be packed into these arrays without wasting space. This places significant burden on implementations since optimal packing is computationally intensive. Implementations may have more internal resources than exposed to the application and so avoid the need to perform packing but this is also considered an expensive solution.

GLSL ES 3.2 therefore relaxes the requirements for packing by specifying a simpler algorithm that may be used. This algorithm specifies a minimum requirement for when a set of variables must be supported by an implementation. The implementation is allowed to support more than the minimum and so may use a more efficient algorithm and/or may support more registers than the virtual target machine.

Vertex outputs and fragment inputs are counted separately. If statically used in the fragment shader, the built-in special variables (`gl_FragCoord`, `gl_FrontFacing` and `gl_PointCoord`) are included when calculating the storage requirements of fragment inputs.

If the vertex and fragment shaders are linked together, inputs and outputs are only counted if they are statically used within the shader. If the vertex and fragment shaders are each compiled into a separable program, all declared inputs and outputs are counted.¹

For the algorithm used, failing resource allocation for a variable must result in an error.

The resource allocation of variables must succeed for all cases where the following packing algorithm succeeds:

- The target architecture consists of a grid of registers, 16 rows by 4 columns for vertex output and fragment input variables. Each register can contain a scalar value, i.e. a float, int or uint.
- Variables with an explicit location are allocated first. When attempting to allocate a location for other variables, if there is a conflict, the search moves to the next available free location.
- Structures are assumed to be flattened. Each data member is treated as if it were at global scope.
- Variables are packed into the registers one at a time so that they each occupy a contiguous sub-rectangle. No splitting of variables is permitted.
- The orientation of variables is fixed. Vectors always occupy registers in a single row. Elements of an array must be in different rows. E.g. `vec4` will always occupy one row; `float[16]` will occupy one column. Since it is not permitted to split a variable, large arrays e.g. `float[32]` will always fail with this algorithm.

¹ GLSL ES 3.1 does not require the implementation to remove vertex outputs which are not statically used in the fragment shader.

12 Counting of Inputs and Outputs

- Non-square matrices of type `matCxR` consume the same space as a square matrix of type `matN` where `N` is the greater of `C` and `R`. Variables of type `mat2` occupies 2 complete rows. These rules allow implementations more flexibility in how variables are stored. Other variables consume only the minimum space required.
- Arrays of size `N` are assumed to take `N` times the size of the base type.
- Variables are packed in the following order:
 1. **mat4** and arrays of **mat4**.
 2. **mat2** and arrays of **mat2** (since they occupy full rows)
 3. **vec4** and arrays of **vec4**
 4. **mat3** and arrays of **mat3**
 5. **vec3** and arrays of **vec3**
 6. **vec2** and arrays of **vec2**
 7. Scalar types and arrays of scalar types
- For each of the above types, the arrays are processed in order of size, largest first. Arrays of size 1 and the base type are considered equivalent. The first type to be packed will be `mat4[4]`, `mat4[3]`, `mat[2]` followed by `mat4`, `mat2[4]`...`mat2[2]`, `mat2`, `vec4[8]`, `ve4[7]`,...`vec4[1]`, `vec4`, `mat3[2]`, `mat3` and so on. The last variables to be packed will be `float` (and `float[1]`).
- For 2,3 and 4 component variables packing is started using the 1st column of the 1st row. Variables are then allocated to successive rows, aligning them to the 1st column.
- For 2 component variables, when there are no spare rows, the strategy is switched to using the highest numbered row and the lowest numbered column where the variable will fit. (In practice, this means they will be aligned to the x or z component.) Packing of any further 3 or 4 component variables will fail at this point.
- 1 component variables (e.g. floats and arrays of floats) have their own packing rule. They are packed in order of size, largest first. Each variable is placed in the column that leaves the least amount of space in the column and aligned to the lowest available rows within that column. During this phase of packing, space will be available in up to 4 columns. The space within each column is always contiguous in the case where no variables have explicit locations.
- For each type, variables with the 'smooth' property are packed first, followed by variables with the 'flat' property.
- Each row can contain either values with the 'smooth' property or the 'flat' property but not both. If this situation is encountered during allocation, the algorithm skips the component location and continues with the next available location. These skipped locations may be used for other values later in the allocation process.
- There is no backtracking. Once a value is assigned a location, it cannot be changed, even if such a change is required for a successful allocation.

12 Counting of Inputs and Outputs

Example: pack the following types:

```
out vec4 a;           // top left
out mat3 b;           // align to left, lowest numbered rows
out mat2x3 c;         // same size as mat3, align to left
out vec2 d[6];        // align to left, lowest numbered rows
out vec2 e[4];        // Cannot align to left so align to z column, highest
                      // numbered rows
out vec2 f;           // Align to left, lowest numbered rows.
out float g[3]        // Column with minimum space
out float h[2];       // Column with minimum space (choice of 3, any
                      // can be used)
out float i;          // Column with minimum space
```

In this example, the variables happen to be listed in the order in which they are packed. Packing is independent of the order of declaration.

	x	y	z	w
0	a	a	a	a
1	b	b	b	
2	b	b	b	
3	b	b	b	
4	c	c	c	
5	c	c	c	
6	c	c	c	
7	d	d	g	
8	d	d	g	
9	d	d	g	
10	d	d		
11	d	d		
12	d	d	e	e
13	f	f	e	e
14	h	i	e	e
15	h		e	e

Some types e.g. `mat4[8]` will be too large to fit. These always fail with this algorithm.

13 Issues

13.1 Compatibility with OpenGL ES 2.0

How should OpenGL ES 3.0 support shaders written for OpenGL ES 2.0?

Option 1: Retain all GLSL ES 1.0 constructs in the new language.

Option 2: Allow GLSL ES 1.0 shaders to run in the OpenGL ES 3.0 API.

RESOLUTION: Option 2. This minimizes the complexity of the language with only a small increase in system complexity. It also leaves open the option of deprecating the old language in future versions of the API.

13.2 Convergence with OpenGL

How much should GLSL ES be influenced by the GLSL specification?

OpenGL ES 3.0 is principally targeted at mobile devices such as smartphones and tablets. As such, it is expected that the major use-cases will include gaming and user-interfaces. It is to be expected that content will be ported to and from desktop devices.

RESOLUTION: In the absence of any other requirements, GLSL ES 3.1 should follow GLSL 4.x. The main exceptions to this are:

- The specification should adhere to the principle that functionality should not be duplicated.
- Functionality specific to mobile devices (such as reduced precision) can be added.
- Improvements found in later versions of GLSL can be considered for inclusion.

13.3 Numeric Precision

Should the Open GL ES 2.0 precision requirements be increased?

Most current implementations support a subset of IEEE 754 32-bit floating point. Many implementations also support reduced precision.

RESOLUTIONS:

- highp float should be specified as a subset of IEEE 754 floating point.
- highp int should be exactly 32 bits.
- lowp and mediump should be retained. Mediump to have increased precision.

Should there be a defined format for mediump?

Option: Yes, this would increase portability and encourage the use of mediump on mobile devices.

Option: No, this would be expensive to implement on devices that do not natively support it.

RESOLUTION: No. The specification should allow efficient implementation of medium float on 16-bit floating point hardware but must also be implementable on devices which only natively support 32-bit floating point.

Should the fragment shader have a default precision?

Vertex shaders have a default high precision because lower precisions are not sufficient for the majority of graphics applications. However, many fragment shader operations do not benefit from high precision and developers should be encouraged to use lower precision where possible as this may increase performance or reduce power consumption. In particular, blend operations normally only require low precision and many texture address calculations can be performed at medium precision.

However OpenGL ES may also be used in higher performance devices where the benefit is limited. Therefore there appears to be no single precision that would be applicable to all situations.

RESOLUTION: No, there will be no default precision for fragment shaders.

13.4 Floating Point Representation and Functionality

Should IEEE 754 representation be mandated?

The internal format used by an implementation might not be visible to an application so it is meaningless to specify this. Certain functionality IEEE 754 must be present though.

RESOLUTION: In general, highp float must behave *as if* it is in IEEE 754 format.

Which features should be mandated?

Most of the IEEE 754 is relatively inexpensive to implement given that 32-bit floating point is a requirement. However some implementations do not implement signed zeros, rounding modes and NaNs because of hardware cost. In addition, there are certain compiler optimizations that the IEEE 754 specification prohibits.

RESOLUTION: Mandate support of signed infinities. Support of signed zeros, NaNs.

Should the support of NaNs be consistent?

Should the specification allow either full IEEE NaN support or no support but nothing in between?

RESOLUTION: No, implementations may have partial support and there is no guarantee of consistency. The only requirement is that `isnan()` must return false if NaNs are not supported.

Should subnormal numbers (also known as 'denorms') be supported?

RESOLUTION: No, subnormal numbers may be flushed to zero at any time.

How should the rounding mode be specified?

Most current implementations support round-to-nearest. Some but not all also support round-to-nearest-even.

RESOLUTION: Within the accuracy specification, the rounding mode should be undefined.

Should there be general invariance rules for numeric formats and operations?

The GLSL ES specification allows the implementation a degree of flexibility. Consequently the results of a computation may be different on different implementations. However, it is not stated whether a single implementation is allowed to vary the results of a given computation, either in different shaders or different parts of the same shader. OpenGL has a general invariance rule that prevents the results of a computation varying if no state (including the choice of shader) is unchanged.

RESOLUTION: Operations and formats are in general considered to be variant.

13.5 Precision Qualifiers

Should the precisions be specified as float16, float32 etc.? This would help portability. It implies different types rather than hints. It will require all implementations to use the same or similar algorithms and reduces the scope for innovation.

RESOLUTION: No, the precision should not specify a format. Standardized arithmetic is not (yet) a requirement for graphics.

Do integers have precision qualifiers? OpenGL ES 3.0 hardware is expected to have native integer support and some implementations may have reduced precision available.

RESOLUTION: Yes, integers have precision qualifiers.

How should wrapping behavior of integers be defined? If an application relies on wrapping on one implementation this may cause portability problems.

Option: The standard should specify either wrapping or clamping. This allows for maximum implementation flexibility.

Option: Mandate wrapping. There is a trend towards more complex shaders and developers will expect integers to behave as in C++.

RESOLUTION: Mandate wrapping.

Are precision qualifiers available in the vertex shader?

RESOLUTION: Yes. Reduced precision may be available in the vertex shader in some implementations and it keeps the languages consistent.

Should different precisions create different types and e.g. require explicit conversion between them?

Option 1: No, they are just hints. But hinting high precision is meaningless if the implementation can ignore it.

Option 2: Yes they are different types. But this introduces complexity.

RESOLUTION: The precision qualifier can significantly affect behavior in many implementations.

highp means 32-bit IEEE 743 floating point is used but **mediump** means that at least medium precision is used (and similarly for **lowp**) so precision qualifiers are more than just hints. As far as the language is concerned it doesn't affect the behavior so they can either be considered as hints or as different types with implicit type conversion. In any case, implementations are free to calculate everything at high precision.

Should precisions be considered when resolving function calls?

RESOLUTION: No, they should be considered more as hints. Function declarations cannot be overloaded based on precision.

How should precisions be propagated in an expression?

Option 1: Only consider the inputs to an operation. For operands that have no defined precision, determination of precision starts at the leaf nodes of the expression tree and proceeds to the root until the precision is found. If necessary this includes the l-value in an assignment. Constant expressions must be invariant and it is expected that they will be evaluated at compile time. Therefore they must be evaluated at the highest precision (either **lowp** or **highp**) supported by the target, or above.

Option 2: Always take the target of the expression into account. The compiler should be able to work out how to avoid losing precision.

RESOLUTION: Option 1. This makes it easier for the developer to specify which precisions are used in a complex expression.

What if there is no precision in an expression?

Option 1: Leave this as undefined.

Option 2: Use the default precision.

RESOLUTION: Use the default precision. It is an error if this is not defined (in the fragment shader).

Do precision qualifiers for uniforms need to match?

Option 1: Yes.

Uniforms are defined to behave as if they are using the same storage in the vertex and fragment processors and may be implemented this way.

If uniforms are used in both the vertex and fragment shaders, developers should be warned if the precisions are different. Conversion of precision should never be implicit.

Option 2: No.

Uniforms may be used by both shaders but the same precision may not be available in both so there is a justification for allowing them to be different.

Using the same uniform in the vertex and fragment shaders will always require the precision to be specified in the vertex shader (since the default precision is highp). This is an unnecessary burden on developers.

RESOLUTION: Yes, precision qualifiers for uniforms must match.

Do precision qualifiers for vertex outputs and the corresponding fragment inputs (previously known as 'varyings') need to match?

Option 1: Yes. Varyings are written by the vertex shader and read by the fragment shader so there are no situations where the precision needs to be different.

Option 2: No, the vertex outputs written by the vertex shader should not be considered to be the same variables as those read by the fragment shader (there can be no shared storage). Hence they can be specified to have different precisions.

RESOLUTION Precision qualifiers for vertex outputs and fragment inputs do not need to match.

lowp int

lowp float has a range of +/- 2.0 but **lowp int** has a range of +/- 256. This becomes problematic if conversion from **lowp float** to **lowp int** is required. Direct conversion i.e. **lowp int = int(lowp float)** loses almost all the precision and multiplying before conversion e.g. **lowp int = int(lowp float * 256)** causes an overflow and hence an undefined result. The only way to maintain precision is to first convert to **mediump float**.

Option 1: Keep this behavior. Accept that conversion of **lowp float** to **low int** loses precision and is therefore not useful.

Options 2: Make **lowp int** consistent with **mediump** and **highp int** by setting its range to +/- 1

Options 3: Redefine the conversion of **lowp float** to **lowp int** to include an 8-bit left shift. The conversion of **lowp int** to **lowp float** then contains an 8-bit right shift.

Option 4: Option 1 but add built-in functions to shift-convert between the two formats.

Option 5: Redefine the **lowp float** to be a true floating point format. It would then be equivalent to a floating point value with a 10 bit mantissa and a 3 bit unsigned exponent.

RESOLUTION: Option 1 Conversion will lose most of the precision.

Precision of built-in texture functions.

Most built-in functions take a single parameter and it is sensible for the precision of the return value to be the same as the precision of the parameter. The texture functions take sampler and coordinate parameters. The return value should be completely independent of the precision of the coordinates. How should the precision of the return value be specified?

RESOLUTION: Allow sampler types to take a precision qualifier. The return value of the texture functions have the same precision as the precision of the sampler parameter.

What should the default precision of sampler types be?

Option 1: **lowp**. This will be faster on some implementations. In general, OpenGL ES should default to fast operation rather than precise operation. It is usually easier to detect and correct a functional error than a performance issue.

Option 2: **lowp** for textures that are expected to contain color values. **highp** for textures that are expected to contain other values e.g. depth.

Option 2: No default precision. Although this requires that the precision be specified in every shader, it will force the developer to consider the requirements.

RESOLUTION: The default precision of all sampler types present in GLSL ES 1.0 should also be **lowp** in GLSL ES 3.0. New sampler types in GLSL ES 3.0 should have no default precision.

13.6 Function and Variable Name Spaces

Do variables and functions share the same name space? GLSL ES doesn't support function pointers so the grammar can always be used to distinguish cases. However this is a departure from C++.

RESOLUTION: Functions and variables share the same name space.

Should redeclarations of the same names be permitted within the same scope? This would be compatible with C. There are several cases e.g.:

1. Redeclaring a function. A function prototype is a declaration but not a definition. A function definition is both a declaration and a definition. Consequently a function prototype and a function definition (of the same function) within the same scope qualifies as redeclaration.
2. Declaring a name as a function and then redeclaring it as a structure.
3. Declaring a name as a variable and then redeclaring it as a structure.

Disallowing multiple function declarations (including allowing a separate function prototype and function definition) would prevent static recursion by design. However it imposes constraints on the structure of shaders.

GLSL ES 1.00 allows a single function definition plus a single optional function declaration.

RESOLUTION: Multiple definitions are disallowed. Multiple function declarations (function prototypes) are allowed. This is in line with C++.

13.7 Local Function Declarations and Function Hiding

Should local functions hide all functions of the same name?

This is considered useful if local function declarations are allowed. However, the only use for local function declarations in GLSL ES is to unhide functions that have been hidden by variable or structure declarations. This is not a compelling reason to include them.

RESOLUTION: Disallow local function declarations.

13.8 Overloading main()

Should it be possible for the user to overload the main() function?

RESOLUTION: No. The main function cannot be overloaded.

13.9 Error Reporting

In general which errors *must* be reported by the compiler?

Some errors are easy to detect. All grammar errors and type matching errors will normally be detected as part of the normal compilation process. Other semantic errors will require specific code in the compiler. The bulk of the work in a compiler occurs after parsing so adding some error detection should not increase the total cost of compilation significantly. However, it is expected that development systems will have sophisticated error and warning reporting and it is not necessary to repeat this process for on-target compilers.

RESOLUTION: All grammar, type mismatch and other specific semantic errors as listed in this specification must be reported. Reporting of other errors or warnings is optional.

Should compilers report if maxima are exceeded, even if the implementation supports them? This could aid portability.

RESOLUTION: No, high-end implementations may quite legitimately go beyond the specification in these areas and mandating the use of the extension mechanism would cause needless complexity. Development systems should issue portability warnings.

Should static recursion be detected?

RESOLUTION: Yes, the compiler will normally generate the necessary control flow graph so detection is easy.

13.10 Structure Declarations

Should structures with the same name and same member variables be considered as the same type?

RESOLUTION: No, follow the C++ rules. Variables only have the same type if they have been declared with the same type and not if they have been declared with different types that have the same name. This does not apply to linking (for uniforms and varying) which has its own rules.

Should structure declarations be allowed in function parameters?

RESOLUTION: No, following the previous resolution it would be impossible to call such a function because it would be impossible to declare a variable with the same structure type.

13.11 Embedded Structure Definitions

Should embedded structure definitions be allowed?

e.g.

```
struct S
{
    struct T
    {
        int a;
    } t;
    int b;
};
```

In order to access the constructor, the structure name would have to be scoped at the same level as the outer level structure. This is inconsistent.

Option 1: Disallow embedded structure definitions.

Option 2: Allow embedded structure definitions but accept that the constructor is not accessible.

Option 3: Scope embedded structure names at the same level as the outermost scope name.

RESOLUTION: Remove embedded structure definitions.

13.12 Redefining Built-in Functions

Should it be possible to redefine or overload built-in functions?

There may be some applications where it is useful to redefine the built-in functions but the language does not include the required functionality for all cases. Built-in functions are likely to be efficiently mapped to the hardware. User-defined functions may not be as efficient but may be able to offer greater precision (e.g. for the trig functions). The application may then want access to both the original and new function. Some user-defined functions would benefit from access to the original function. Once the new function has been declared, the original function is hidden so both these use cases are impossible with the current specification.

Option 1: Allow both redefinition and overloading of built-in functions.

Option 2: Disallow redefinition of built-in functions. Allow them to be overloaded. This may be useful where it is required to extend the functionality of a built-in function. However it creates a subtle incompatibility with the desktop:

```
int sin(int x) {return x;}
void main()
{
    float a = sin(1.0); // legal in GLSL ES, not legal in desktop GLSL.
}
```

It is also a potential source of backwards-incompatibility if a future version of the language introduces new overloads.

Option 3: Remove the ability to redefine or overload functions.

RESOLUTION: Disallow both overloading and redefining built-in functions. There is no compelling use case.

13.13 Global Scope

How should the scoping levels for user-defined and built-in names be defined?

GLSL ES 1.00 and most versions of GLSL have a global scope for user-defined functions and variables and a distinct 'outer' scope where the built-in functions reside. This is different from C++. Since GLSL ES 3.00 does not allow the redefinition of built-in functions, a single global scope is sufficient.

RESOLUTION: A single global scope will be used for user-defined and built-in names.

13.14 Constant Expressions

Should user and built-in functions be allowed in constant expressions? e.g.

```
const float a = sin(1.0);
```

The compiler must be able to evaluate all possible constant expressions as they can potentially be used to size arrays and functions resolution is dependent on array size. Compile-time evaluation of built-in functions is expensive in terms of code size. The complexity of compile-time evaluation of user-defined functions is potentially unbounded.

RESOLUTION: Allow built-in functions to be included in constant expressions. Redefinition of built-in functions is prohibited. User-defined functions are not allowed in constant expressions.

13.15 Varying Linkage

In the vertex shader, a particular varying may be either 1) not declared, 2) declared but not written, 3) declared and written but not in all possible paths or 4) declared and written in all paths. Likewise a varying in a fragment shader may be either a) not declared, b) declared but not read, c) declared and read in some paths or d) declared and read in all paths. Which of these 16 combinations should generate an error?

The compiler should not attempt to discover if a varying is read or written in all possible paths. This is considered too complex for OpenGL ES.

The same vertex shader may be paired with different fragment shaders. These fragment shaders may use a subset of the available input varyings. This behavior should be supported without causing errors. Therefore if the vertex shader writes to a varying that the fragment shader doesn't declare or declared but doesn't read then this is not an error.

If the vertex shader declares but doesn't write to a varying and the fragment shader declares and reads it, is this an error?

RESOLUTION: No.

RESOLUTION: The only error case is when a varying is declared and read by the fragment shader but is not declared in the vertex shader.

13.16 `gl_Position`

Is it an error if the vertex shader doesn't write to `gl_Position`? Whether a shader writes to `gl_Position` cannot always be determined e.g. if there is dependence on an attribute.

Option 1: No it is not an error. The behavior is undefined in this case. Development systems should issue a warning in this case but the on-target compiler should not have to detect this.

Option 2: It is an error if the vertex shader does not statically write to `gl_Position`

Option 3: It is an error if there is any static path through the shader where `gl_Position` is not written.

RESOLUTION: No error (option 1). The nature of the undefined behavior must be specified.

13.17 Preprocessor

Is the preprocessor necessary?

Arguments for removing or simplifying the preprocessor:

- The preprocessor is moderately complex to implement. In particular, function-like macros may have arbitrary complexity and require significant resources to compile.
- The C++ standard does not fully specify the preprocessor. In particular, the situations where preprocessor tokens are subject to macro expansion are not fully defined. Neither is the effect of macro definitions encountered during macro expansion.
- Over-use of the preprocessor is a common source of programming errors because there is limited compile-time checking.

Arguments for retaining the preprocessor:

- The extension mechanism relies on the preprocessor so this would need to be replaced.
- The **`#define`**, **`#ifdef`**, **`#ifndef`**, **`#elif`** and **`#endif`** constructs are commonly used for managing different versions and for include guards.
- There is no template mechanism in GLSL ES so macros are often used instead.

GLSL ES 1.00 removed token pasting and other functionality.

RESOLUTION: Keep the basic preprocessor as defined in the GLSL ES 1.00 specification.

13.18 Character set

GLSL ES 1.00 only allowed a subset of the ascii character set to be used in shaders. That included names and comments. The written languages of many countries include other characters or use a completely different character set. This makes it difficult or impossible to write comments in those languages.

Where should the new characters be allowed? It would be possible to decide independently for comments, identifiers and macros. For macros, they could be allowed as part of macro definitions but prohibited in the final output of macro expansion.

RESOLUTION: The new characters are only allowed inside comments.

Which character set should be used to define the new characters.

UTF-8 has the advantage that it is backwards-compatible with ASCII. All ASCII characters are valid UTF-8 single-byte characters and UTF-8 multi-byte characters all have the highest bit set to '1' in each byte. The disadvantage is that UTF-8 is variable length.

RESOLUTION: UTF-8

How should the extended character set be specified?

Options include full UTF-8 or by explicitly listing the allowed characters.

RESOLUTION: Full UTF-8

Should the compiler check for the presence of invalid UTF-8 byte sequences?

Since any multi-byte characters will only occur within comments and so not required further processing, it would be inexpensive to check for valid UTF-8 characters. Conversely, there appears to be no advantage to doing so. The issue of validity is only of concern to text editors.

RESOLUTION: The compiler must not check for invalid UTF-8 characters. Bytes '0' and newline characters will be interpreted as such wherever they occur.

How does the #version directive interact with the use of UTF-8 in comments?

Following C++, the 'phases of translation' specification defines comment processing to be performed before macro directives are processed. However UTF-8 is legal in GLSL ES 3.00, identified by #version 300 but not in GLSL ES 1.00, identified by #version 100 (or by absence of a #version directive). Therefore the #version behavior in GLSL ES 1.00 would require compilation to be dependent on a directive occurring later in the shader source.

Option: The shader is processed in 2 passes. The first determines the shader version and the second performs compilation as before.

Option: Replace the current version directive mechanism with a byte or character sequence that must always occur at the start of the shader. This is similar to other standards that have multiple versions e.g. HTTP.

Option: Make UTF-8 characters an optional feature of GLSL ES 1.00

RESOLUTION: Replace the version directive in GLSL ES 1.0 with a character sequence that must always occur at the start of the shader.

13.19 Line Continuation

Should the line continuation character '\ be included in the specification?

Line continuation was deliberately excluded from previous versions of GLSL and GLSL ES in order to discourage excessive use of the preprocessor. However, function-like macros are commonly used because there is no 'template' mechanism, which would allow functions to be parametrized by a type. Long macro definitions are therefore not uncommon and the line-continuation character may aid readability.

Given that shader source is stored in a list of character strings, the newline character can be omitted and this has the same effect as a newline followed by a line-continuation.

RESOLUTION: Include line-continuation.

How does this interact with #version?

RESOLUTION: Same issue as with UTF-8 in general. Line-continuation to be made optional in GLSL ES 1.00

13.20 Phases of Compilation

Should the preprocessor run as the very first stage of compilation or after conversion to preprocessor tokens as with C/C++?

The cases where the result is different are not common.

```
#define e +1
int n = 1e;
```

According to the c++ standard, '1e' should be converted to a preprocessor token which then fails conversion to a number. If the preprocessor is run first, '1e' is expanded to '1+1' which is then parsed successfully.

RESOLUTION: Follow c++ rules.

13.21 Maximum Number of Varyings

How should gl_MaxVaryingFloats be defined? Originally this was specified as 32 floats but currently some desktop implementations fail to implement this correctly. Many implementations use 8 vec4 registers and it is difficult to split varyings across multiple registers without losing performance.

Option 1: Specify the maximum as 8 4-vectors. It is then up to the application to pack varyings. Other languages require the packing to be done by the application. Developers have not reported this as a problem.

Option 2: Specify the maximum according to a packing rule. The developer may use a non-optimal packing so it is better to do this in the driver. Requiring the application to pack varyings is problematic when shaders are automatically generated. It is easier for the driver to implement this.

RESOLUTION: The maximum will be specified according to a packing rule.

Should attributes and uniforms follow this rule?

RESOLUTION: Attributes should not follow this rule. They will be continued to be specified as vec4s.

RESOLUTION: Uniforms should not follow this rule for GLSL ES 3.00. Implementations are expected to virtualize such resources.

Should the built-in special variables (`gl_FragCoord`, `gl_FrontFacing`, `gl_PointCoord`) be included in this packing algorithm? Built-in special variables are implemented in a variety of ways. Some implementations keep them in separate hardware, some do not.

RESOLUTION: Any built-in special variables that are statically used in the shader should be included in the packing algorithm.

Should `gl_FragCoord` be included in the packing algorithm? The x and y components will always be required for rasterization. The z and w components will often be required.

RESOLUTION: `gl_FragCoord` is included in the count of varyings.

How should **mat2** varyings be packed?

Option 1: Pack them as 2x2.

Option 2: Pack them as 4 columns x 1 row. This is usually more efficient for an implementation.

Option 3: Allocate a 4 column x 2 row space. This is inefficient but allows flexibility in how implementations map them to registers.

Option 4: As above but pack 2 **mat2** varyings into each 4 column x 2 row block. Any unpaired **mat2** takes a whole 4x2 block.

RESOLUTION: Option 3

Should **mat3** take 3 whole rows?

This would again allow flexibility in implementation but it wastes space that could be used for floats or float arrays.

RESOLUTION: No, **mat3** should take a 3x3 block.

Should **vec3** take a whole row?

RESOLUTION: No.

Should `gl_MaxVertexUniformComponents` be changed (from desktop GLSL) to reflect the packing rules?

RESOLUTION: Rename `gl_MaxVertexUniformComponents` to `gl_MaxVertexUniformVectors`. Rename `gl_MaxFragmentUniformComponents` to `gl_MaxFragmentUniformVectors`.

13.22 Array Declarations

Unsigned array declarations.

Desktop GLSL allows arrays to be declared without a size and these can then be accessed with constant integral expressions. The size never needs to be declared. This was to support `gl_TexCoord` e.g.

```
varying vec4 gl_TexCoord[];
...
gl_FragColor = texture (tex, gl_TexCoord[0].xy);
```

This allows `gl_TexCoord` to be used without having to declare the number of texture units.

`gl_TexCoord` is part of the fixed functionality so unsigned arrays should be removed for GLSL ES

RESOLUTION: Remove unsigned array declarations.

Which forms of array declarations should be permitted?

```
float a[5];
...
float b[] = a; // b is explicitly size 5
```

or

```
float a[] = float[] (1.0, 2.0, 3.0);
```

RESOLUTION: All above constructs are valid. However, any declaration that leaves the size undefined is disallowed as this would add complexity and there are no use-cases.

13.23 Invariance

How should invariance between shaders be handled?

Version 1.10 of desktop GLSL uses `frtransform()` to guarantee that `gl_Position` can be guaranteed to be calculated the same way in different vertex shaders. This relies on the fixed function that has been removed from ES. It is also very restrictive in that it only allows vertex transforms based on matrices. It does not apply to other values such as those used to generate texture coordinates.

Option 1: Specify all operations to be invariant. No, this is too restrictive. Optimum use of resources becomes impossible for some implementations.

Option 2: Add an invariance qualifier to functions that require invariance. No, this does not work as the inputs to the functions and operations performed on the outputs may not be invariant.

Option 3: Add an invariance qualifier to all variables (including shader outputs).

RESOLUTION: Add an invariance qualifier to variables but permit its use only for outputs from the vertex and fragment shaders. Add a global invariance option for use when complete invariance is required.

Should the invariance qualifier be permitted on parameters to texture functions?

Many algorithms rely on two or more textures being exactly aligned, either within a single invocation of a shader or using multi-pass techniques. This could be guaranteed by using the invariant qualifier on variables that are used as parameters to the texture function.

Using the global invariance pragma also guarantees alignment of the textures. It is not clear whether allowing finer control of invariance is useful in practice. Compilers may revert to global invariance and there may be other specific cases that need to be considered.

RESOLUTION: Use of a variable as a parameter to a texture function does not imply that it may be qualified as invariant.

Do invariance qualifiers for declarations in the vertex and fragment shaders need to match?

Option 1: Only allow invariance declarations on outputs. If a vertex shader output is declared as **invariant**, it implies that the corresponding input to the fragment shader is also invariant.

Option 2: Specify that they must match.

RESOLUTION: Only allow invariant declarations on outputs.

Should this rule apply if the varying is declared but not used?

RESOLUTION: Yes, this rule applies for declarations, independent of usage.

How does this rule apply to the built-in special variables?

Option 1: It should be the same as for varyings. But `gl_Position` is used internally by the rasterizer as well as for `gl_FragCoord` so there may be cases where rasterization is required to be invariant but `gl_FragCoord` is not.

Option 2: `gl_FragCoord` and `gl_PointCoord` can be qualified as invariance if and only if `gl_Position` and `gl_PointSize` are qualified invariant, respectively.

RESOLUTION: Option 1.

Can undefined values be made invariant?

If a type is implemented by a larger native type and due to lack of initialization, a variable of that type has an illegal value, it is possible for variant behavior to occur.

For example suppose a boolean is represented by a 32-bit integer with 'false' represented as 0 and 'true' represented as '1'. If the compiler uses both an 'equals 0' and an 'equals 1' test, the following may occur:

```
bool b; // The implementation sets this to an illegal value e.g. 3

if (b) // implementation tests 'b == 1' which is false
{
    f();
}
else // implementation tests 'b == 0' which is also false
{
    g();
}
```

Neither f() nor g() are executed which is unexpected behavior. Such cases could be made invariant but would for example require the compiler to initialize undefined values which is a performance cost.

RESOLUTION: Undefined values cannot be made invariant. These shaders are malformed and therefore have undefined behavior.

13.24 Invariance Within a shader

How should invariance within a shader be specified?

Compilers may decide to recalculate a value rather than store it in a register (rematerialization). The new value may not be exactly the same as the original value.

Option 1: Prohibit this behavior.

Option 2: Use the invariance qualifier on variables to control this. This is consistent with the desktop.

RESOLUTION: Values within a shader are invariant by default. The invariance qualifier or pragma may be used to make them invariant.

Should constant expressions be invariant? In the following example, it is not defined whether the literal expression should always evaluate to the same value.

```
precision mediump int;
precision mediump float;
const int size = int(ceil(4.0/3.0 - 0.333333));
int a[size];
for (int i=0; i<int(ceil(4.0/3.0 - 0.333333)); i++) {a [i] = i;}
```

Implementations must usually be able to evaluate constant expressions at compile time since they can be used to declare the size of arrays. Hardware may compute a less accurate value compared with maths libraries available in C. It would however be expected that functions such as sine and cosine return similar results whether or not they are part of a constant expression. This suggests that the implementation might want to evaluate these functions only on the hardware. However, there are no situations, even with global invariance, where compile time evaluation and runtime evaluation must match exactly.

RESOLUTION: Yes, constant expressions must be invariant.

13.25 While-loop Declarations

What is the purpose of allowing variable declarations in a while statement?

```
while (bool b = f()) {...}
```

Boolean b will always be true until the point where it is destroyed. It is useful in C++ since integers are implicitly converted to booleans.

RESOLUTION: Keep this behavior. Will be required if implicit type conversion is added to a future version.

A similar issue exists in for-loops. The grammar allows constructs such as

```
for(;bool x = a < b;) ;
```

13.26 Cross Linking Between Shaders

Should it be permissible for a fragment shader to call a function defined in a vertex shader or *vice versa*?

RESOLUTION: No, there is no need for this behavior.

13.27 Visibility of Declarations

At what point should a declaration take effect?

```
int x=1;
{
    int x=2, y=x; // case A
    int z=z;      // case B
}
```

Option 1: The name should be visible immediately after the identifier. Both cases above are legal. In case A, y is initialized to the value 2. This is consistent with c++. For case B, the use case is to initialize a variable to point to itself e.g. void* p = &p; This is not relevant to GLSL ES.

Option 2: The name should be visible after the initializer (if present), otherwise immediately after the identifier. In case A, y is initialized to 2. Case B is an error (assuming no prior declaration of z).

Option 3: The name should be visible after the declaration. In case A, y is initialized to 1. Case B is an error if z is has no prior declaration.

RESOLUTION: Option 2. Declarations are visible after the initializer if present, otherwise after the identifier.

13.28 Language Version

What version number should the language have? This version of the language is based on version 3.30 of the desktop GLSL. However it includes a number of features that are in version 4.20 but not 3.30. The previous version of GLSL ES was version 1.00 so this version could be called version 2.00.

RESOLUTION: Follow the desktop GLSL convention so that the language version matches the API version. Hence this version will be called 3.00

13.29 Samplers

Should samplers be allowed as l-values? The specification already allows an equivalent behavior:

Current specification:

```
uniform sampler2D sampler[8];
int index = f(...);
vec4 tex = texture(sampler[index], xy); // allowed
```

Using assignment of sampler types:

```
uniform sampler2D s;
s = g(...);
vec4 tex = texture(s, xy); // not allowed
```

RESOLUTION: Dynamic indexing of sampler arrays is now prohibited by the specification. Restrict indexing of sampler arrays to constant integral expressions.

13.30 Dynamic Indexing

For GLSL ES 1.00, support of dynamic indexing of arrays, vectors and matrices was not mandated because it was not directly supported by some implementations. Software solutions (via program transforms) exist for a subset of cases but lead to poor performance. Should support for dynamic indexing be mandated for GLSL ES 3.00?

RESOLUTION: Mandate support for dynamic indexing of arrays except for sampler arrays, fragment output arrays and uniform block arrays.

Should support for dynamic indexing of vectors and matrices be mandated in GLSL ES 3.00?

RESOLUTION: Yes.

Indexing of arrays of samplers by constant-index-expressions is supported in GLSL ES 1.00. A constant-index-expression is an expression formed from constant-expressions and certain loop indices, defined for a subset of loop constructs. Should this functionality be included in GLSL ES 3.20?

RESOLUTION: Arrays of all opaque types can now be accessed with dynamically uniform expressions.

13.31 Maximum Number of Texture Units

The minimum number of texture units that must be supported in the fragment shader is currently 2 as defined by `gl_MaxTextureImageUnits = 8`. Is this too low for GLSL ES 3.0?

Option 1: Yes, the number of texturing units is the limiting factor for fragment shaders. The number of texture units was increased from 1 to 2 going from OpenGL ES 1.0 to OpenGL ES 1.1 and increased to 8 for OpenGL ES 2.0

RESOLUTION: Increase to 16

13.32 On-target Error Reporting

Should compilers be required to report any errors at compile time or can errors be deferred until link time?

RESOLUTION: If a program cannot be compiled, on-target compilers are only required to report that an error has occurred. This error may be reported at compile time or link time or both. Development systems must generate grammar errors at compile time.

13.33 Rounding of Integer Division

Should the rounding mode be specified for integer division?

The rounding mode for division is related to the definition of the remainder operator. The important relation in most languages (but not relevant in this version of GLSL ES) is:

$$(a / b) * b + a \% b = a \quad (a \text{ and } b \text{ are integers})$$

Usually the remainder operator is defined to have the same sign as the dividend which implies that divide must round towards zero. (Note that the modulo function is not the same as the remainder function. Modulo is defined to have the same sign as the divisor).

The remainder operator was not part of GLSL ES 1.00, so it was not necessary to specify the rounding mode. In GLSL ES 3.00, the remainder operator is included but the results are undefined if either or both operands are negative.

RESOLUTION: The rounding mode is undefined for this version of the specification.

13.34 Undefined Return Values

If a function is declared with a non-void return type, any return statements within the definition must specify a return expression with a type matching the return type. However if the function returns without executing a return statement the behavior is undefined. Should the compiler attempt to check for these cases and report them as an error?

Example:

```
int f()
{
    // no return statement
}

...

int a = f();
```

Option 1: An undefined value is returned to the caller. No error is generated. This is what most c++ compilers do in practice (although the c++ standard actually specifies 'undefined behavior').

Option 2: There must be a return statement at the end of all function definitions that return a value.

No, this requires statements to be added that may be impossible to execute.

Option 3: A return statement at the end of a function definition is required only if it is possible for execution to reach the end of the function:

E.g.

```
int f(bool b)
{
    if (b)
        return 1;
    else
        return 0;
    // No error. The execution can never reach the end of the function so
    // the implicit return statement is never executed.
}
```

This becomes impossible to determine in the presence of loops.

Option 4: All finite static paths through a function definition must end with a return statement. A static path is a path that could potentially be taken if each branch in the code could be controlled independently.

RESOLUTION: Option 1: The function returns an undefined value.

13.35 Precisions of Operations

Should the precision of operations such as add and multiply be defined?

These are not defined by the C++ standard but it is generally assumed that C++ implementations will use IEEE 754 arithmetic. This is not true for GPUs which generally support only a subset of IEEE 754. In addition, many operations such as the transcendental functions are considered too expensive to implement with more than 10 significant bits of precision. Division is commonly implemented by reciprocal and multiplication.

RESOLUTION: Include a table of precisions for operations.

13.36 Compiler Transforms

What compiler transforms should be allowed?

C++ prohibits compiler transforms of expressions that alter the final result. (Note that C++ allows higher precisions than specified to be used but this is a different issue.) GPUs commonly make use of such transforms, for example when mapping sequential code to vector-based architectures.

RESOLUTION: A specified set of transforms (in addition to those permitted by C++) are allowed.

13.37 Expansion of Function-like Macros in the Preprocessor

When expanding macros, each macro can only be applied once to the original token or any token generated from that token. To implement this, the expansion of function-like macros requires a list of applied macros for each token to be maintained. This is a large overhead.

RESOLUTION: Follow the C++ specification.

What should the behavior be if a directive is encountered during expansion of function-like macros?

This is currently specified as undefined in C++ although several compilers implement the expected behavior.

RESOLUTION: Leave as undefined behavior.

13.38 Should Extension Macros be Globally Defined?

For each extension there is an associated macro that the shader can use to determine if an extension is available on a given implementation. Should this macro be defined globally or should it be defined when the extension is (successfully) enabled?

Both alternatives are usable since attempting to enable an unimplemented extension only results in a warning.

Option 1: Globally defined

```
#ifdef GL_OES_<extension-name>
    #extension GL_OES_<extension-name> : enable
    ...
#endif
```

Option 2: Defined as part of #extension

```
#extension GL_OES_<extension-name> : enable // warning if not available
#ifdef GL_OES_<extension-name>
    ...
#endif
```

RESOLUTION: The macros are defined globally. There should be a warning-free path for all legal cases.

13.39 Minimum Requirements

GLSL ES 1.00 specified a set of minimum requirements that effectively made parts of the specification optional. The purpose was to enable low cost implementations while allowing higher performance devices to expose features without recourse to extensions. That flexibility came at the cost of portability. Should the minimum requirements section be included as part of GLSL ES 3.00?

RESOLUTION: No, except for the section on counting of varyings.

13.40 Packing Functions

These functions are used to pack and unpack a 32-bit bit-vector into various types.

Should the conversions be based on the precision (lowp, mediump, highp)? e.g.

```
highp uint    packFloat2x16(mediump vec2 v);
```

RESOLUTION: No. Since mediump can be implemented using more than 16 bits, packing and then unpacking a mediump value might result in a different value on some platforms but not on others.

Should conversion to and from 8-bit types be supported?

RESOLUTION: No. It is not clear which low precision types to support. e.g. lowp is nominally 10 bit.

Which variant of snorm should be used?

Option 1: The range is [-32768, +32767]. Zero is not representable. Uses all the available values. Sometimes known as the 'attribute snorm format'.

Option 2: The range is [-32767, +32767]. Zero is representable. Does not use all the available values. Sometimes known as the 'texture snorm format'.

RESOLUTION: Option 2. It is important that zero is representable. Option 1 is simpler to implement but this is not considered significant for current hardware. The API specification will be amended to use this format for all snorm to float and float to snorm conversions.

13.41 Boolean logical vector operations

The logical binary operators and (&&), or (||), and exclusive or (^) operate only on two boolean expressions and result in a boolean expression. Should they be extended to operate on boolean vectors?

The 2nd operand is conditionally evaluated for these operators.

```
bvec4 f();
bvec4 g();

f() && g(); // g() gets 'run' for some components but not others.
           // This isn't well defined.
```

RESOLUTION: No, these should not be part of the language.

13.42 Range Checking of literals

Should an error be generated if a literal integer is outside the range of a 32-bit integer?

This can be easily checked by the compiler. However, there is a complication because the literal does not include the minus sign for negative constants. Signed integers can be distinguished from unsigned integers by the 'u' suffix but the value 0x8000000 is only valid if preceded by a unary minus.

Option: Check only that the numeric part of a literal integer (signed or unsigned) is representable by 32 bits.

Option: Include any preceding unary minus and check that the literal is within the range of a signed or unsigned integer as appropriate.

Option: Extend the checking to any constant integral expression.

RESOLUTION: It is an error to have a literal unsigned integer outside the range of a 32-bit integer.

Should this apply to floating-point numbers?

The GLSL spec allows an arbitrary number of digits before the decimal point. It therefore possible for a float literal to have an arbitrarily large number of characters but still be representable e.g.

```
1<1 million zeros>.0e-10000000
```

1. Parsing constraints. Should the number of characters in each field be limited in some way?

1. Should the mantissa be limited to e.g. 16 characters?
2. Should the unsigned part of the mantissa be required to fit into a 32 bit integer?

2. Range checks.

1. If the value is larger than 3.40282347e38, should it be required to return INF? Or return an error?

RESOLUTION: No limit on the number of characters in the mantissa or exponent in a float literal.

RESOLUTION: Values larger than representable in a float 32 must return INF (+ or - as appropriate). Values with a magnitude too small to be representable in a float 32 must return zero.

13.43 Sequence operator and constant expressions

Should the following construct be allowed?

```
float a[2,3];
```

The expression within the brackets uses the sequence operator (',') and returns the integer 3 so the construct is declaring a single-dimensional array of size 3. In some languages, the construct declares a two-dimensional array. It would be preferable to make this construct illegal to avoid confusion.

One possibility is to change the definition of the sequence operator so that it does not return a constant-expression and hence cannot be used to declare an array size.

RESOLUTION: The result of a sequence operator is not a constant-expression.

13.44 Version Directive

The version directive in GLSL ES 1.00 has been found to be unsuitable in cases where certain features of the language specification are changed. The existing mechanism relies on a preprocessor directive but, following the order of operations specified by the 'phases of translation' section in the C++ specification, it is difficult or perhaps impossible to change features of the language that are processed before such directives are invoked. Such features include the introduction of the line-continuation character ('\') and the extension of the character set.

There are several options for an improved version mechanism. All specify the version in the first line of the shader and require that the version directive is followed by a newline.

Option 1: Add a byte sequence to the start of the shader. This would allow any change to be made to the language, including changing the character set. This mechanism is often used in file formats for images.

Option 2: Add a character string sequence to the start of the shader. Define it to appear to be a preprocessor directive e.g.

```
#version 300 es
```

Option 3: As option 2 but allow some flexibility in the format so that extra white-space would still be allowed.

Option 4: As option 2 but use a distinctive non-preprocessor format e.g.

```
version-300-es
```

Option 5: As option 4 but include the characters 'glsl' to aid identification e.g.

```
glsl-version-300-es
```

RESOLUTION: Option 3. The version directive is a string, present as the only non-white-space in the first line of the shader. It is very unlikely that the character set will be changed in an incompatible way from UTF-8 in the future. Option 3 is the closest in appearance to the current mechanism.

13.45 Use of Unsigned Integers

Should functions that can only return a positive value e.g. `textureSize()` and the `length()` method, return signed or unsigned values?

Option 1: Unsigned integer. This allows for some degree of compile-time checking. For example it would be impossible to accidentally access an array element with a negative index in a typical initialization loop such as:

```
float a[5];
for (uint i=0u; i<a.length (); i++)
    a[i] = 0.0;
```

Option 2: Signed integer. This allows greater flexibility in calculating array indices without the need for type conversions e.g.

```
float a[SIZE];
...
int index = a.length() - 3; // Library code. SIZE may not be known when
                           // this code is written
if (index >= 0)             // would not work with an unsigned integer
    f(a[index]);
```

RESOLUTION: Option 2. The principle is that integers that represent values and hence may form part of arithmetic expressions should always be signed, even if it is known that they will always be positive. Values that represent bit vectors should always be unsigned.

The extra checking made available by the use of unsigned integers for values known to be positive is minimal. It would be preferable to include a *range* mechanism in a future version of the language.

13.46 Vertex Attribute Aliasing

Vertex inputs (attributes) can be assigned a location in 3 ways:

- By the location qualifier in the shader
- By the API (BindAttribLocation)
- Automatically by the linker (default if the location is not specified explicitly)

These methods may be mixed e.g. some locations may be defined by the shader and others automatically by the linker.

Option 1: Disallow aliasing. The linker would be required to detect and report any aliasing.

Option 2: Permit aliasing.

Issue: How do inputs with different types alias?

Option 2a: Type conversion is performed

Option 2b: A 'reinterpret cast' is used i.e. the bit pattern is unchanged.

Issue: This is well-defined for highp values but lowp integers, lowp floats and mediump floats have undefined bit representations.

Option 3: Leave undefined. Implementations may choose to detect errors, may convert them according to any of the above methods or may generate arbitrary values.

There are some valid uses for aliasing. An 'uber shader' (i.e. a large shader that consists of multiple selectable smaller shaders) might have too many vertex inputs if they all have unique locations but could map two or more inputs to the same location if it is known that they will not be used within the same shader invocation. However, this technique appears not to be widely used. Furthermore, it risks applications making use of undefined type conversions that may work in some implementations but not others.

RESOLUTION: Aliasing is disallowed and the linker must report an error.

Issue: Under which conditions are two inputs with conflicting locations considered to be aliased?

Option 1: Declared but not referenced.

Option 2: Declared and *statically used*.

Option 3: Declared and not removed by compiler optimization.

In general, the behavior of GLSL ES should not depend on compiler optimizations which might be implementation-dependent. Name matching rules in most languages, including C++ from which GLSL ES is derived, are based on declarations rather than use.

RESOLUTION: The existence of aliasing is determined by declarations present after preprocessing.

13.47 Does a vertex input Y collide with a fragment uniform Y?

If a vertex shader declares

```
in vec3 y;
```

and a fragment shader declares

```
uniform float y;
```

Should this be a link error?

The original intention was that uniforms could be shared across shader stages. Hence there is a single name space for uniforms and uniforms with the same name but in different shaders must have the same type and precision. However, a single name space does not imply a single scope and it is the scope that defines where a name is visible. In the above example, the uniform name 'y' is in the uniform name space and in the global scope of the fragment shader but is not in scope in the vertex shader. The vertex input 'y' exists in the vertex global name space and the vertex global scope.

If the vertex shader had declared a uniform 'y' with type 'vec3', that would be an error.

Within shaders, there is a one-to-one correspondence between (regions of) scopes and name spaces. However, when uniforms are declared, they are conceptually inserted into two name spaces: the global name space of the shader and a separate program-level uniform name space. This does not apply to shader input names which are only inserted into the global scope of the shader. Consequently, there is no conflict between a uniform name declared in the fragment shader and an input name declared in the vertex shader.

RESOLUTION: There is no collision and hence no error in this case.

13.48 Counting Rules for Flat and Smooth Varyings

Should the algorithm assume flat and smooth varyings can be packed into the same **vec4** register?

RESOLUTION: No. Some implementations disallow this.

Does the algorithm need to specify the order of packing of flat and smooth varyings of the same type?

RESOLUTION: Yes. If flat and smooth varyings are interleaved, there is less chance that a subsequent array can be packed without violating the resolution above. Therefore, to make the algorithm deterministic, the order should be specified.

13.49 Array of Arrays: Ordering of Indices

Consider an array of size 2 of an array of size 3 of **float**. C++ and GLSL ES both allow the following syntax:

```
float x[2][3]; // x is an array of size 2 of array of size 3 of float
```

However, GLSL ES also allows an alternative syntax:

```
float [m][n] x;
```

What should the order be in this case?

Option 1:

```
float[3][2] x; // Array of size 2 of array of size 3 of float
```

Each dimension of the size is next to the type which it is sizing. The '[3]' is applied to float. Essentially the whole declaration is read in reverse: A float within an array of size 3 within an array of size 2 defines x. The declaration is associative so that it is equivalent to (float[3])[2]. This would also allow a future typedef construct:

```
typedef float[3] float3;
float3[2] x; // equivalent to float[3][2] x;
```

Note that this causes a contradiction when defining an anonymous formal parameter. The following should mean the same but they don't:

```
void f (float /* p */ [3][2]); // anonymous formal parameter
                                // array of size 3 of array of size 2
void f (float [3][2] /* p */); // anonymous parameter of type float[3][2]
                                // i.e. array of size 2 of array of size 3
```

One of these forms would have to be removed from the language if this option is chosen.

Option 2:

```
float[2][3] x; // Array of size 2 of array of size 3 of float
```

The order of the indices is always the same and is read left to right. This is consistent with the C++ syntax:

```
float (*x)[3] = new float[2][3];
```

Option 3:

Replace the current postfix array types with a prefix notation:

```
[2][3]float x; // Array of size 2 of array of size 3 of float
```

The type specifier is similar to the syntax in other languages where declarations take the form:

var-name is array *size-or-range* of *base-type*

Also note that if the elements in the declaration are 'rotated', the syntax is consistent with C++:

```
[2][3]float x;
[3]float x[2];
float x[2][3];
```

This notation can therefore be classed as 'prefix' (with respect to the type), as opposed to the GLSL 'postfix' notation which has the dimensions following the type.

Option 4:

Disallow sizes in the non-C++ style syntax for the cases where the ordering is ambiguous.

```
float a[2][3]; // Allowed. Standard C++-style syntax
float[3] b;    // Allowed. No ambiguity and ensures backwards compatibility
float[][] d = a; // Allowed. Dimensions are defined by 'a'.

float[2][3] c; // Disallowed
float[2][] d = <initializer>; // Disallowed
```

Anonymous formal parameters must use the C++-style notation:

```
void f (float /* p */ [2][3]);
```

When returning an array of arrays, the size must be deduced from return statements within the function. It is an error if there is no such return statement.

Resolution: Option 2

Note that for all the options above, the arrangement of parameters in the initializer is the same. The following is always an array of size 2 of array of size 3 of float:

```
float[][] (float[] (1.0, 2.0, 3.0), float[] (4.0, 5.0, 6.0));
```

13.50 Precision of Evaluation of Compile-time Expressions

Should compile-time and run-time expressions be evaluated at the same precision? Can this be extended to allow expressions to be invariant, independent of whether they are evaluated at compile-time or run-time?

The precision at which compile-time expressions are evaluated has always been stated to be `highp`. Originally `highp` was specified in terms of a minimum set of requirements so the compiler could evaluate these expressions at single or double precision or any other precision greater or equal to the `highp` minima. The definition of `highp` was subsequently changed to be a subset of IEEE 754.

As currently specified, the run-time precision is not full IEEE so compile-time and run-time evaluation of the same expression cannot be guaranteed to be exact. However, it is still desirable to minimize differences where possible.

Option 1: Mandate that compile-time evaluation must be done at the same precision as run-time evaluation.

This is rejected due to the complexity of emulating hardware (since it is not full IEEE 754).

Option 2: Mandate that compile time evaluation must be done using IEEE single precision.

Standard 'C' libraries are frequently used to implement built-in functions such as transcendentals. Usually these libraries are implemented with double precision.

13.51 Matching of Memory Qualifiers in Function Parameters

Should an exact match of memory access qualifiers be enforced when calling a function?

The qualifiers **`coherent`**, **`volatile`**, **`readonly`** and **`writeln`** subset the allowed behavior when applied to a variable. Therefore adding these qualifiers when calling a function cannot produce unexpected behavior. Conversely, removing them should not be allowed. For example, removing the `coherent` property from a variable declared with **`coherent`** would allow an implementation to perform out of order accesses or even omit accesses altogether. Since there would be no defined ordering, it would be possible for a write to complete a long time after the write was executed. The variable could therefore not be considered coherent at any time after the function is called.

The `restrict` qualifier adds to the allowable behavior by permitting (but not requiring) compiler optimizations that assume the underlying storage is not modified via another variable. It could therefore be considered dangerous to add the `restrict` qualifier when calling a function, especially since the qualification is not visible in the calling code. However, the aliasing is under the control of the developer and there are potential use cases e.g.

```
void copy(restrict image2D from, restrict image2D to);
// Function definition assumes that 'from' and 'to' do not alias.

image2D a;
image2D b; // Never aliases with a
image2D c; // Could alias either a or b, so neither can be declared restrict
copy(a,b); // This will behave correctly. Should it be legal?
```

Resolution: Allow **`restrict`** to be added or removed when calling a function.

Should the built-in function prototypes have parameters qualified with memory access qualifiers such as **`restrict`** or **`coherent`**?

Option 1: Yes and each function has one signature. There is no attempt to overload each function based on the qualifiers. This minimizes the number of functions the compiler needs to implement but either removes some use cases or removes some possible optimizations.

Option 2: Yes but there are multiple overloads for each function. This would make built-ins a special case since functions cannot normally be overloaded based on qualifiers.

Option 3: No. The compiler is able to specialize each function based on the qualification of the actual parameters. This is consistent with the way precision qualifiers are handled.

Resolution: No. The built-in functions are effectively overloaded with all combinations of qualifiers.

14 Acknowledgments

This specification is based on the work of those who contributed to the OpenGL ES 3.0 Language Specification, the OpenGL ES 2.0 Language Specification, and the following contributors to this version:

Acorn Pooley, NVIDIA	Chris Tserng, TI	Hans-Martin Will, Vincent
Alberto Moreira, Qualcomm	Clay Montgomery, TI	Hwanyong Lee, Huone
Aleksandra Krstic, Qualcomm		I-Gene Leong, NVIDIA
Alon Or-bach, Nokia & Samsung	Daniel Kartch, NVIDIA	Ian Romanick, Intel
Andrzej Kacprowski, Intel	Daniel Koch, Transgaming & NVIDIA	Ian South-Dickinson, NVIDIA
Arzhange Safdarzadeh, Intel	Daoxiang Gong, Imagination Technologies	Ilan Aelion-Exch, Samsung
Aske Simon Christensen, ARM	Dave Shreiner, ARM	Inkyun Lee, Huone
Avi Shapira, Graphic Remedy	David Garcia, AMD	Jacob Strm, Ericsson
Barthold Lichtenbelt, NVIDIA	David Jarmon, Vivante	James Adams, Broadcom
Ben Bowman, Imagination Technologies	Derek Cornish, Epic Games	James Jones, Imagination Technologies
Ben Brierton, Broadcom	Dominick Witzak, Mobica	James McCombe, Imagination Technologies
Benj Lipchak, Apple	Eben Upton, Broadcom	Jamie Gennis, Google
Benson Tao, Vivante	Ed Plowman, ARM	Jan-Harald Fredriksen, ARM
Bill Licea-Kane, Qualcomm	Eisaku Ohbuchi, DMP	Jani Vaisanen, Nokia
Brent Insko, Intel	Elan Lennard, ARM	Jarkko Kemppainen, Symbio
Brian Murray, Freescale	Erik Faye-Lund, ARM	
Bruce Merry, ARM		Jeff Bolz, NVIDIA
Carlos Santa, TI	Graham Connor, Imagination Technologies	Jeff Leger, Qualcomm
Cass Everitt, Epic Games & NVIDIA	Graham Sellers, AMD	Jeff Vigil, Qualcomm
Cemil Azizoglu, TI	Greg Roth, NVIDIA	Jeremy Sandmel, Apple
Chang-Hyo Yu, Samsung	Guillaume Portier, Hi Corporation	Jeremy Thorne, Broadcom
Chris Dodd, NVIDIA	Guofang Jiao, Qualcomm	Jim Hauxwell, Broadcom
Chris Knox, NVIDIA		Jinsung Kim, Huone
		Jiyoung Yoon, Huone

14 Acknowledgments

John Kessenich, LunarG	Max Kazakov, DMP	Rune Holm, ARM
Jon Kennedy, 3DLabs	Mika Pesonen, Nokia	Sami Kyostila, Nokia
Jon Leech, Khronos	Mike Cai, Vivante	Sean Ellis, ARM
Jonathan Putsman, Imagination Technologies	Mike Weiblen, Zebra Imaging & Qualcomm	Shereef Shehata, TI
Joohoon Lee, Samsung	Mila Smith, AMD	Sila Kayo, Nokia
JoukoKylmäoja, Symbio	Nakhoon Baek, Kyungpook Univeristy	Slawomir Cygan, Intel
Jrn Nystad, ARM	Nate Huang, NVIDIA	Slawomir Grajewski, Intel
Jussi Rasanen, NVIDIA	Neil Trevett, NVIDIA	Steve Hill, STM & Broadcom
Kalle Raita, drawElements	Nelson Kidd, Intel	Steven Olney, DMP
Kari Pulli, Nokia	Nick Haemel, NVIDIA	Suman Sharma, Intel
Keith Whitwell, VMware	Nick Penwarden, Epic Games	Tapani Palli, Nokia
Kent Miller, Netlogic Microsystems	Niklas Smedberg, Epic Games	Teemu Laakso, Symbio
Kimmo Nikkanen, Nokia	Nizar Romdan, ARM	Tero Karras, NVIDIA
Konsta Karsisto, Nokia	Oliver Wohlmuth, Fujitsu	Timo Suoranta, Imagination Technologies & Broadcom
Krzysztof Kaminski, Intel	Pat Brown, NVIDIA	Tom Cooksey, ARM
Larry Seiler, Intel	Paul Ruggieri, Qualcomm	Tom McReynolds, NVIDIA
Lars Remes, Symbio	Per Wennersten, Ericsson	Tom Olson, TI & ARM
Lee Thomason, Adobe	Petri Talala, Symbio	Tomi Aarnio, Nokia
Lefan Zhong, Vivante	Phil Huxley, ZiiLabs	Tommy Asano, Takumi
Marcus Lorentzon, Ericsson	Philip Hatcher, Freescale & Intel	Wes Bang, Nokia
Mark Butler, Imagination Technologies	Piers Daniell, NVIDIA	YanJun Zhang, Vivante
Mark Callow, Hi Corporation	Pyry Haulos, drawElements	
Mark Cresswell, Broadcom	Piotr Tomaszewski, Ericsson	
Mark Snyder, Alt Software	Piotr Uminski, Intel	
Mark Young, AMD	Rami Mayer, Samsung	
Mathieu Robart, STM	Rauli Laatikainen, RightWare	
Matt Netsch, Qualcomm	Rob Barris, NVIDIA	
Matt Russo, Matrox	Rob Simpson, Qualcomm	
Maurice Ribble, Qualcomm	Roj Langhi, Vivante	

15 Normative References

1. The OpenGL® ES Graphics System Version 3.10
2. The OpenGL® ES Shading Language Versions 1.00, 3.00
3. The OpenGL® Graphics System: A Specification (Versions 3.3 – 4.2)
4. International Standard ISO/IEC 14882:1998(E). Programming Languages – C++
5. International Standard ISO/IEC 646:1991. Information technology - ISO 7-bit coded character set for information interchange
6. The Unicode Standard Version 6.0 – Core Specification
7. IEEE 754-2008. IEEE Standard for Floating-Point Arithmetic